



NexaLearn — AI-Powered E-Learning Platform Backend

A Graduation Project Report Submitted to
Faculty of Engineering, Cairo University
in Partial Fulfillment of the requirements of the degree of
Bachelor of Science in Computer Engineering.

Presented by

Youssef Bahy Youssef
Ali Afifi

Ahmed Kamal
Rufida Kassem

Supervised by

Dr. Lydia Wahid

Department of Computer Engineering

Academic Year: 2025/2026

All rights reserved. This report may not be reproduced in whole or in part,
by photocopying or other means, without the permission of the authors/department.

Abstract

The global shift toward digital education has intensified demand for backend platforms that can orchestrate authentication, structured course delivery, secure payments, media processing, automated assessment, and real-time learner engagement within a single cohesive system. Existing solutions often address these concerns in isolation: commercial Learning Management Systems (LMS) provide turnkey features but limit customisation and data ownership; legacy open-source platforms rely on monolithic designs that resist modern scalability requirements; and generic web application templates omit domain-specific concerns such as enrollment state machines, payment webhook idempotency, video transcoding pipelines, and AI-assisted content generation. This fragmentation leaves educators and engineering teams without a complete, production-ready reference architecture for end-to-end e-learning operations.

This graduation project addresses that gap by designing and implementing **NexaLearn**, a full-featured e-learning backend built with NestJS and TypeScript. The primary objective is to deliver a modular monolith organised according to Domain-Driven Design (DDD), comprising more than fifteen bounded contexts—including Authentication, Courses, Enrollment, Payments (Stripe), Media (Mux video), AI Pipeline (quiz generation via Whisper transcription and large language models), Assessment, Discussion, Progress, Reviews, Certificates, Search, Streak, Notifications, and Instructor Analytics—each encapsulated with clear domain boundaries, application use cases, and infrastructure adapters.

The adopted approach follows Clean Architecture and CQRS-lite separation of commands and queries, complemented by an event-driven integration model and a transactional outbox pattern to guarantee reliable, at-least-once event delivery with exactly-once side effects across contexts. Security is enforced through JSON Web Token (JWT) authentication with refresh-token rotation, role-based access control (RBAC), and HMAC-signed service-to-service calls for internal pipeline communication. Persistence is managed with PostgreSQL and the Prisma ORM; Redis supports caching, session management, and rate limiting; Stripe handles checkout and refund flows; Mux manages direct-to-cloud video upload and transcoding; and AWS S3 stores supplementary assets. The system is containerised with Docker for reproducible deployment.

The resulting platform exposes a comprehensive REST API covering enrollment lifecycles, lesson progress projection, quiz attempts, discussion moderation, certificate generation, full-text course search, and instructor analytics dashboards. Integration testing with Jest and Testcontainers validates critical paths including concurrent outbox processing, payment activation flows, and access-control enforcement. NexaLearn demonstrates that a rigorously structured NestJS backend can unify the functional breadth of a modern e-learning marketplace while preserving maintainability, testability, and operational readiness.

Keywords: E-Learning, NestJS, Domain-Driven Design, Clean Architecture, Transactional Outbox, Stripe, Mux, Artificial Intelligence, PostgreSQL

الملخص

يشهد قطاع التعليم الإلكتروني نمواً متسارعاً يفرض الحاجة إلى منصات خلفية متكاملة قادرة على إدارة المصادقة، ومحتوى الدورات، ومعالجة الفيديو، والمدفوعات، والاختبارات المولدة بالذكاء الاصطناعي، وتفاعل المتعلمين في نظام واحد متماسك. ومع ذلك، تفتقر الحلول المتاحة حالياً — سواء المنصات التجارية المغلقة أو الأنظمة مفتوحة المصدر القديمة أو القوالب العامة — إلى التكامل الكامل بين هذه المتطلبات، مما يحد من قابلية التوسع والتخصيص والجهازية للإنتاج.

يهدف هذا المشروع إلى تصميم وتنفيذ **NexaLearn**، وهي منصة تعليم إلكتروني متكاملة مبنية بإطار عمل **NestJS** بلغة **TypeScript**، وفق مبادئ التصميم القائم على المجال (**DDD**) والهندسة النظيفة (**Clean Architecture**). يتكون النظام من أكثر من خمسة عشر سياقاً محدداً، من بينها: المصادقة، الدورات، التسجيل، المدفوعات (**Stripe**)، الوسائط (**Mux**)، خط أنابيب الذكاء الاصطناعي (توليد الاختبارات عبر **Whisper** ونماذج اللغة الكبيرة)، التقييم، المناقشات، التقدم، المراجعات، الشهادات، البحث، الاستمرارية، الإشعارات، وتحليلات المدربين.

يعتمد النهج المتبع على فصل أوامر **CQRS-lite**، وبنية موجهة بالأحداث، ونقط صندوق الصادر المعاملاتي (**Transactional Outbox**) لضمان تسليم الأحداث بموثوقية عالية. تُطبق المصادقة عبر **JWT** مع تدوير رموز التحديث، والتحكم في الوصول القائم على الأدوار (**RBAC**)، واستدعاءات موقعة بـ **HMAC** بين الخدمات. تُستخدم **PostgreSQL** مع **Prisma ORM** للتخزين، و **Redis** للتخزين المؤقت وإدارة الجلسات، و **Stripe** للمدفوعات، و **Mux** لمعالجة الفيديو، و **AWS S3** للملفات، مع **Docker** للنشر.

أنتج المشروع واجهة **REST API** شاملة تغطي دورة حياة التسجيل، وتتبع التقدم، ومحاولات الاختبارات، والمناقشات، وإصدار الشهادات، والبحث، ولوحات تحليلات المدربين. وتم التحقق من المسارات الحرجة عبر اختبارات تكامل باستخدام **Jest** و **Testcontainers**. يُظهر **NexaLearn** إمكانية بناء منصة تعليم إلكتروني حديثة قابلة للصيانة والاختبار والتوسع ضمن بنية **NestJS** منظمة.

Acknowledgment

We would like to express our sincere gratitude to our supervisor, Dr. Lydia Wahid, for the continuous guidance, constructive feedback, and academic support provided throughout every phase of this graduation project. Their expertise in software engineering and system architecture was instrumental in refining the design decisions, validating the architectural approach, and maintaining the technical rigor expected of a production-grade platform.

We are deeply indebted to the Faculty of Engineering at Cairo University, and particularly to the Department of Computer Engineering, for providing the academic environment, laboratory resources, and intellectual foundation that made this work possible. The curriculum and research culture of the department encouraged us to pursue a solution that is not merely functional, but also scalable, maintainable, and aligned with industry best practices.

Our appreciation extends to all faculty members and teaching assistants who contributed, directly or indirectly, to our education in software design, databases, distributed systems, and project management. We also thank our colleagues and peers whose technical discussions, code reviews, and collaborative debugging sessions helped improve the quality of the NexaLearn implementation.

Finally, we owe a profound debt of gratitude to our families and friends for their patience, encouragement, and unwavering support during the long hours of analysis, development, testing, and documentation required to complete this report.

Contents

Abstract	iii
Abstract (Arabic)	v
Acknowledgment	vi
Table of Contents	vii
List of Figures	xv
List of Listings	xvii
List of Tables	xix
List of Abbreviations	xxi
List of Symbols	xxii
Contacts	xxiii
1 Introduction	1
1.1 Motivation and Justification	2
1.1.1 The Post-COVID E-Learning Market	2
1.1.2 Fragmentation of Existing Open-Source and Template Solutions	3
1.1.3 Architectural Challenges Facing E-Learning Backend Develop-	
ers	4
1.1.4 The AI Imperative in Modern E-Learning	5
1.1.5 Justification of the NexaLearn Approach	6
1.2 Project Objectives and Problem Definition	7
1.2.1 Problem Definition	7
1.2.2 Project Objectives	13
1.3 Project Outcomes	16
1.3.1 Primary Software Artefact: NestJS REST API	16
1.3.2 PostgreSQL Database Schema and Migration History	17
1.3.3 Deployment and Operations Stack	18

CONTENTS

1.3.4	OpenAPI Documentation and Developer Experience	18
1.3.5	Automated Test Suite	18
1.3.6	Architecture and Project Documentation	19
1.3.7	Quantitative Summary of Deliverables	19
1.4	Document Organization	20
1.4.1	Chapter 2: Market Visibility Study	21
1.4.2	Chapter 3: Literature Survey	22
1.4.3	Chapter 4: System Design and Architecture	22
1.4.4	Chapter 5: System Testing and Verification	23
1.4.5	Chapter 6: Conclusions and Future Work	23
1.4.6	References and Appendices	23
1.4.7	Reading Paths for Different Audiences	24
2	Market Visibility Study	25
2.1	Targeted Customers	25
2.1.1	Primary Segment: EdTech Startups and Independent Developers	26
2.1.2	Secondary Segment: Universities and Educational Institutions	27
2.1.3	Tertiary Segment: Corporate Training and L&D Departments	28
2.1.4	Segment Prioritisation and Go-to-Market Implications	28
2.2	Market Survey	29
2.2.1	Udemy for Business and Udemy-Like Marketplace Platforms	29
2.2.2	Teachable and Thinkific	31
2.2.3	Moodle (Open-Source LMS)	32
2.2.4	Coursera for Campus	33
2.2.5	LearnDash (WordPress LMS Plugin)	34
2.2.6	Competitive Comparison Summary	36
2.3	Business Case and Financial Analysis	38
2.3.1	Business Case and Adoption Assumptions	38
2.3.2	Capital Expenditure (Capex)	39
2.3.3	Operating Expenditure (Opex)	39
2.3.4	Revenue Model and Year 1 Projection	40
2.3.5	Twenty-Four-Month Cash Flow Forecast	41
2.3.6	Break-Even Analysis	42
2.3.7	Financial Conclusions	44
3	Literature Survey	45
3.1	Background on Domain-Driven Design and Clean Architecture	45
3.1.1	Domain-Driven Design Fundamentals	45
3.1.2	Clean Architecture Layers	47

CONTENTS

3.1.3	Application to NexaLearn's Bounded Contexts	47
3.2	Background on Event-Driven Architecture and the Outbox Pattern . .	48
3.2.1	Event-Driven Architecture in Distributed Systems	48
3.2.2	The Transactional Outbox Pattern	49
3.2.3	NexaLearn's Outbox Implementation	49
3.2.4	Idempotency and Exactly-Once Processing	49
3.2.5	CQRS-Lite Read Models	50
3.3	Background on AI in E-Learning	51
3.3.1	OpenAI Whisper for Speech Recognition	51
3.3.2	LLM-Based Quiz Generation	51
3.3.3	HMAC-Signed Callback Authentication	52
3.4	Background on Video Streaming Technology	53
3.4.1	HTTP Live Streaming (HLS)	53
3.4.2	Mux Video Platform	53
3.4.3	Video Asset Lifecycle in NexaLearn	54
3.4.4	MP4 Renditions for the AI Pipeline	54
3.5	Comparative Study of Previous Work	55
3.5.1	Dossena et al. (2022) — DDD in E-Learning Platforms	55
3.5.2	Yang et al. (2023) — Microservices for E-Learning	55
3.5.3	Kumar et al. (2023) — LLM Quiz Generation	56
3.5.4	Mehta et al. (2023) — Event-Driven Educational Analytics . . .	56
3.5.5	Synthesis and Gap Analysis	56
3.6	Implemented Approach and Justification	58
3.6.1	Framework Selection: NestJS over Spring Boot and Django . .	58
3.6.2	ORM Selection: Prisma over TypeORM	59
3.6.3	Transactional Outbox over Direct Message Queue Publication	59
3.6.4	Modular Monolith over Microservices	60
3.6.5	Summary of Architectural Decisions	61
4	System Design and Architecture	62
4.1	Overview and Assumptions	62
4.2	System Architecture	63
4.2.1	High-Level Architecture Block Diagram	63
4.2.2	Module Dependency Graph	64
4.2.3	Module Summary	65
4.2.4	Entity-Relationship Diagram	67
4.3	Authentication and Session Management Module	68
4.3.1	Functional Description	68
4.3.2	JWT Token Architecture	69

CONTENTS

4.3.3	Redis Session Store Design	69
4.3.4	Refresh Token Rotation and Reuse Detection	70
4.4	Course Management Module	71
4.4.1	Functional Description	71
4.4.2	DDD Aggregate Design: Course–Module–Lesson Hierarchy	71
4.4.3	Optimistic Concurrency Control	73
4.4.4	Idempotency for Lesson and Module Operations	73
4.5	Media and Video Processing Module	74
4.5.1	Functional Description	74
4.5.2	Mux Integration and VideoAsset State Machine	74
4.5.3	Signed Playback URL Generation	75
4.6	AI Pipeline Integration Module	76
4.6.1	Functional Description	76
4.6.2	HMAC Signature Verification	77
4.6.3	HandleGenerationCallbackUseCase — Exactly-Once Processing	77
4.6.4	Quiz Review Workflow	78
4.7	Assessment and Quiz Module	80
4.7.1	Functional Description	80
4.7.2	Quiz Authoring Aggregate	80
4.7.3	Quiz Attempt Lifecycle	82
4.7.4	Asynchronous Short-Answer Grading	83
4.8	Payment Module	83
4.8.1	Functional Description	83
4.8.2	Stripe Integration and Webhook Handling	83
4.8.3	Payment Intent State Machine	84
4.8.4	Reconciliation and Expiration Services	85
4.9	Enrollment Module	85
4.9.1	Free vs. Paid Enrollment Flow	85
4.9.2	Idempotency and Concurrent Enrollment Protection	86
4.10	Discussion and Q&A Module	88
4.10.1	Thread and Post Aggregate Design	88
4.10.2	Instructor Reply Notifications	89
4.11	Security Design	90
4.11.1	HMAC Service-to-Service Authentication	90
4.11.2	Rate Limiting	91
4.11.3	Audit Logging	91
5	System Testing and Validation	93

CONTENTS

5.1	Testing Setup	93
5.1.1	Runtime and Toolchain	93
5.1.2	Test Environment Configuration	94
5.1.3	Testcontainers Integration	94
5.1.4	HTTP Assertions	95
5.2	Testing Plan and Strategy	96
5.2.1	Testing Principles	96
5.2.2	Test Pyramid Distribution	97
5.2.3	Unit Testing	97
5.2.4	Integration Testing	100
5.2.5	End-to-End Testing	103
5.3	Test Schedule	106
5.4	Comparative Results to Previous Work	107
5.4.1	Key Findings	108
5.4.2	Limitations of the Comparison	108
6	Conclusions and Future Work	109
6.1	Faced Challenges	109
6.1.1	Dual-Write Problem in Event Publishing	109
6.1.2	BigInt Serialisation in JSON	110
6.1.3	Optimistic Concurrency Conflicts in Course Mutations	110
6.1.4	AI Pipeline Callback Idempotency	111
6.1.5	Cross-Bounded-Context Communication Without Tight Coupling	112
6.1.6	Refresh Token Reuse Detection Under Concurrent Requests .	112
6.1.7	Stripe Webhook Exactly-Once Processing	113
6.1.8	Video MP4 Rendition Timing	113
6.1.9	Rate Limiting State Across Multiple Instances	114
6.1.10	Complex Quiz Grading with Async Short-Answer Questions . .	114
6.2	Gained Experience	114
6.2.1	Domain-Driven Design in Practice	115
6.2.2	Event-Driven Architecture and Reliable Messaging	115
6.2.3	Third-Party API Integration	116
6.2.4	Designing for Idempotency and Exactly-Once Semantics . . .	116
6.2.5	Infrastructure and DevOps	117
6.2.6	Team Collaboration and Version Control	117
6.3	Conclusions	117
6.3.1	What Was Achieved	117
6.3.2	Key Innovations	118
6.3.3	Limitations Acknowledged	118

CONTENTS

6.3.4	Validation Against Objectives	119
6.4	Future Work	120
6.4.1	Microservices Extraction	120
6.4.2	WebSocket-Based Real-Time Features	121
6.4.3	Redis-Backed Distributed Rate Limiting	121
6.4.4	OAuth2 Social Login	121
6.4.5	Course Recommendation Engine	121
6.4.6	Multi-Tenant Support	122
6.4.7	Certificate Blockchain Verification	122
6.4.8	Mobile Push Notifications	122
6.4.9	React/Next.js Frontend	123
6.4.10	GraphQL API Layer	123
References		124
A Development Platforms and Tools		126
A.1	Hardware Platforms	126
A.1.1	Development Machines	126
A.1.2	Deployment Server Specifications	126
A.2	Software Tools	127
A.2.1	Node.js 22 and NestJS 11	127
A.2.2	PostgreSQL 16	127
A.2.3	Prisma ORM 6	128
A.2.4	Redis 7	128
A.2.5	Docker and Docker Compose	129
A.2.6	Stripe API	129
A.2.7	Mux Video API	130
A.2.8	AWS S3	130
A.2.9	Resend Email Service	131
A.2.10	Jest and Testcontainers	131
A.2.11	Swagger/OpenAPI	131
B Use Cases		133
B.1	Use Case Descriptions	133
B.1.1	UC-01: Student Signup and Email Verification	133
B.1.2	UC-02: Instructor Course Creation and Publishing	135
B.1.3	UC-03: Student Enrollment in Paid Course	136
B.1.4	UC-04: Video Upload and AI Quiz Generation	138
B.1.5	UC-05: Student Takes Quiz Attempt	140

CONTENTS

B.1.6	UC-06: Instructor Reviews AI-Generated Quiz	142
B.1.7	UC-07: Student Watches Video with Progress Tracking	142
B.1.8	UC-08: Discussion Thread Creation and Instructor Reply	143
B.1.9	UC-09: Certificate Generation After Course Completion	144
B.1.10	UC-10: Instructor Views Revenue Analytics Dashboard	146
C	User Guide	148
C.1	API Authentication Guide	149
C.1.1	Obtaining JWT Tokens	149
C.1.2	Token Storage Best Practices	150
C.2	Instructor Guide	151
C.2.1	Creating a Course	151
C.2.2	Adding Modules and Lessons	151
C.2.3	Uploading Videos	152
C.2.4	Managing AI-Generated Quizzes	153
C.3	Student Guide	153
C.3.1	Enrolling in a Course	153
C.3.2	Watching Lessons and Tracking Progress	154
C.3.3	Taking Quizzes	155
C.3.4	Discussion	155
C.4	Admin Guide	156
C.4.1	Managing Categories	156
C.4.2	Viewing Audit Logs	156
C.4.3	Managing Instructor Applications	157
D	Code Documentation	158
D.1	Key Design Patterns	158
D.1.1	Transactional Outbox Publisher	158
D.1.2	Projection Event Ledger for Idempotency	159
D.1.3	Optimistic Concurrency Utility	160
D.2	Domain Entity Design	161
D.3	Use Case Structure	163
E	Feasibility Study	166
E.1	Technical Feasibility	166
E.2	Operational Feasibility	166
E.3	Economic Feasibility	167
E.3.1	Development Costs	167
E.3.2	Production Infrastructure Costs	167
E.3.3	Licensing	168

CONTENTS

E.4 Schedule Feasibility	168
E.5 Risk Analysis	170

List of Figures

1.1	NexaLearn High-Level System Overview	6
1.2	NexaLearn Bounded Contexts and Integration Relationships	12
1.3	Summary of NexaLearn Project Deliverables and Their Relationships	20
2.1	NexaLearn Target Customer Segments and Value Proposition Alignment	27
2.2	Competitive Positioning: Architectural Openness versus Backend Customisation Depth	37
2.3	Global E-Learning Market by Segment (2025). Data sources: HolonIQ, Statista Global EdTech Market Report, Grand View Research.	37
2.4	Break-Even Analysis: Cumulative Cash Flow over Twenty-Four Months (Figure 2.1)	43
2.5	Revenue Projection: Years 1–3 (Enterprise Support Model). Revenue assumes \$5,000/contract with 50, 70, and 73 cumulative clients respectively; expenses include Capex amortisation, Opex, and staff costs as detailed in Tables 2.3 and 2.4.	44
4.1	NexaLearn System Architecture — Clean Architecture Layers	64
4.2	NexaLearn Module Interaction Diagram	65
4.3	Entity-Relationship Diagram (Core Domain)	68
4.4	Auth Session State Machine	71
4.5	Course Status Lifecycle	74
4.6	Video Asset State Machine	74
4.7	Sequence Diagram: Video Upload, Transcoding & Playback	76
4.8	Sequence Diagram: AI Pipeline — Transcript & Quiz Generation	79
4.9	AI Generation Job State Machine	80
4.10	Quiz Attempt State Machine	82
4.11	Payment Intent State Machine	84
4.12	Enrollment State Machine	87
4.13	Sequence Diagram: Paid Enrollment Flow	88
4.14	Discussion Thread State Machine	90

LIST OF FIGURES

5.1	NexaLearn Testing Pyramid: unit tests form the base (263 tests), integration tests occupy the middle layer (89 tests), and E2E tests sit at the apex (34 tests).	96
5.2	E2E test sequence diagram for the student enrollment flow, showing HTTP calls across Auth, Payment, and Enrollment modules.	104
5.3	Gantt chart visualisation of the testing schedule, showing the parallel execution of unit tests across bounded contexts followed by sequential integration and E2E phases.	107
6.1	Video Asset State Machine with MP4 Rendition Tracking	113
6.2	Skills gained during the NexaLearn project, self-assessed across eight dimensions on a 1–5 scale.	115
6.3	Priority roadmap for future NexaLearn development, showing short-term (months 1–6), medium-term (months 7–12), and long-term (beyond 12 months) initiatives.	123
C.1	Swagger UI documentation at <code>/api/docs</code> showing the Auth module endpoints. The interactive UI allows testing endpoints directly from the browser.	148

List of Listings

3.1	AI-generated quiz JSON schema	52
4.1	Refresh Token Reuse Detection	70
4.2	Course Aggregate Root	71
4.3	Optimistic Concurrency Update	73
4.4	Generating a Signed Mux Playback URL	75
4.5	HMAC Callback Verification	77
4.6	Exactly-Once Callback Processing	78
4.7	QuizEntity Aggregate Root Showing Invariants	80
4.8	Idempotent Stripe Webhook Processing	84
4.9	Concurrent Enrollment Protection	86
4.10	Discussion Thread Aggregate	88
4.11	Generic HMAC Webhook Guard	90
5.1	Jest configuration excerpt for unit and integration tests	93
5.2	Test environment variables	94
5.3	Testcontainers lifecycle for integration tests	94
5.4	E2E application bootstrap	95
5.5	E2E test for AI quiz generation	105
6.1	Outbox event written atomically with business data	109
6.2	BigInt serialisation interceptor	110
6.3	Optimistic concurrency retry utility	111
6.4	Projection event ledger for idempotent callback processing	111
6.5	Atomic refresh token consumption	112
A.1	NexaLearn Docker Compose services	129
C.1	Register a new user	149
C.2	Verify email address	149
C.3	Login and obtain tokens	149
C.4	Response with tokens	149
C.5	Authenticated request	150
C.6	Refresh tokens	150
C.7	Create a course	151
C.8	Add a module	151
C.9	Add a lesson	151

LIST OF LISTINGS

C.10 Request upload URL	152
C.11 Check video status	152
C.12 Attach video to lesson	153
C.13 View AI-generated quiz	153
C.14 Publish quiz	153
C.15 List published courses	153
C.16 View course	153
C.17 Enroll in free course	154
C.18 List enrolled courses	154
C.19 View course curriculum	154
C.20 Access lesson	154
C.21 Update lesson progress	154
C.22 Complete lesson	154
C.23 View quiz	155
C.24 Submit quiz attempt	155
C.25 View graded attempt	155
C.26 Create thread	155
C.27 Post reply	156
C.28 Resolve thread	156
C.29 Create category	156
C.30 Query audit logs	156
C.31 Promote user to instructor	157
D.1 Transactional outbox publisher	158
D.2 Using outbox in a use case	159
D.3 Projection event ledger	159
D.4 Optimistic concurrency retry	160
D.5 QuizEntity aggregate root	161
D.6 HandleGenerationCallbackUseCase	163

List of Tables

3	Project team and supervisor contact information	xxiii
1.1	High-Level Comparison of Platform Categories Relevant to NexaLearn Motivation	4
1.2	Stakeholder Groups and Their Architectural Expectations in NexaLearn	11
1.3	Mapping of Project Objectives to Bounded Contexts and Verification Methods	16
1.4	Quantitative Summary of NexaLearn Project Outcomes	19
1.5	Roadmap of Report Chapters and Primary Focus Areas	21
2.1	Customer Segment Summary and NexaLearn Fit	29
2.2	Competitive Platform Comparison (Table 2.1)	36
2.3	Capital Expenditure (Capex) Breakdown — Month 1	39
2.4	Monthly Operating Expenditure (Opex) Components	40
2.5	Twenty-Four-Month Cash Flow Forecast (USD)	41
3.1	Comparative Summary of Prior Work and NexaLearn Coverage	57
3.2	Architectural Decisions and Justification	61
4.1	Module Summary — Bounded Contexts, Entities, Events, and Dependencies	65
5.1	Test Case Distribution by Bounded Context	97
5.2	Unit Test Coverage Summary	98
5.3	Integration Test Files and Scope	101
5.4	End-to-End Test Scenarios	103
5.5	Testing Schedule Across Development Phases	106
5.6	Testing Methodology Comparison: NexaLearn vs. Open-Source LMS Backends	107
6.1	Project Achievements Summary	118
6.2	Validation Against Project Objectives	120
E.1	Estimated Monthly Infrastructure Costs (Medium Deployment)	168
E.2	Project Timeline and Milestones	169

LIST OF TABLES

E.3 Risk Assessment Matrix 170

List of Abbreviations

Abbreviation	Definition
ACID	Atomicity, Consistency, Isolation, Durability
AI	Artificial Intelligence
API	Application Programming Interface
AWS	Amazon Web Services
CAP	Consistency, Availability, Partition tolerance
CQRS	Command Query Responsibility Segregation
DDD	Domain-Driven Design
DRY	Don't Repeat Yourself
DXA	Domain eXperience Architecture
FK	Foreign Key
GP	Graduation Project
HMAC	Hash-based Message Authentication Code
HTTP	Hypertext Transfer Protocol
JWT	JSON Web Token
LLM	Large Language Model
MVC	Model–View–Controller
OCC	Optimistic Concurrency Control
ORM	Object-Relational Mapping
PK	Primary Key
RBAC	Role-Based Access Control
REST	Representational State Transfer
S3	Simple Storage Service (Amazon S3)
SDK	Software Development Kit
SOLID	Single responsibility, Open–closed, Liskov substitution, Interface segregation, Dependency inversion
SQL	Structured Query Language
SSO	Single Sign-On
TTL	Time-To-Live
UUID	Universally Unique Identifier

List of Symbols

Symbol	Definition
$\mathcal{BC}$	Bounded context (DDD)
$\mathcal{D}$	Domain layer in Clean Architecture
λ	Rate-limiting threshold (requests per window)
σ	Standard deviation (analytics aggregates)
τ	Event processing latency
E	Domain event
H	HMAC digest for service-to-service authentication
I	Infrastructure adapter
P	Aggregate persistence port
Q	Query model / read-side projection
R	Refresh token rotation interval
S	Object storage bucket (AWS S3)
T	Transactional outbox record
U	Use case (application service)
$\parallel$	Database concurrency conflict (optimistic locking)

Contacts

Table 3: Project team and supervisor contact information

Name	Email	Phone
Youssef Bahy Youssef	youssef.bahy@eng.cu.edu.eg	+20 100 000 0000
Ahmed Kamal	ahmed.kamal@eng.cu.edu.eg	+20 100 000 0000
Ali Afifi	ali.afifi@eng.cu.edu.eg	+20 100 000 0000
Rufida Kassem	rufida.kassem@eng.cu.edu.eg	+20 100 000 0000
Dr. Lydia Wahid (Supervisor)	lydia.wahid@eng.cu.edu.eg	+20 100 000 0000

This page is intentionally left blank.

Chapter 1:

Introduction

This chapter introduces **NexaLearn**, a production-oriented e-learning backend platform developed as the graduation project of the Department of Computer Engineering at Cairo University, Faculty of Engineering. NexaLearn is an AI-augmented, commerce-enabled learning platform backend implemented as a NestJS modular monolith in TypeScript, backed by PostgreSQL, Redis, and a suite of managed external services including Stripe, Mux, and AWS S3. The system is organised into sixteen domain bounded contexts and four infrastructure modules, grouped into three high-level capability areas: the *Learning Core* (authentication, courses, enrollment, media, progress, assessment, certificates), *Engagement & Commerce* (discussion, reviews, streaks, notifications, payments, search), and *Platform & Operations* (instructor analytics, administration, audit logging, event infrastructure, AI pipeline orchestration, and shared data access).

The purpose of this chapter is fivefold. First, it establishes *why* such a platform is needed in the current e-learning market and why existing alternatives—commercial SaaS products, legacy open-source LMS platforms, and generic web boilerplates—fail to provide a cohesive reference for modern backend engineering. Second, it defines the *problems* that NexaLearn explicitly sets out to solve, including unreliable asynchronous integrations, architectural erosion across many bounded contexts, financial and AI callback integrity, and privacy-preserving analytics. Third, it states measurable *objectives* that govern design and implementation decisions throughout the project. Fourth, it enumerates the *outcomes*—REST APIs, database schema, deployment stack, OpenAPI documentation, integration tests, and architecture records—that constitute the deliverable artefacts. Fifth, it explains how the remainder of this report is organised so that evaluators and future readers can navigate from market context through literature, detailed design, verification, and conclusions.

NexaLearn is conceived not merely as a collection of REST endpoints, but as a *reference architecture*: a codebase and accompanying documentation that demonstrate how Domain-Driven Design (DDD), Clean Architecture, CQRS-lite read/write separation, event-driven integration via a transactional outbox, and production reliability patterns (idempotency keys, optimistic concurrency control, HMAC-verified webhooks, refresh-token rotation) can be applied to the full vertical stack of an on-

line learning marketplace. A student enrolling in a paid course, watching transcoded video lessons, attempting AI-generated quizzes, participating in moderated discussions, earning a certificate, and leaving a review triggers a coordinated chain of domain events spanning more than half of the bounded contexts in the system. Chapter 1 frames the engineering ambition required to implement such workflows correctly, setting the stage for the technical depth presented in subsequent chapters.

1.1. Motivation and Justification

1.1.1. The Post-COVID E-Learning Market

The global e-learning industry has undergone a structural and likely permanent transformation since the COVID-19 pandemic forced universities, training providers, and enterprises to deliver instruction remotely at unprecedented scale. What was previously treated as a supplementary channel—useful for corporate compliance modules or distance-learning electives—has become a primary delivery mechanism for degree programmes, professional certification, and independent creator economies. Industry analyses value the global e-learning market at approximately \$250–400 billion in the mid-2020s, with compound annual growth rates consistently projected in the double digits through the end of the decade. Growth drivers include corporate upskilling mandates, higher-education digitisation, mobile-first learners in emerging markets, and the proliferation of creator-led course businesses.

Leading commercial platforms illustrate the scale that modern backends must support. **Udemy** reports tens of millions of registered learners and hundreds of millions of course enrolments cumulatively; peak traffic during promotional events stresses payment, entitlement, and content-delivery subsystems simultaneously. **Coursera** partners with universities worldwide and must reconcile academic calendars, verified credentials, and enterprise B2B contracts within shared infrastructure. **LinkedIn Learning**, **edX**, and regional platforms such as **Udacity** and **Pluralsight** similarly combine video libraries, assessments, progress tracking, and monetisation in unified user experiences. Although NexaLearn is scoped as a graduation project rather than a hyperscale production deployment, it is deliberately engineered against the *functional* and *architectural* expectations these platforms establish: stateless horizontally scalable API tiers, durable workflow orchestration, secure payment settlement, entitlement-aware media delivery, and analytics that inform instructor and operator decisions.

The post-COVID era also normalised hybrid instructional models in which asynchronous video modules coexist with synchronous sessions, formative AI-generated quizzes supplement summative exams, and discussion forums extend classroom dialogue. Backend systems must therefore support not only CRUD operations

on course catalogues, but long-running asynchronous pipelines (transcoding, transcription, question generation), fine-grained authorisation (instructor, student, admin, teaching assistant roles), and cross-module eventual consistency without sacrificing user-visible correctness.

1.1.2. Fragmentation of Existing Open-Source and Template Solutions

Despite market maturity, the software engineering community lacks a cohesive, open, and pedagogically transparent backend that unifies the full breadth of e-learning concerns within a single architecturally disciplined codebase. The gap manifests in three distinct categories of alternative, each insufficient as a holistic reference.

Commercial closed platforms. Udemy, Coursera, Teachable, Thinkific, and Kajabi deliver polished end-user experiences but expose only narrow public APIs. Internal service boundaries, saga implementations, event schemas, and consistency models remain trade secrets. Independent engineering teams cannot inspect how enrollment activation is coupled to payment webhooks, how video DRM integrates with progress checkpoints, or how AI-generated assessments are validated before publication. Vendor lock-in and revenue-sharing models further discourage universities and startups seeking data sovereignty.

Legacy open-source LMS platforms. Moodle, Sakai, and Canvas (open-core) evolved over decades as plugin-oriented PHP or Java monoliths. They prioritise feature breadth—forums, gradebooks, SCORM players, attendance plugins—over modular domain clarity. Cross-plugin coupling, shared global state, and inconsistent API layers make it difficult to extract modern patterns such as bounded contexts, transactional outboxes, or typed domain events. Moodle’s architecture, while battle-tested at institutional scale, reflects design decisions from an era before widespread REST microservices, TypeScript type safety, and cloud-native video pipelines.

Specialised micro-templates and generic boilerplates. The NestJS ecosystem offers numerous starter repositories providing JWT authentication, Prisma CRUD, and Docker Compose scaffolding. Separate commercial services address individual verticals: Stripe for payments, Mux or Cloudflare Stream for video, Algolia for search, SendGrid for email. None of these templates demonstrate how fifteen or more bounded contexts should collaborate under shared transactional guarantees, nor do they implement e-learning-specific state machines (pending enrollment awaiting payment, video asset awaiting transcoding, quiz attempt awaiting async grading). Developers assembling these pieces ad hoc frequently produce systems that *appear* functional in demos but fail under duplicate webhook delivery, concurrent course edits, or partial pipeline failures.

Table 1.1 summarises the principal limitation motivating NexaLearn: no widely

available open reference spans authentication, structured course aggregates, payment-backed enrollment, Mux-integrated media, Whisper/LLM AI pipelines, assessment, social learning, certificates, search, streaks, notifications, and instructor analytics within one DDD-governed NestJS codebase.

Table 1.1: High-Level Comparison of Platform Categories Relevant to NexaLearn Motivation

Category		Representative Examples	Primary Strengths	Gap NexaLearn Addresses
Commercial SaaS		Udemy, Coursera, Teachable	Polished UX, scale, content marketplace	No source access; opaque architecture; vendor lock-in
Legacy LMS	OSS	Moodle, Open edX, Sakai	Institutional adoption; rich plugin ecosystems	Monolithic plugin sprawl; dated stacks; weak DDD
Generic starters	NestJS	Community boilerplates	Fast bootstrap; JWT + ORM patterns	No e-learning domain; no outbox; no Stripe/Mux/AI
Vertical APIs	SaaS	Stripe, Mux, OpenAI	Best-in-class single concern	No unified domain model or cross-context reliability

1.1.3. Architectural Challenges Facing E-Learning Backend Developers

Developers who attempt to build e-learning platforms from scratch encounter recurring architectural challenges that introductory web development curricula rarely address. These challenges form the technical core of NexaLearn’s motivation.

Cross-bounded-context workflow orchestration. A representative student journey traverses Search (discovery), Enrollment (intent), Payments (settlement), Courses/Media (consumption), Progress (checkpointing), Assessment (validation), Certificates (credential), Reviews (feedback), Notifications (engagement), and Streak (retention). Each step is owned by a distinct module with its own aggregates and invariants. Naive implementations invoke synchronous service calls across module boundaries, creating availability cascades and tight coupling. Event-driven designs alleviate synchronous coupling but introduce delivery semantics questions: what happens if the notification dispatcher fails after enrollment activates?

Eventual consistency and external callbacks. Video platforms emit webhooks when uploads complete or renditions become playable; Stripe notifies applications of `payment_intent.succeeded` and `charge.refunded`; AI workers callback with transcription segments or generated question banks minutes to hours later. Network re-

tries cause duplicate deliveries; maintenance windows cause delayed out-of-order arrivals. Without idempotent handlers backed by persistent idempotency keys or outbox deduplication, systems exhibit double enrollment activation, orphaned Mux assets detached from lessons, or students charged without course access—failures that erode trust and create support burden.

Concurrency and aggregate integrity. Instructors reorder lessons while students mark progress; admins publish course revisions while enrollment counts feed analytics dashboards. Lost updates on course structures, duplicate quiz attempts, and race conditions on seat-limited offerings require optimistic or pessimistic concurrency strategies encoded at the domain layer, not patched ad hoc in controllers.

Security and compliance surface area. Authentication must resist credential stuffing (rate limiting), token theft (refresh rotation with reuse detection), and privilege escalation (RBAC with instructor-scoped analytics). Payment endpoints must verify webhook signatures. AI pipeline callbacks must authenticate via HMAC to prevent forged quiz injection. Audit trails must record administrative actions without storing secrets or raw payment instrument data.

1.1.4. The AI Imperative in Modern E-Learning

A further motivation distinct from pre-2020 LMS development is the rapid integration of artificial intelligence into educational workflows. Learners expect searchable transcripts of video lectures; instructors seek automated generation of formative multiple-choice and short-answer questions aligned with lesson content; platform operators explore async grading assistance for open-ended responses to scale human review. These capabilities are not cosmetic feature flags—they imply dedicated pipeline orchestration, job state tracking, cost-aware retry policies, content moderation boundaries, and privacy controls over copyrighted or personally identifiable material submitted to third-party models.

Speech-to-text models such as **OpenAI Whisper** produce transcripts that feed retrieval and accessibility features. Large language models (LLMs) synthesise question banks when prompted with transcript segments and instructor constraints (difficulty, count, question types). Such pipelines are inherently *asynchronous* and *non-deterministic*: failures, partial results, and model revisions must be handled gracefully without corrupting published assessment entities. Treating AI as a bolt-on script rather than a first-class **AI Pipeline bounded context** leads to duplicated callback endpoints, inconsistent correlation identifiers, and security gaps when unsigned webhook payloads mutate production quiz data.

NexaLearn addresses this imperative by modelling AI jobs as domain entities with explicit lifecycle states (queued, processing, completed, failed), securing service-to-service calls with HMAC signatures, and applying idempotent callback processing

so retried pipeline notifications never duplicate questions or overwrite instructor edits.




Figure 1.1: NexaLearn High-Level System Overview

Figure 1.1 illustrates NexaLearn’s placement within the broader ecosystem: client applications (web, mobile, admin dashboards) consume a REST API exposed by a modular monolith; internal modules communicate via domain events persisted to a transactional outbox; external integrations include Stripe (payments), Mux (video), AWS S3 (supplementary assets), Resend or equivalent (transactional email), Redis (sessions, cache, rate limits), and external AI workers (Whisper transcription, LLM question generation). The diagram reflects a deliberate *modular monolith* stance—operational simplicity with internal seams ready for future extraction if scale demands.

1.1.5. Justification of the NexaLearn Approach

The justification for undertaking NexaLearn as a graduation project rests on four pillars aligned with both academic evaluation criteria and industry practice.

Reference implementation gap. No widely cited open codebase demonstrates end-to-end e-learning backend composition with DDD, Clean Architecture, transactional outbox, Stripe/Mux/AI integrations, and Testcontainers-backed integration tests in TypeScript. NexaLearn fills this gap for students, instructors evaluating the project, and future teams seeking a reproducible starting point.

Production-grade reliability as a learning objective. Computer Engineering graduates must understand patterns that distinguish demo-quality software from operable systems: idempotency, exactly-once side effects, webhook signature verifi-

cation, optimistic locking, and structured audit logging. NexaLearn embeds these patterns in working code with tests, not merely in slides.

Modular monolith pragmatism. Microservice decomposition introduces distributed tracing, service mesh, and cross-service transaction complexity disproportionate to a graduation project timeline. A well-structured modular monolith delivers clear module boundaries and event-driven decoupling while remaining deployable via Docker Compose—satisfying the faculty requirement that interested readers could replicate the system without a Kubernetes cluster.

Capstone coherence for Cairo University. The project integrates database design, API engineering, security, cloud service integration, AI pipeline orchestration, and software architecture—competencies expected of Computer Engineering graduates. The resulting artefact is defensible in viva examination, extensible by subsequent cohorts, and portfolio-worthy for industry recruitment.

In summary, NexaLearn is motivated by market scale, architectural fragmentation in existing alternatives, the non-trivial reliability demands of async e-learning workflows, and the emerging necessity of AI-augmented content pipelines—and justified as a rigorous, open reference that future engineers can study, deploy, and improve.

1.2. Project Objectives and Problem Definition

1.2.1. Problem Definition

The central research and engineering problem addressed by this graduation project is stated formally below.

Given a multi-stakeholder e-learning platform serving concurrent students, instructors, and administrators, how can a team design and implement a complete, open, production-ready backend that coordinates authentication, structured course management, video media lifecycles, AI-generated educational content, financial transactions, learner engagement features, and privacy-preserving instructor analytics—while preserving domain clarity across fifteen or more bounded contexts, guaranteeing reliable asynchronous integration with external providers, and maintaining ACID consistency for business-critical state transitions?

This overarching problem decomposes into six interrelated sub-problems, each corresponding to observable failure modes in real-world e-learning systems and motivating explicit architectural responses in NexaLearn.

Lack of Complete Vertical Coverage in Open Backends

Most publicly available backend templates and partial LMS implementations excel at isolated vertical slices—user registration with JSON Web Tokens (JWT), file upload to object storage, or a basic quiz engine—but fail to implement the *full lifecycle* of a commercial e-learning product within one cohesive domain model. Consider a single paid enrollment: the system must validate course availability and pricing, create a pending enrollment record, initiate a Stripe PaymentIntent, await a signed webhook confirming payment, atomically activate enrollment, initialise progress projections for all lessons, grant Mux playback entitlements, enqueue welcome notifications, update search-facing enrollment statistics, and record audit events—all while remaining resilient to duplicate client retries and partial failures.

When these concerns are implemented as an undifferentiated “services” layer without bounded contexts, implicit coupling emerges rapidly. Payment services begin mutating enrollment tables directly; progress services bypass assessment prerequisites; search reads denormalised columns that drift when outbox processors lag. The sub-problem is therefore not merely absent features, but the absence of a *ubiquitous language* and *context map* assigning ownership, invariants, and published integration contracts to each concern. NexaLearn responds by defining sixteen domain modules—Auth, Courses, Enrollment, Payments, Media, AI Pipeline, Assessment, Discussion, Progress, Reviews, Certificates, Search, Streak, Notifications, Instructor, and Admin—plus infrastructure modules for Prisma, Events, and Audit.

Unreliable Asynchronous Event Processing

E-learning backends are inherently event-heavy. Representative external event sources include:

- **Mux webhooks** — `video.asset.ready`, upload cancellation, encoding errors;
- **Stripe webhooks** — `payment_intent.succeeded`, `payment_intent.payment_failed`, `charge.refunded`, `checkout.session.completed`;
- **AI pipeline callbacks** — transcription completion, generated question payloads, grading results for short-answer items;
- **Internal domain events** — `EnrollmentActivated`, `LessonCompleted`, `QuizAttemptSubmitted`, `CertificateIssued`.

Network partitions, at-least-once delivery guarantees from cloud providers, and application restarts during consumer processing make duplicate and out-of-order de-

livery *normal* rather than exceptional. Naive “publish to message broker after commit” approaches suffer the dual-write problem: database commits succeed while message publication fails, or vice versa. NexaLearn must achieve *at-least-once transport with exactly-once side effects* without two-phase commit across heterogeneous systems. The chosen pattern—**transactional outbox** with row-level locking via `SELECT FOR UPDATE SKIP LOCKED`—persists events in the same PostgreSQL transaction as domain mutations, then processes them idempotently through dedicated handlers that update projections, dispatch notifications, and trigger downstream workflows.

Architectural Erosion Across Many Bounded Contexts

Domain-Driven Design recommends decomposing complex domains into bounded contexts with explicit upstream/downstream relationships, anti-corruption layers for external models, and context maps documenting integration styles (shared kernel, customer-supplier, conformist, open-host service). In practice, naively creating fifteen or more NestJS modules without enforced dependency rules yields circular imports, anaemic entities that mirror database rows without behaviour, and repository leakage where controllers query foreign context tables directly.

The sub-problem is organisational and technical: maintain Clean Architecture layering (domain innermost; application use cases; infrastructure adapters outermost) across twenty NestJS modules while preserving developer velocity. Dependency constraints require that Courses never import Payments repositories; Discussion never writes Progress tables; Instructor analytics reads through authorised query services rather than ad hoc joins. Without discipline, the modular monolith collapses into the monolithic “ball of mud” it was intended to prevent.

Exactly-Once Semantics for Financial and AI Integrations

Monetary transactions and machine-generated assessments carry elevated correctness, regulatory, and reputational requirements. Failure modes include:

- duplicate Stripe webhook → double enrollment activation or duplicate revenue recognition;
- lost refund webhook → student retains access after chargeback;
- replayed AI callback → duplicate quiz questions or overwritten instructor edits;
- forged unsigned callback → injection of malicious assessment content.

Conversely, legitimately retried webhooks after transient 503 responses must succeed without manual operations intervention. NexaLearn must implement *idempotent, auditable integration endpoints* with Stripe signature verification, HMAC-authenticated AI pipeline callbacks, structured correlation identifiers (`enrollmentId`, `pipelineJobId`, `idempotencyKey`), and persistent deduplication stores recording processed event fingerprints.

Privacy-Preserving Instructor Analytics

Instructors require aggregated insights—module completion funnels, median quiz scores, revenue by course offering, streak participation rates, discussion engagement—to improve content and pricing decisions. Platform administrators require security dashboards (failed login spikes, token reuse alerts). However, indiscriminate SQL joins across `User`, `Enrollment`, `LessonProgress`, and `Payment` tables risk exposing personally identifiable information (PII), enabling re-identification in small cohorts, or allowing Instructor A to infer performance data for courses they do not own.

The sub-problem is to deliver *actionable analytics through dedicated application services* that enforce role-based access control (RBAC), aggregate at safe granularities, parameterise queries by `instructorId` derived from authenticated identity rather than client-supplied filters, and never return row-level student identifiers in instructor-facing revenue or engagement exports unless explicitly authorised for institutional admin roles.

Stakeholder Expectations and Non-Functional Constraints

Beyond functional gaps, the problem includes satisfying diverse stakeholder non-functional requirements simultaneously. Table 1.2 maps primary stakeholders to expectations that constrain architecture.

Table 1.2: Stakeholder Groups and Their Architectural Expectations in NexaLearn

Stakeholder	Primary Concerns	Architectural Implications
Students	Fair access after payment; reliable video playback; accurate grades	Idempotent enrollment activation; signed Mux URLs; atomic quiz finalisation
Instructors	Content control; AI assistance; performance insights	OCC on course aggregates; AI job tracking; scoped analytics queries
Platform admins	Security; auditability; moderation	RBAC; Audit module; admin dashboards; rate limiting
Developers	Maintainability; testability; clear boundaries	DDD modules; outbox tests; OpenAPI docs; Docker Compose
Evaluators	Reproducibility; self-contained documentation	Migrations; integration tests; this report




Figure 1.2: NexaLearn Bounded Contexts and Integration Relationships

Figure 1.2 visualises the sixteen domain bounded contexts and their primary event-driven (solid) and query/API-driven (dashed) relationships. Three high-level groupings—Learning Core, Engagement & Commerce, Platform & Operations—provide cognitive structure without implying separate deployable microservices. The Events module and transactional outbox sit centrally because nearly every context publishes or consumes domain integration events.

Collectively, these sub-problems define a problem space where superficial feature accumulation fails; only explicit domain modelling, reliability patterns, and layered architecture can produce a backend that remains correct under concurrency, trustworthy under duplicate webhooks, and comprehensible to engineers who inherit the codebase after graduation.

1.2.2. Project Objectives

In direct response to the problems defined in Section 1.2.1, NexaLearn pursues eight primary objectives. Each objective is stated as an implementable goal, accompanied by success criteria verifiable through code inspection, automated tests, or API demonstration. Objectives map to bounded contexts and cross-cutting mechanisms described in detail in Chapter 4.

1. Implement a secure, JWT-based authentication system with refresh token rotation and Redis session management.

The Auth bounded context shall deliver complete identity lifecycle management: user signup with email verification tokens, password reset via time-limited secrets, login issuing short-lived access JWTs and opaque refresh tokens, logout invalidating refresh token families, and token refresh with *rotation*—each refresh retires the previous token and detects reuse indicative of theft.

Success criteria: Argon2 password hashing; refresh tokens stored as hashes; Redis-backed session metadata and optional blocklists; global rate limiting via `@nestjs/throttler` with Redis storage; RBAC roles (STUDENT, INSTRUCTOR, ADMIN) enforced by guards on protected routes; integration tests covering login, refresh rotation, and reuse detection.

2. Build a course management system with DDD aggregates supporting modules, lessons, and optimistic concurrency.

The Courses context shall model `Course` as an aggregate root containing ordered `Module` entities, each containing ordered `Lesson` entities with publication states (DRAFT, PUBLISHED, ARCHIVED). Use cases include create/update course metadata, add/reorder modules and lessons, attach prerequisites, and publish revisions visible to enrolled students.

Success criteria: Domain invariants enforced in entities (e.g., cannot publish empty course); optimistic concurrency via `version` column-throws on conflict; domain events (`CoursePublished`, `LessonAdded`) emitted to outbox; instructor ownership validated on every mutation; no direct Prisma calls from controllers.

3. Integrate the Mux video platform for upload, transcoding, and signed playback URLs.

The Media context shall orchestrate Mux direct uploads, persist `VideoAsset` entities tracking provider identifiers and status (AWAITING_UPLOAD, UPLOADED, TRANSCODING, READY, FAILED), process signed Mux webhooks idempotently, as-

sociate ready assets with lessons, and issue signed playback tokens restricted to enrolled students.

Success criteria: Webhook signature verification; idempotent handler keyed by Mux event identifier; playback authorisation integrated with Enrollment/Progress guards; instructor dashboard listing assets with error diagnostics; integration test simulating webhook delivery.

4. **Build an AI pipeline integration for automatic quiz generation and transcript creation using Whisper.**

The AI Pipeline context shall accept instructor-initiated jobs bound to lesson video assets, submit HMAC-signed requests to external transcription (Whisper) and LLM question-generation workers, track job status transitions, and process asynchronous callbacks idempotently to create or update Assessment entities (transcripts, draft questions) without duplicating content on retries.

Success criteria: PipelineJob entity with lifecycle states; correlation IDs linking jobs to lessons and instructors; callback authentication; instructor review gate before student-visible quiz publication; failure records with retry policy.

5. **Implement a full payment flow using Stripe with webhook handling and idempotency.**

The Payments and Enrollment contexts shall cooperate in a documented saga: create PENDING enrollment, initiate Stripe PaymentIntent or Checkout Session with metadata (enrollmentId, courseId, userId), transition to ACTIVE upon verified payment_intent.succeeded webhook, support refunds and expiration of unpaid pending enrollments, and persist idempotency keys on all mutating client endpoints.

Success criteria: Stripe webhook signature validation; processed-event deduplication table; enrollment activation exactly once per payment; refund revokes access; Testcontainers integration test covering full payment flow.

6. **Create a discussion and Q&A system with real-time notifications.**

The Discussion context shall provide threaded conversations scoped to lessons, with instructor moderation workflows (PENDING, ACCEPTED, REJECTED post states), @mention parsing, and domain events triggering Notifications for replies and moderation outcomes via email (Resend) and persistent in-app notification records.

Success criteria: Students cannot view unmoderated posts from peers where policy requires acceptance; instructors notified within outbox-driven pipeline; spam rate limits applied; audit trail for moderation decisions.

7. Provide instructor analytics with course performance, revenue breakdown, and student engagement metrics.

The Instructor context shall expose aggregated dashboards per owned course: enrollment counts by status, lesson completion percentages, quiz score distributions (aggregated, not row-level unless admin), gross revenue and refund totals from Payments read models, streak participation rates, and review rating summaries—all scoped by authenticated instructor identity.

Success criteria: No cross-instructor data leakage in queries; response times acceptable via indexed read models or parallel aggregation; admin-only endpoints separated from instructor routes.

8. Ensure production-grade reliability: transactional outbox, idempotency keys, and optimistic locking.

Cross-cutting infrastructure shall persist domain events in an `OutboxMessage` table within the same ACID transaction as state changes; process messages via concurrent workers using `SELECT FOR UPDATE SKIP LOCKED`; mark processed handlers with deduplication keys; apply `Idempotency-Key` HTTP header support on mutating endpoints; employ optimistic locking on contested aggregates in Courses, Enrollment, and Payments.

Success criteria: Integration test demonstrating concurrent outbox workers without duplicate side effects; documented idempotency semantics in OpenAPI; zero known lost-event paths in payment and enrollment flows.

Table 1.3 traces each objective to the primary bounded context(s) and the verification artefact type used in Chapter 5.

Table 1.3: Mapping of Project Objectives to Bounded Contexts and Verification Methods

#	Objective (abbrev.)	Primary Context(s)	Verification
1	JWT auth + Redis sessions	Auth	Integration + E2E tests
2	Course aggregates + OCC	Courses	Unit + integration tests
3	Mux video lifecycle	Media, Enrollment	Webhook integration test
4	AI quiz/transcript pipeline	AI Pipeline, Assessment	Callback idempotency test
5	Stripe payment saga	Payments, Enrollment	Payment flow int-spec
6	Discussion + notifications	Discussion, Notifications	Event handler tests
7	Instructor analytics	Instructor, Payments	Authorisation + query tests
8	Outbox + idempotency	Events (infra)	Concurrent outbox int-spec

Collectively, these objectives define a backend that is feature-complete for a representative e-learning marketplace and exemplifies the non-functional qualities—security, consistency, observability, and maintainability—expected of production systems. Section 1.3 maps objectives to tangible deliverables; Chapters 4 and 5 supply design rationale and empirical verification evidence.

1.3. Project Outcomes

The successful completion of this graduation project yields tangible deliverables demonstrating functional completeness, architectural discipline, and operational readiness. Outcomes span executable software, persistent data models, deployment automation, machine-readable API contracts, automated test suites, and human-readable architecture documentation. Each outcome is traceable to one or more objectives in Section 1.2.2.

1.3.1. Primary Software Artefact: NestJS REST API

The centrepiece deliverable is a fully functional NestJS 10 application written in TypeScript, exposing more than **150 REST endpoints** across **twenty registered modules**: sixteen domain modules (Auth, Courses, Enrollment, Payments, Media, AI Pipeline, Assessment, Discussion, Progress, Reviews, Certificates, Search, Streak, Notifications, Instructor, Admin) and four infrastructure modules (Prisma, Events, Au-

dit, plus shared Common utilities). Controllers remain thin, delegating to application-layer use cases; domain rules live in entities and domain services; infrastructure concerns (Prisma repositories, Stripe client, Mux adapter, email sender) implement ports defined inward.

Representative endpoint families include:

- **Auth** — signup, verify-email, login, refresh, logout, forgot/reset password, profile;
- **Courses** — CRUD courses/modules/lessons, reorder, publish, instructor-scoped listing;
- **Enrollment** — create pending enrollment, activate from payment, cancel, list my enrollments;
- **Payments** — create PaymentIntent, Stripe webhook receiver, refund request;
- **Media** — request Mux upload URL, webhook handler, assign video to lesson, signed playback;
- **Assessment** — create quizzes/questions, start/submit attempts, async grading callbacks;
- **Progress** — mark lesson complete, course progress summary, certificate eligibility;
- **Discussion** — create threads/posts, moderation accept/reject, list by lesson;
- **Search** — full-text course search, category management;
- **Instructor/Admin** — analytics dashboards, audit log queries, user administration.

All public endpoints are documented via Swagger decorators; request DTOs use `class-validator` for input validation; authorisation guards enforce JWT presence and role permissions consistently.

1.3.2. PostgreSQL Database Schema and Migration History

The persistent data model comprises **58 Prisma models** mapped to normalised PostgreSQL tables, including core entities (User, Course, Module, Lesson, Enrollment, Payment, VideoAsset, Quiz, QuizAttempt, LessonProgress, DiscussionThread, Review, Certificate, OutboxMessage, PipelineJob) and supporting join/index tables for RBAC, idempotency records, refresh tokens, and audit events. Foreign-key constraints preserve referential integrity; strategic indexes support search vectors, instructor analytics queries, and outbox polling patterns.

The full **Prisma migration history** chronicles schema evolution from initial bootstrap through incremental feature additions (payments hardening, outbox indexes, AI job tracking). Migrations enable reproducible provisioning via `prisma migrate deploy` in Docker Compose, CI pipelines, and evaluator local setups—satisfying the faculty criterion that an interested reader could rebuild the database layer without manual DDL execution.

1.3.3. Deployment and Operations Stack

NexaLearn ships with **Docker Compose** orchestration defining services for the API container, PostgreSQL 15, Redis 7, and an **Nginx** reverse proxy. Nginx terminates TLS in production-oriented configurations, forwards proxied requests to NestJS, enforces client body size limits for uploads, and applies rate limiting backed by Redis when configured. Environment variables externalise secrets (Stripe keys, Mux tokens, JWT secrets, database URLs) following twelve-factor principles.

The deployment outcome enables evaluators to launch the entire stack with minimal manual dependency installation and provides a credible path toward cloud deployment (e.g., AWS ECS, Railway, or DigitalOcean App Platform) without architectural rework.

1.3.4. OpenAPI Documentation and Developer Experience

NestJS Swagger integration auto-generates an **OpenAPI 3.0 specification** served at `/api/docs`, documenting paths, request/response schemas, authentication schemes (Bearer JWT), and standard error envelopes. This contract reduces frontend integration friction, supports codegen for client SDKs, and serves as a living reference validated against controller decorators during continuous integration.

1.3.5. Automated Test Suite

Verification artefacts include unit tests for domain entities and use-case logic, and **17+ integration test files** leveraging **Testcontainers** to spin ephemeral PostgreSQL instances (and Redis where required). Critical scenarios under automated regression include:

- concurrent outbox processor workers without duplicate side effects;
- Stripe webhook idempotency and enrollment activation;
- payment expiration of pending enrollments;
- progress projection updates upon `LessonCompleted` events;
- course access control denying unenrolled playback;
- Mux webhook asset state transitions;

- auth refresh token rotation and reuse invalidation.

These tests provide confidence that cross-bounded-context workflows—the primary source of production bugs in event-heavy systems—behave correctly under conditions difficult to exercise manually.

1.3.6. Architecture and Project Documentation

Beyond this LaTeX report, the repository includes inline module README fragments, Prisma schema comments, sequence descriptions for payment and AI flows, and TikZ/architecture diagrams. Appendix D provides selected code walkthroughs; Chapter 4 presents the authoritative system design narrative. Together they satisfy the faculty expectation that the project be *self-contained*: a motivated reader could understand, extend, or rebuild a similar platform without oral defence supplementation.

1.3.7. Quantitative Summary of Deliverables

Table 1.4 consolidates quantitative metrics for quick reference by evaluators.

Table 1.4: Quantitative Summary of NexaLearn Project Outcomes

Deliverable	Quantity	Notes
REST API endpoints	150+	Across auth, courses, commerce, media, AI, analytics
NestJS modules	20	16 domain + 4 infrastructure
Bounded contexts	16	Grouped into 3 high-level modules
PostgreSQL models/tables	58	Prisma-managed with FK constraints
Integration test files	17+	Testcontainers-backed PostgreSQL
External integrations	5+	Stripe, Mux, AWS S3, Redis, AI workers
Docker Compose services	4	API, PostgreSQL, Redis, Nginx




Figure 1.3: Summary of NexaLearn Project Deliverables and Their Relationships

Figure 1.3 illustrates dependencies among deliverables: the API implements business logic against the schema; migrations version the schema; Docker Compose deploys API and dependencies; OpenAPI documents the API surface; integration tests validate cross-module behaviour; and this report plus appendices capture design intent. Each outcome closes the loop between motivation (Section 1.1), problem definition (Section 1.2.1), and objectives (Section 1.2.2).

1.4. Document Organization

This report is structured to guide evaluators, future maintainers, and interested engineers from commercial and academic motivation through detailed engineering design, empirical verification, and critical reflection. Each chapter fulfils a distinct role in the narrative arc; together they satisfy the Cairo University Faculty of Engineering graduation project template while tailoring content to the NexaLearn backend domain. Table 1.5 provides a concise roadmap; subsections below expand the purpose and internal organisation of each chapter.

Table 1.5: Roadmap of Report Chapters and Primary Focus Areas

Ch.	Title	Primary Focus
1	Introduction	Motivation, problems, objectives, outcomes, report structure
2	Market Visibility Study	Customers, competitors, business case, financial analysis
3	Literature Survey	DDD, Clean Architecture, technologies, prior work comparison
4	System Design and Architecture	Modules, diagrams, use cases, implementation patterns
5	System Testing and Verification	Test strategy, module/integration results, schedules
6	Conclusions and Future Work	Challenges, experience gained, limitations, extensions
–	References	Numbered citations in order of appearance
A–E	Appendices	Tools, use cases, user guide, code docs, feasibility study

1.4.1. Chapter 2: Market Visibility Study

Chapter 2 positions NexaLearn within the commercial e-learning landscape. It opens with an introductory paragraph framing the market for backend platforms and independent course infrastructure. Section 2.1 identifies **target customers** in three segments: universities and educational institutions requiring data sovereignty and customisation; corporate learning and development (L&D) departments needing SSO integration, compliance tracking, and ROI analytics; and independent course creators or edtech startups seeking white-label backends without revenue-sharing SaaS constraints.

Section 2.2 presents a **market survey** of competing platforms—Udemy for Business, Coursera, Teachable, Thinkific, Moodle, Open edX, and generic NestJS boilerplates—with dedicated subsections analysing strengths, weaknesses, architectural openness, and implications for NexaLearn differentiation. Section 2.3 develops the **business case and financial analysis**, estimating capital expenditure (development salaries, infrastructure setup), operational expenditure (cloud hosting, monitoring, support), projected sales volumes over five years, pricing strategy relative to competitors, monthly cash-flow projections, and the projected break-even date. This chapter demonstrates that NexaLearn is not only technically viable but commercially conceivable if productised.

1.4.2. Chapter 3: Literature Survey

Chapter 3 supplies the theoretical and technological foundations required to understand design decisions in Chapters 4 and 5. An introductory paragraph outlines the chapter’s two-part structure: engineering background, then comparative prior work.

Background sections (e.g., Sections 3.1–3.2) cover pivotal topics including **Domain-Driven Design** (entities, aggregates, bounded contexts, context maps), **Clean Architecture** and dependency inversion, **CQRS-lite** command/query separation, **event-driven architecture**, the **transactional outbox pattern**, and relevant platform technologies: NestJS module system, PostgreSQL ACID semantics, Prisma ORM, Redis caching patterns, Stripe payment lifecycles, Mux video pipelines, and AI speech/LLM integration patterns.

Section 3.3 presents a **comparative study of previous work**, reviewing recent academic publications and industry articles on e-learning platform architecture, microservice reliability, and LMS modernisation from approximately the past three years. Section 3.4 concludes with the **implemented approach**, justifying NexaLearn’s modular monolith over premature microservice decomposition, and documenting adaptations (CQRS-lite rather than full event sourcing, outbox rather than Kafka in v1) appropriate to graduation project scope.

1.4.3. Chapter 4: System Design and Architecture

Chapter 4 is the technical core, answering *what was built* and *how it was built*. It opens with assumptions (single-tenant deployment, English primary locale, external AI provider availability) and non-functional requirements (security, horizontal scalability of stateless API tier, observability via structured logging and audit trails).

Section 4.2 presents **system architecture** through block diagrams, deployment topology, and the three high-level module groupings:

- **Learning Core** — Auth, Courses, Enrollment, Media, Progress, Assessment, Certificates;
- **Engagement & Commerce** — Discussion, Reviews, Streak, Notifications, Payments, Search;
- **Platform & Operations** — Instructor Analytics, Admin, Audit, Events/Outbox, AI Pipeline, Prisma/Common infrastructure.

Subsequent sections (4.3 onward) decompose each major module following the faculty template substructure: *functional description*, *modular decomposition* (domain/application/infrastructure layers), *design constraints*, and supplementary discussion with diagrams, sequence charts, ERD excerpts, and representative TypeScript code snippets. Event flows—enrollment activation, outbox processing, Mux

webhook handling, AI callback ingestion—receive dedicated treatment because they embody cross-cutting reliability concerns introduced in Chapter 1.

1.4.4. Chapter 5: System Testing and Verification

Chapter 5 documents verification strategy and results. Section 5.1 describes the **testing setup**: Docker Compose environments, Testcontainers PostgreSQL/Redis, mocked Stripe/Mux/AI providers where external calls are impractical in CI, and Jest as the test runner. Section 5.2 explains the **testing plan and strategy**, distinguishing unit tests (domain logic), integration tests (repository + outbox + webhook handlers against real database), and end-to-end HTTP tests (full request/response cycles).

Subsections 5.2.1 and 5.2.2 cover **module testing** and **integration testing** respectively, presenting representative test cases, expected versus actual outcomes, and defects discovered during development. Section 5.3 outlines the **testing schedule** aligned with development milestones. Section 5.4 provides **comparative results** where applicable—e.g., outbox throughput under concurrent workers, payment activation latency—in tabular or graphical form with interpretive commentary linking results back to Objective 8 (production-grade reliability).

1.4.5. Chapter 6: Conclusions and Future Work

Chapter 6 synthesises the project narrative. Section 6.1 discusses **faced challenges**: Stripe webhook edge cases (delayed duplicates, metadata mismatches), coordinating Prisma migrations across team members, managing AI pipeline latency and cost, and balancing DDD purity with delivery deadlines. Section 6.2 reflects on **gained experience** in NestJS modular design, event-driven reliability, payment integration security, and technical writing.

Section 6.3 states **conclusions**: the extent to which objectives were met, distinctive features of NexaLearn (unified open reference, outbox-tested reliability, AI pipeline integration), and acknowledged limitations (single-region deployment, no bundled frontend in scope, dependence on external AI providers). Section 6.4 proposes **future work**: multi-tenant isolation, GraphQL adjunct APIs, recommendation engines, SCORM/xAPI import, localised content, Kubernetes deployment manifests, and third-party security audit.

1.4.6. References and Appendices

References follow Chapter 6, numbered in order of first citation throughout the report (faculty template style [1], [2], ...).

Appendix A — Development Platforms and Tools: hardware/software inventory, NestJS, Prisma, Docker, Stripe/Mux SDKs, AI tooling. **Appendix B** — Use Cases: formal use case diagrams and descriptions for student enrollment, instructor pub-

lishing, admin moderation. **Appendix C** — User Guide: step-by-step operational instructions for instructors and administrators with screenshots. **Appendix D** — Code Documentation: selected module walkthroughs and directory structure. **Appendix E** — Feasibility Study: expanded financial projections supporting Chapter 2.

1.4.7. Reading Paths for Different Audiences

Readers may navigate selectively:

- **Faculty evaluators** — Chapter 1 (context), Chapter 4 (design depth), Chapter 5 (verification), Chapter 6 (reflection);
- **Future student maintainers** — Chapter 4, Appendices A and D, plus repository README;
- **Industry reviewers** — Chapter 1 objectives/outcomes, Chapter 3 patterns, Chapter 4 outbox/payment/AI sections;
- **Business stakeholders** — Chapter 2 market study, Appendix E feasibility.

This organisational structure ensures NexaLearn is documented as a complete, reproducible engineering contribution—not an isolated GitHub repository—and closes Chapter 1 by orienting the reader toward the detailed chapters that follow.

Chapter 2:

Market Visibility Study

The e-learning backend market has matured substantially over the past decade, driven by post-pandemic digitisation, corporate upskilling budgets, and the growth of independent creator economies. Yet the market remains bifurcated: commercially successful platforms operate as closed ecosystems with opaque architectures, while open-source alternatives often lack the unified domain modelling, AI-assisted assessment pipelines, and payment-backed enrollment flows that modern products require. This chapter positions NexaLearn within that landscape by identifying the customer segments most likely to adopt an open, DDD-governed NestJS backend; surveying five representative competing platforms across the SaaS, open-source, and plugin-based spectrum; and presenting a business case with capital expenditure, operating expenditure, revenue projections, and a twenty-four-month cash-flow analysis demonstrating commercial viability under an open-core enterprise support model.

The analysis assumes NexaLearn is productised after graduation as an open-source core with optional commercial support—a model familiar from companies such as GitLab, Supabase, and PostHog that monetise expertise, SLAs, and enterprise features while preserving developer adoption through free self-hosted deployment. Financial figures are illustrative projections suitable for feasibility assessment rather than audited forecasts; they are grounded in conservative adoption assumptions, publicly available cloud pricing, and enterprise support price points comparable to specialised backend infrastructure vendors serving the edtech segment.

2.1. Targeted Customers

NexaLearn is architected as a backend *platform* rather than a turnkey consumer-facing application. Its modular monolith design, OpenAPI-documented REST surface, Docker Compose deployment, and explicit bounded contexts make it most valuable to organisations that employ software engineers and require deep customisation, data sovereignty, or integration with existing identity, billing, and content systems. Three customer segments—primary, secondary, and tertiary—emerge from this positioning, each with distinct pain points that NexaLearn’s architecture

addresses directly.

2.1.1. Primary Segment: EdTech Startups and Independent Developers

The **primary target segment** comprises edtech startups, independent software developers, and small product teams building online learning marketplaces, niche certification platforms, or vertical training products (e.g., medical continuing education, creative skills, language tutoring). These customers typically possess TypeScript or Node.js competency but lack the calendar months required to implement production-grade payment sagas, Mux-integrated video lifecycles, transactional outbox event processing, and AI pipeline orchestration from first principles.

Pain points. Generic NestJS boilerplates provide authentication scaffolding but omit enrollment state machines, Stripe webhook idempotency, and instructor analytics. Stitching together Stripe, Mux, SendGrid, and OpenAI APIs without a unified domain model produces fragile glue code that fails under duplicate webhook delivery or concurrent course edits. Venture-backed startups face investor scrutiny on time-to-market; bootstrapped founders face direct opportunity cost for every month spent reinventing outbox patterns instead of differentiating on curriculum or user experience.

How NexaLearn benefits this segment. NexaLearn delivers a *reference implementation* with fifteen bounded contexts, 150+ documented endpoints, Testcontainers integration tests, and architecture documentation sufficient to fork and customise. Teams can deploy via Docker Compose on day one, replace branding and frontend independently, extend the Courses or Assessment contexts for domain-specific rules, and inherit reliability patterns (idempotency keys, optimistic locking, HMAC callbacks) without research overhead. The open-source core eliminates licensing fees during MVP validation; enterprise support becomes relevant only when SLAs, security reviews, or dedicated migration assistance are required at scale.

Representative use cases include a regional coding bootcamp launching a proprietary alumni learning portal, a startup white-labelling course delivery for influencer creators, and an agency building custom LMS backends for multiple clients from a shared NexaLearn fork.

Figure 2.2 — Customer Segments




Figure 2.1: NexaLearn Target Customer Segments and Value Proposition Alignment

Figure 2.1 illustrates how NexaLearn’s architectural strengths—open source, DDD modularity, AI pipeline integration, Stripe/Mux readiness—map most directly to developer-led edtech teams while remaining applicable to institutional and corporate buyers with internal engineering capacity.

2.1.2. Secondary Segment: Universities and Educational Institutions

The **secondary segment** includes universities, faculties, and continuing-education departments seeking a **self-hosted** learning management backend under institutional data-governance policies. Public universities in Egypt and the broader MENA region increasingly require on-premise or sovereign-cloud deployment where student records, payment metadata, and proctored assessment data must not reside on unreviewed third-party SaaS infrastructure subject to foreign jurisdiction.

Pain points. Legacy Moodle deployments satisfy self-hosting requirements but burden IT departments with PHP plugin maintenance, inconsistent security patching, and limited TypeScript talent pools among student developers contributing to customisations. Commercial platforms such as Coursera for Campus offer polished experiences but constrain integration with existing Student Information Systems (SIS), restrict modifications to enrollment logic, and impose per-seat licensing that scales expensively across tens of thousands of registered students.

How NexaLearn benefits this segment. NexaLearn runs entirely on institution-controlled PostgreSQL and Redis instances, exposes REST APIs for SIS integration (enrollment sync, grade export), and supports RBAC roles aligned with academic hi-

erarchies (student, instructor, department admin). The transactional audit module assists compliance officers reviewing administrative actions. AI-assisted quiz generation augments faculty workflows without mandating a specific frontend—faculties may build localised Arabic/English clients atop the API. Institutions with Computer Science departments may further extend NexaLearn as a living capstone platform, as demonstrated by this graduation project.

Adoption pathway. Pilot deployment for a single department (e.g., Computer Engineering online electives), evaluation against Moodle on maintainability and security criteria, and phased rollout with enterprise support covering migration tooling and SLA-backed incident response.

2.1.3. Tertiary Segment: Corporate Training and L&D Departments

The **tertiary segment** comprises mid-to-large enterprises operating internal **white-label** training programmes: employee onboarding, compliance certification (OSHA, ISO, financial regulations), sales enablement, and leadership development. These buyers prioritise SSO integration (SAML/OIDC), branded learner experiences, completion tracking for HR analytics, and cost predictability over marketplace discovery features irrelevant to internal audiences.

Pain points. Enterprise LMS suites (Cornerstone, SAP SuccessFactors) involve multi-year contracts, lengthy implementation cycles, and limited API access for custom mobile apps. Lightweight SaaS tools (Teachable, Thinkific) target creators rather than IT departments and lack SSO, audit trails, and horizontal scaling patterns for workforces numbering tens of thousands. Building in-house from scratch duplicates NexaLearn’s sixteen-context scope unless engineering teams exceed twenty back-end developers.

How NexaLearn benefits this segment. NexaLearn provides white-label readiness: self-hosted deployment behind corporate VPN or private cloud, custom domain and email templates via Notifications context, instructor analytics aggregated for L&D managers without exposing individual performance to unauthorised managers (privacy-preserving query services described in Chapter 1), and Streak/Progress modules supporting compliance deadlines. The open-core model allows enterprises to audit source code during security review—a requirement absent from closed SaaS alternatives. Enterprise support at \$5,000/year per deployment instance is orders of magnitude below enterprise LMS licensing for comparable customisation freedom.

2.1.4. Segment Prioritisation and Go-to-Market Implications

Table 2.1 summarises segment prioritisation for commercialisation. The primary edtech developer segment drives open-source adoption metrics (stars, forks, com-

munity contributions) that reduce customer acquisition cost for enterprise support sales to secondary and tertiary buyers.

Table 2.1: Customer Segment Summary and NexaLearn Fit

Segment	Priority	Decision Maker	Key NexaLearn Advantage
EdTech startups & indie developers	Primary	CTO / lead developer	Full source access; DDD reference; fast Docker deploy
Universities & institutions	Secondary	IT director / dean	Self-hosted; audit trail; SIS API integration
Corporate L&D	Tertiary	HR IT / L&D manager	White-label; SSO-ready; compliance analytics

In summary, NexaLearn’s architecture is optimised for technically sophisticated buyers who value *control, clarity, and extensibility* over turnkey drag-and-drop simplicity—a deliberate trade-off that aligns with the platform’s graduation project origins and open-source commercialisation strategy developed in Section 2.3.

2.2. Market Survey

This section surveys five representative platforms spanning corporate marketplace models, creator SaaS tools, open-source LMS software, institutional partnerships, and WordPress plugin ecosystems. Each subsection follows a consistent analytical frame: platform description, publicly known technical architecture, strengths (*pros*), weaknesses relative to NexaLearn (*cons*), and strategic implications for positioning. The survey is based on publicly available documentation, vendor pricing pages, and industry analyses current to the 2025–2026 academic reporting period; proprietary internal architectures of closed platforms are inferred from observable API surfaces and engineering blog posts where available.

A consolidated comparison appears in Table 2.2 following the competitor subsections.

2.2.1. Udemy for Business and Udemy-Like Marketplace Platforms

Description. Udemy is the world’s largest consumer marketplace for online courses, hosting more than 200,000 courses and tens of millions of learners. **Udemy for Business (UFB)** is the enterprise-facing product offering curated course libraries, analytics dashboards, and LMS integrations for corporate training teams. Similar marketplace-centric platforms include LinkedIn Learning and Pluralsight, which combine proprietary content libraries with subscription access models rather than en-

abling organisations to operate fully custom backends.

Technical architecture (public knowledge). Udemy operates a proprietary microservices architecture deployed on cloud infrastructure, with internal services for content delivery, recommendation, payment settlement, and anti-piracy DRM not exposed to customers. Public APIs for UFB focus on user provisioning (SCIM), progress reporting (xAPI subsets), and catalogue embedding—not on course authoring, payment routing, or video pipeline customisation. Video delivery relies on proprietary CDN and transcoding stacks optimised for global scale rather than customer-configurable providers such as Mux.

Pros.

- **Massive scale and reliability** — battle-tested infrastructure serving millions of concurrent learners worldwide;
- **Content marketplace** — immediate access to thousands of professional courses without content creation overhead;
- **Brand recognition** — reduces procurement friction for corporate L&D buyers familiar with the consumer product;
- **Analytics maturity** — established completion and skills reporting for enterprise HR integrations.

Cons.

- **Closed-source and non-customisable backend** — no access to enrollment logic, event pipelines, or assessment engines;
- **Revenue-sharing and licensing costs** — creators on consumer Udemy receive 37–97% revenue share depending on acquisition channel; UFB involves per-seat enterprise licensing beyond customer control;
- **No DDD/Clean Architecture reference** — zero educational value for engineering teams studying modern backend patterns;
- **Limited AI quiz generation** — AI features (if present) are platform-controlled black boxes, not extensible Whisper/LLM pipelines;
- **Vendor lock-in** — migration of course data, progress records, and payment history off-platform is intentionally difficult;
- **No self-hosting** — incompatible with university data-sovereignty requirements addressed by NexaLearn.

2.2.2. Teachable and Thinkific

Description. **Teachable** and **Thinkific** are leading SaaS platforms targeting independent course creators, coaches, and small academies. They provide drag-and-drop course builders, landing pages, email marketing integrations, affiliate programmes, and built-in checkout experiences. Both platforms prioritise time-to-first-sale over backend extensibility, positioning themselves as “business-in-a-box” solutions for non-technical creators.

Technical architecture (public knowledge). Both platforms operate multi-tenant closed backends (likely Ruby on Rails or similar monolithic/SaaS stacks internally, though not publicly confirmed). Creators interact primarily through web dashboards; limited public REST APIs support enrolment webhooks, coupon management, and third-party automation via Zapier. Video hosting is bundled—Teachable migrated to native video hosting; Thinkific offers built-in hosting with CDN delivery. Payment processing integrates Stripe and PayPal on behalf of creators, with platforms acting as merchant-of-record or payment facilitator depending on plan tier.

Pros.

- **Rapid creator onboarding** — first course publishable within hours without engineering staff;
- **Integrated marketing tooling** — sales pages, coupons, and affiliate tracking reduce need for external tools;
- **Stripe integration out-of-the-box** — creators receive payment processing without PCI compliance burden;
- **Established user bases** — large creator communities, template libraries, and support documentation;
- **Predictable SaaS pricing** — subscription tiers (\$39–\$499/month) versus custom development costs.

Cons.

- **Closed-source with restricted API access** — cannot modify enrollment state machines, outbox patterns, or assessment algorithms;
- **Platform transaction fees** — Teachable charges 5% transaction fees on Basic plans plus payment processing; Thinkific charges 0–5% depending on tier;
- **No AI-native quiz generation pipeline** — assessments are manually authored; no Whisper/LLM integration comparable to NexaLearn’s AI Pipeline context;

- **Video processing opacity** — creators cannot swap Mux for alternative encoders or implement custom DRM policies;
- **No DDD architecture** — unsuitable as an engineering reference for Computer Engineering curricula;
- **Vendor lock-in** — course content, student lists, and payment history export with incomplete API coverage;
- **Limited white-label depth** — custom domains available but backend remains platform-controlled multi-tenant infrastructure.

2.2.3. Moodle (Open-Source LMS)

Description. **Moodle** is the dominant open-source Learning Management System globally, powering institutions from primary schools to research universities. It provides course management, forums, quizzes, gradebooks, SCORM players, and thousands of community plugins. Moodle's GPL licence and self-hosting support have made it the default choice for budget-conscious institutions worldwide.

Technical architecture (public knowledge). Moodle is a PHP monolith using a plugin architecture layered atop a relational database (typically MySQL/MariaDB or PostgreSQL). Core subsystems include mod (activity modules), auth plugins, enrol plugins, and theme engines. Background tasks run via cron; caching supports Redis/Memcached. Version 4.x introduces modernised UI (Boost theme) but retains the fundamental plugin-sprawl architecture. There is no native event-driven outbox, bounded context separation, or TypeScript type safety across module boundaries.

Pros.

- **Open source and self-hostable** — full source access, no per-seat licensing, data sovereignty;
- **Massive installed base** — extensive community, plugin ecosystem, and institutional familiarity;
- **Feature breadth** — gradebooks, workshops, Wikis, SCORM, competencies, and proctoring plugins;
- **Academic credibility** — decades of pedagogical feature accumulation validated in accreditation contexts;
- **Global localisation** — interface packs in 100+ languages.

Cons.

- **Legacy monolithic architecture** — plugin interdependencies create maintenance burden; no DDD or Clean Architecture guidance;

- **PHP stack divergence** — modern edtech startups prefer TypeScript/Node.js talent pools over PHP plugin development;
- **No native Stripe/Mux/AI pipeline integration** — requires custom plugins of varying quality; payment and video remain fragmented;
- **AI features via third-party plugins only** — no unified Whisper/LLM quiz generation context; inconsistent security models across plugins;
- **Operational complexity at scale** — large deployments require careful cron tuning, caching layers, and database optimisation without cloud-native defaults comparable to Docker Compose NestJS stacks;
- **Developer experience** — lacks OpenAPI-first REST design, Prisma-type safety, and Testcontainers integration testing patterns exemplified by NexaLearn.

2.2.4. Coursera for Campus

Description. Coursera for Campus (and related Coursera for Business products) provides universities and enterprises access to Coursera’s catalogue of courses and degrees, combined with tools for instructors to author private courses, manage programmes, and track learner progress. Coursera partners with top global universities for branded content, positioning the platform as a premium institutional channel rather than a self-hosted backend framework.

Technical architecture (public knowledge). Coursera operates a proprietary cloud-native platform originally built on Scala/Play Framework microservices (per early engineering publications), with specialised services for video streaming, peer grading, item response theory assessments, and identity verification. Integrations include LTI for LMS embedding, API access for enterprise reporting, and SSO via SAML. The platform is entirely SaaS-hosted on Coursera infrastructure; institutions do not receive backend source code or database access.

Pros.

- **Premium content library** — immediate access to courses from Stanford, Google, IBM, and partner universities;
- **Institutional brand association** — credibility for universities offering co-branded programmes;
- **Scalable video and assessment infrastructure** — handles massive concurrent MOOC enrolments;
- **Guided projects and hands-on labs** — mature interactive content formats;

- **Some AI integration** — Coursera Coach and AI-assisted learning features emerging in 2024–2025 product releases.

Cons.

- **Fully closed-source** — zero backend customisation; institutions are tenants, not owners;
- **High institutional licensing costs** — per-student or per-programme fees scale expensively compared to self-hosted NexaLearn;
- **No DDD reference or API completeness** — integration limited to prescribed LTI and reporting endpoints;
- **AI features not extensible** — institutions cannot deploy custom Whisper/LLM pipelines or modify quiz generation logic;
- **Payment integration opaque** — Coursera manages monetisation for partner content; private course payment flows constrained;
- **Vendor lock-in and data portability limits** — migrating authored content and learner records off-platform is restricted;
- **Limited relevance for custom edtech startups** — business model assumes catalogue leverage, not greenfield platform construction.

2.2.5. LearnDash (WordPress LMS Plugin)

Description. LearnDash is a commercial WordPress plugin transforming a standard WordPress site into an LMS with courses, lessons, topics, quizzes, certificates, and drip-feed content scheduling. It targets creators and small businesses already invested in the WordPress ecosystem, leveraging familiar admin UI and the vast WordPress plugin marketplace for commerce (WooCommerce), membership (MemberPress), and page building (Elementor).

Technical architecture (public knowledge). LearnDash extends WordPress's PHP/MySQL architecture with custom post types for courses and lessons, WordPress hooks for lifecycle events, and Gutenberg blocks for frontend rendering. Quizzes use proprietary question engines; video lessons embed via shortcodes supporting YouTube, Vimeo, or self-hosted MP4. Payments integrate through WooCommerce or native Stripe/PayPal add-ons. The architecture inherits WordPress's request-per-page model, transients API for caching, and wp-cron for scheduled tasks—fundamentally different from NestJS microservice-ready modular monoliths.

Pros.

- **Low entry cost for WordPress users** — leverages existing hosting, themes, and SEO workflows;
- **GPL-licensed core plugin** — source readable; large WordPress developer community for extensions;
- **Integrated commerce via WooCommerce** — mature payment plugin ecosystem;
- **Certificate and drip content features** — adequate for solo creators and small academies;
- **Affordable licensing** — approximately \$199/year for single-site plugin licence versus enterprise LMS contracts.

Cons.

- **WordPress architectural constraints** — not designed for high-concurrency API-first backends; scaling requires caching plugins and database tuning rather than horizontal stateless API tiers;
- **No DDD, outbox, or event-driven reliability patterns** — enrollment/payment race conditions handled ad hoc via WordPress post meta;
- **Video processing limited** — no Mux-grade transcoding pipeline; relies on external hosts or raw MP4 uploads without adaptive bitrate streaming defaults;
- **AI quiz generation absent natively** — requires separate WordPress AI plugins with inconsistent quality and security;
- **Security surface area** — WordPress's plugin ecosystem is a frequent attack target; institutional IT departments often resist WordPress for regulated training data;
- **Not a developer reference architecture** — unsuitable for Computer Engineering graduation projects demonstrating NestJS, Prisma, and transactional outbox patterns;
- **Vendor dependency on LearnDash updates** — major WordPress or LearnDash version incompatibilities can break production sites.

2.2.6. Competitive Comparison Summary

Table 2.2 consolidates the market survey into a single reference matrix. NexaLearn is included as the proposed alternative to highlight differentiated capabilities. Rating conventions: **Yes** = first-class native support; **Partial** = available via plugins, limited APIs, or opaque platform features; **No** = not meaningfully supported; **N/A** = not applicable to platform business model.

Table 2.2: Competitive Platform Comparison (Table 2.1)

Platform	Open Source	AI Features	Video Processing	Payment Integration	DDD Architecture	Price (indicative)
Udemy / UFB	No	Partial	Proprietary CDN	Proprietary	No	Per-seat enterprise; creator revenue share
Teachable	No	No	Built-in hosting	Stripe (built-in)	No	\$39–\$499/mo + transaction fees
Thinkific	No	No	Built-in hosting	Stripe/PayPal	No	\$0–\$499/mo + transaction fees
Moodle	Yes	Partial (plugins)	Plugin-dependent	Plugin-dependent	No	Free (GPL) + hosting/support costs
Coursera Campus	No	Partial (platform AI)	Proprietary	Proprietary	No	Institutional licence (per student)
LearnDash	Partial (GPL plugin)	No (plugins only)	Basic embed/host	WooCommerce/Stripe	No	\$199/yr plugin + WordPress hosting
NexaLearn (proposed)	Yes	Yes (Whisper/LLM)	Mux integration	Stripe (native)	Yes	OSS core + \$5,000/yr enterprise support




Figure 2.2: Competitive Positioning: Architectural Openness versus Backend Customisation Depth

Figure 2.3 contextualises NexaLearn’s target market within the broader e-learning landscape. Corporate e-learning and higher education dominate in market size, consistent with NexaLearn’s primary and secondary customer segments (Section 2.1). The STEM and coding segment—at \$11B—represents a growing overlap with the developer-adoption funnel described in the business case (Section 2.3).




Figure 2.3: Global E-Learning Market by Segment (2025). Data sources: HolonIQ, Statista Global EdTech Market Report, Grand View Research.

Figure 2.2 visualises the strategic whitespace NexaLearn occupies: higher architectural openness and DDD discipline than closed SaaS platforms, and more modern API-first engineering than legacy Moodle/WordPress stacks. No surveyed competitor simultaneously offers open source, native Stripe/Mux/AI pipeline integration, and an explicit DDD modular monolith reference—validating the market gap identified in Chapter 1.

In conclusion, the market survey demonstrates that while each competing platform excels within its niche—Udemy’s catalogue scale, Teachable’s creator UX, Moodle’s institutional adoption, Coursera’s academic partnerships, LearnDash’s WordPress accessibility—none delivers NexaLearn’s combination of *open-source extensibility*, *production reliability patterns*, and *AI-augmented backend completeness*. The following section translates this positioning into a quantified business case and financial analysis.

2.3. Business Case and Financial Analysis

This section presents a feasibility-level business case for commercialising NexaLearn under an **open-core model**: the backend source code and documentation remain freely available to maximise developer adoption, while **enterprise support contracts** (\$5,000 per organisation per year) monetise SLAs, priority incident response, migration assistance, and security advisory services. The financial analysis covers adoption assumptions, capital expenditure (Capex), operating expenditure (Opex), revenue projections, a twenty-four-month cash-flow forecast, and break-even analysis. All figures are expressed in US dollars (USD) and rounded to the nearest hundred for readability.

2.3.1. Business Case and Adoption Assumptions

The business case rests on two complementary adoption curves:

Developer community adoption (open-source funnel). NexaLearn targets **500 developer adoptions** in Year 1, defined as unique organisations or teams cloning the repository, deploying via Docker Compose, and registering for community communications (GitHub stars, documentation analytics, or newsletter sign-ups). Developer adoption grows **40% year-over-year** thereafter (700 in Year 2, 980 in Year 3), following typical open-source SaaS funnel dynamics where visibility compounds through conference presentations, university capstone references, and technical blog content. Developer adopters do not generate direct revenue but reduce enterprise customer acquisition cost by demonstrating production readiness and supplying community-contributed patches.

Enterprise support revenue (monetisation). Enterprise clients are organisations requiring contractual SLAs, dedicated onboarding, and security questionnaire

support—primarily secondary (universities) and tertiary (corporate L&D) segments from Section 2.1. Year 1 targets **50 enterprise support contracts** at **\$5,000 per year** each, yielding **\$250,000** annual gross revenue under the assumption that clients prepay annually upon contract signing. Year 2 targets a 40% increase in enterprise clients (70 cumulative active contracts, 20 net new after renewals), consistent with developer funnel conversion rates of approximately 2–4% from qualified DevOps/evaluation contacts to paid support.

Pricing rationale. The \$5,000/year enterprise tier aligns with support pricing for specialised infrastructure products serving small-to-medium engineering teams (e.g., hosted observability, database support retainers). It is deliberately an order of magnitude below enterprise LMS licensing (often \$50,000–\$500,000 annually) while reflecting real costs of engineer-time for incident response and migration consulting.

2.3.2. Capital Expenditure (Capex)

One-time capital expenditure occurs primarily in Month 1 during commercial entity setup and production infrastructure provisioning, as detailed in Table 2.3:

Table 2.3: Capital Expenditure (Capex) Breakdown — Month 1

Item	Cost (USD)	Description
Production server setup	\$10,000	AWS/GCP baseline: RDS PostgreSQL, ECS/EC2 API nodes, ALB, S3 buckets, initial CDN configuration
Development tools	\$2,000	CI/CD pipelines (GitHub Actions minutes), monitoring (Datadog/Grafana Cloud starter), domain and SSL certificates
Third-party service setup	\$500	Stripe Connect/platform configuration, Mux account provisioning, webhook endpoint certification, sandbox testing
Total Capex	\$12,500	One-time expenditure in Month 1

Capex excludes graduation-project development labour (sunk academic cost). Commercialisation scenarios assume a dedicated three-person engineering team funded through Opex salaries beginning Month 3 in the cash-flow model below.

2.3.3. Operating Expenditure (Opex)

Recurring monthly operating expenditure comprises cloud infrastructure, video encoding pass-through costs, payment processing fees, and—under the commer-

cialised scenario—developer salaries, as shown in Table 2.4:

Table 2.4: Monthly Operating Expenditure (Opex) Components

Item	USD/month	Notes
Cloud hosting	\$300	API containers, PostgreSQL (managed), Redis, Nginx; scales modestly with traffic in Year 1
Mux video encoding	\$200	Pass-through encoding costs for demo and pilot customer content; variable with usage
Stripe processing fees	Variable	Approximately 2.9% + \$0.30 per transaction if NexaLearn processes subscription billing; modelled at 1% of gross revenue
Developer salaries (commercialised)	\$15,000	One senior + one mid-level engineer + part-time DevOps from Month 3 onward
Fixed subtotal	\$15,500	Excluding variable Stripe fees

Months 1–2 reflect **pre-commercialisation** infrastructure burn at \$500/month (domain, staging environment, third-party sandbox accounts) before full salary commitment begins in Month 3 coinciding with the first enterprise sales outreach.

2.3.4. Revenue Model and Year 1 Projection

Revenue derives exclusively from **enterprise support contracts** in this model; open-source adopters pay \$0. Each contract is priced at **\$5,000 per year**, invoiced annually in advance upon signature. Year 1 revenue target:

$$R_{Y1} = 50 \text{ clients} \times \$5,000 = \$250,000 \text{ gross annual revenue} \quad (2.1)$$

Sales ramp assumes gradual enterprise procurement cycles: two contracts in Month 3, accelerating to six–seven new contracts per month by Months 7–8 as case studies from early adopters mature, reaching 50 cumulative active contracts by Month 12. No consumer SaaS subscription tier is modelled in Year 1 to avoid channel conflict with open-source adoption; a managed-hosting tier may be introduced in Year 3 as optional expansion revenue.

2.3.5. Twenty-Four-Month Cash Flow Forecast

Table 2.5 presents the monthly cash-flow forecast for Months 1–24. **Revenue** includes annual prepayments from new enterprise contracts signed that month (\$5,000 per new client). **Expenses** include Opex and, in Month 1, Capex. **Net Cash Flow** equals Revenue minus Expenses; **Cumulative Cash Flow** tracks running profitability. Variable Stripe fees are approximated at 1% of revenue and embedded in Expenses.

Table 2.5: Twenty-Four-Month Cash Flow Forecast (USD)

Month	New Ent.	Cum. Ent.	Dev. Adopters*	Revenue	Expenses	Cum. Cash Flow
1	0	0	25	0	13,000	–13,000
2	0	0	55	0	500	–13,500
3	2	2	90	10,000	15,650	–19,150
4	3	5	130	15,000	15,650	–19,800
5	3	8	175	15,000	15,650	–20,450
6	4	12	230	20,000	15,700	–16,150
7	5	17	295	25,000	15,750	–6,900
8	6	23	370	30,000	15,800	+7,300
9	7	30	400	35,000	15,850	+26,450
10	7	37	430	35,000	15,850	+45,600
11	6	43	465	30,000	15,800	+59,800
12	7	50	500	35,000	15,850	+79,950
13	2	52	540	36,000	15,860	+100,090
14	2	54	580	36,000	15,860	+120,230
15	3	57	615	39,000	15,890	+143,340
16	3	60	650	39,000	15,890	+166,450
17	2	62	680	37,000	15,870	+187,580
18	3	65	700	39,000	15,890	+210,690
19	2	67	730	37,000	15,870	+231,820
20	2	69	760	37,000	15,870	+252,950
21	1	70	790	35,500	15,855	+272,595
22	1	71	820	35,500	15,855	+292,240
23	1	72	850	35,500	15,855	+311,885
24	1	73	880	35,500	15,855	+331,530

*Developer adopters (cumulative community metric, not revenue-bearing). Assumes Month 1–2 pre-commercialisation Opex \$500/mo; Capex \$12,500 in Month 1; full Opex \$15,500/mo from Month 3. Enterprise revenue = new contracts × \$5,000 annual prepayment. Break-even cumulative cash flow achieved in Month 8.

Year 1 revenue total (Months 3–12): \$250,000 from 50 new enterprise contracts, consistent with Equation 2.1. Months 13–24 include **renewal revenue** from Year 1 clients (\$5,000 per renewal) blended with new-client signings, yielding approximately \$440,000 gross revenue in Year 2 and reaching **73 cumulative enterprise clients** by Month 24. Developer adopter counts track toward 700 by Month 18 and 880 by Month 24 (40% Year 1 → Year 2 growth trajectory).

2.3.6. Break-Even Analysis

Break-even point occurs when cumulative cash flow crosses from negative to positive territory. Per Table 2.5, cumulative cash flow is $-\$6,900$ at the end of Month 7 and **+\$7,300 at the end of Month 8**. Linear interpolation places the break-even crossing at approximately **Month 8, Week 2**—eight months after commercial launch (Month 3) and ten months after initial Capex deployment. Break-even is achieved ahead of Year 1 revenue target completion (Month 12), indicating that the open-core enterprise support model reaches self-sustaining cash flow before the first full sales cycle concludes.

Sensitivity considerations:

- **Slower enterprise sales** (50% of projected signings) delays break-even to approximately Month 14–16 but remains viable given low fixed infrastructure costs relative to salary burden;
- **Higher cloud/Mux usage** (\$500/month additional) shifts break-even by less than one month due to revenue leverage from enterprise prepayments;
- **Deferred salary hiring** (Month 6 instead of Month 3) advances break-even to Month 6 but risks insufficient support capacity for early enterprise clients.




Figure 2.4: Break-Even Analysis: Cumulative Cash Flow over Twenty-Four Months (Figure 2.1)

Figure 2.4 illustrates the cumulative cash-flow trajectory: initial Capex and pre-revenue months produce a trough of approximately \$20,450 (Month 5), followed by recovery as enterprise prepayments accelerate in Months 6–8. The shaded break-even line at Month 8 demarcates transition from investment phase to positive cumulative cash generation. By Month 24, cumulative cash flow exceeds \$331,000, demonstrating that NexaLearn’s commercialisation path—while modest compared to venture-scale SaaS unicorns—is financially plausible for a bootstrapped open-source infrastructure product serving the edtech developer segment.

Figure 2.5 summarises the three-year financial trajectory. Year 1 revenue of \$250,000 and expenses of approximately \$170,000 yield a net profit of \$80,000, consistent with the cumulative cash flow at Month 12 (\$79,950). Year 2 projects \$440,000 in revenue against \$210,000 in expenses (net \$230,000), and Year 3 projects \$616,000 revenue against \$260,000 in expenses (net \$356,000), assuming 40% enterprise client growth and controlled cost scaling. The widening profit margin reflects the high-margin nature of enterprise support revenue once the initial team is established.




Figure 2.5: Revenue Projection: Years 1–3 (Enterprise Support Model). Revenue assumes \$5,000/contract with 50, 70, and 73 cumulative clients respectively; expenses include Capex amortisation, Opex, and staff costs as detailed in Tables 2.3 and 2.4.

2.3.7. Financial Conclusions

The business case supports three conclusions aligned with Chapter 1 motivation:

1. **Market demand exists** for an open, DDD-architected e-learning backend, evidenced by 500 projected developer adoptions and 50 enterprise support clients in Year 1;
2. **Commercialisation via open-core enterprise support** reaches break-even within eight months under conservative sales ramp assumptions, with Year 1 gross revenue of \$250,000 against total Year 1 expenses of approximately \$170,000 (Capex + Opex);
3. **Competitive pricing at \$5,000/year** undercuts enterprise LMS and SaaS alternatives while funding ongoing engineering maintenance, satisfying feasibility study requirements in Appendix E.

This financial analysis complements the competitive survey in Section 2.2 and the customer segmentation in Section 2.1, completing the market visibility study and establishing the commercial context for the technical literature and system design presented in Chapters 3 and 4.

Chapter 3:

Literature Survey

This chapter surveys the foundational technologies, architectural patterns, and academic literature underpinning the NexaLearn platform. The survey is organised into six thematic sections that correspond to the principal technical pillars of the system. Section 3.1 reviews Domain-Driven Design (DDD) and Clean Architecture, establishing the decomposition strategy used to partition NexaLearn into fifteen bounded contexts with explicit dependency rules. Section 3.2 examines event-driven architecture and the transactional outbox pattern, which together guarantee reliable asynchronous integration across contexts and with external providers. Section 3.3 surveys recent advances in AI-assisted e-learning—specifically OpenAI Whisper for speech-to-text transcription and large language model (LLM) pipelines for automated quiz generation—detailing how these technologies reduce instructor workload while maintaining content quality through human-in-the-loop review. Section 3.4 covers HTTP Live Streaming (HLS), the Mux video platform, and signed URL access control, which collectively enable adaptive-bitrate playback and fine-grained entitlement enforcement. Section 3.5 presents a structured comparison of prior academic and industrial e-learning backend approaches, highlighting gaps in domain modelling, event reliability, and AI integration that NexaLearn addresses. Finally, Section 3.6 justifies the specific technology and pattern choices—NestJS, Prisma, the outbox pattern, and modular monolith deployment—synthesising the surveyed literature into a coherent architectural rationale.

3.1. Background on Domain-Driven Design and Clean Architecture

3.1.1. Domain-Driven Design Fundamentals

Domain-Driven Design (DDD), as originally formulated by Evans [1], provides a set of strategic and tactical modelling patterns for managing complexity in software systems whose value derives from intricate domain logic. The central premise is that the *domain model*—a software embodiment of the business concepts, rules, and invariants—must be the primary artefact driving design decisions, rather than implementation technology or database schemas. Evans introduces the concept of a **ubiquitous language**: a shared vocabulary used by domain experts, developers, and all project stakeholders, embedded consistently in code, diagrams, and documentation. This language prevents the translation errors that arise when business concepts are renamed or restructured between specification documents and implementation.

At the strategic level, DDD prescribes decomposing large problem spaces into **bounded**

contexts—explicit boundaries within which a particular ubiquitous language applies. Each bounded context owns its domain model, database schema, and internal logic; communication with other contexts occurs through well-defined integration contracts. A **context map** documents the relationships between bounded contexts, specifying whether the relationship is upstream/downstream (customer-supplier), conformist (the downstream adopts the upstream model without change), or shared-kernel (two contexts share a subset of the domain model). This decomposition prevents the architectural erosion described in Section 1.2.1 by enforcing that code in one context cannot directly access repositories or entities owned by another.

At the tactical level, DDD defines a vocabulary of building blocks:

Entity. An object defined by its identity rather than its attributes. Two entities with identical field values are distinguishable by their identifier. In NexaLearn, `CourseEntity` and `QuizEntity` are entities whose identity persists across state changes.

Value Object. An immutable object defined solely by the values of its attributes. `QuestionOption`, `Address`, or monetary amounts are typical value objects. Equality is structural rather than identity-based.

Aggregate. A cluster of entities and value objects treated as a single consistency boundary. Aggregates have a root entity (the **aggregate root**) that enforces invariants and mediates all external access to the aggregate’s internals. In NexaLearn, `QuizEntity` is the aggregate root for a graph that includes `QuestionEntity` children and `QuestionOption` value objects; all operations that add, remove, or reorder questions must go through the quiz aggregate root, ensuring invariants such as “a quiz must have at least one question” and “question ordering within a quiz is contiguous.”

Domain Event. A record of something that happened in the domain that domain experts care about. `EnrollmentActivated`, `LessonCompleted`, and `PaymentSettled` are domain events in NexaLearn. Events are immutable facts published to interested contexts or external subscribers.

Repository. A mechanism for encapsulating storage, retrieval, and search behaviour, presenting a collection-like interface to the domain layer. Repositories per aggregate root abstract the underlying database technology—PostgreSQL, Redis, or in-memory—from domain logic.

Domain Service. A stateless operation that does not naturally belong to an entity or value object. Domain services orchestrate multiple aggregates or coordinate with external systems while keeping business rules in the domain layer.

NexaLearn’s codebase applies these tactical patterns systematically. The `QuizEntity` aggregate root, for example, owns a list of `questions`, enforces the invariant that minimum passing score is between 0 and 100, and exposes behaviour methods such as `addQuestion()` rather than allowing direct mutation of the `questions` array.

3.1.2. Clean Architecture Layers

Martin's Clean Architecture [2] refines the layering concerns earlier expressed in the Hexagonal Architecture and Onion Architecture patterns. The defining principle is the **Dependency Rule**: source code dependencies must point *inward*, from outer layers toward inner layers. Nothing in an inner layer can know about anything in an outer layer.

Applied to NexaLearn, the four layers are:

1. **Domain Layer (innermost).** Contains entities, value objects, aggregate roots, domain events, and repository *interfaces* (ports). This layer has zero dependencies on frameworks, databases, or external libraries. The ubiquitous language lives here: `Course`, `Enrollment`, `Payment`, `VideoAsset` are expressed as pure TypeScript classes with behaviour and invariants.
2. **Application Layer.** Contains use-case interactors that orchestrate domain objects to fulfil a single user story. Each use case (e.g., `SubmitQuizAttemptUseCase`, `ActivateEnrollmentUseCase`) depends on repository interfaces from the domain layer but has no knowledge of how those repositories are implemented. The application layer may also define application-specific DTOs and events.
3. **Infrastructure Layer.** Implements the repository interfaces declared in the domain layer. Prisma-based repository implementations, Mux API clients, Stripe SDK wrappers, and outbox event publishers all reside here. This is also where the transactional outbox publisher, job queues, and scheduled task processors live.
4. **Presentation Layer (outermost).** Controllers, HTTP request/response DTOs, validation pipes, guards, and interceptors. NestJS controllers in NexaLearn are thin: they parse HTTP input, delegate to an application use case, and format the response. No business logic resides in controllers.

The Dependency Rule is enforced through TypeScript's module system and NestJS module boundaries. A custom ESLint rule in the NexaLearn codebase disallows infrastructure-layer modules from importing any NestJS decorator or Prisma client code into domain or application layer files.

3.1.3. Application to NexaLearn's Bounded Contexts

NexaLearn decomposes its domain into fifteen bounded contexts, each mapped to a NestJS module (or set of related modules). The decomposition follows Dossena et al.'s [3] recommendations for applying DDD to e-learning platforms, where enrolment, assessment, content management, and analytics each form natural bounded contexts because they involve distinct stakeholders, disparate lifecycle rules, and different persistence trade-offs. The **Courses** context, for example, owns course catalogue data, curriculum structure, and lesson content. It exposes read-only projections for the **Progress** context via materialised views and reacts to `EnrollmentActivated` events published by the **Enrollment** context. The **Assessment** context owns quizzes, questions, attempts, and scoring logic; its aggregate root

(QuizEntity) enforces invariants such as maximum attempts per learner and question ordering. This bounded-context-per-module approach, combined with Clean Architecture layering within each module, yields a codebase where a developer can understand the Quiz domain entirely within the Assessment module without traversing cross-cutting infrastructure concerns.

The modular monolith deployment strategy (detailed in Section 3.6) packages all fifteen contexts into a single NestJS application process, but the DDD boundaries and Clean Architecture layers ensure that the system remains maintainable as it grows. Each context can be extracted into an independent microservice with minimal refactoring if scaling or team autonomy requirements change, following the split-pattern described by Yang et al. [4].

Ultimately, DDD and Clean Architecture provide NexaLearn with three benefits directly motivated by the problem statement of Chapter 1. First, **testability**: domain entities with no framework dependencies can be unit-tested without NestJS test containers or database connections. Second, **maintainability**: a developer can reason about the Enrollment aggregate without understanding Stripe webhook processing details. Third, **bounded context isolation**: the Courses context can evolve its schema and domain logic without cascading changes to Discussion or Streak, provided the published event contracts remain stable.

3.2. Background on Event-Driven Architecture and the Outbox Pattern

3.2.1. Event-Driven Architecture in Distributed Systems

Event-Driven Architecture (EDA) is a software paradigm in which components communicate by producing and consuming *events*—immutable records of something that has occurred, rather than by issuing direct command invocations [5]. In e-learning systems, events arise naturally from learner interactions (lesson viewed, quiz submitted, discussion post created), system state changes (enrolment activated, certificate issued), and external provider callbacks (payment succeeded, video ready). EDA decouples event producers from consumers: a single enrolment activation event may be consumed by the Progress context (to initialise progress projections), the Media context (to issue playback entitlements), the Notifications context (to send a welcome email), and the Search context (to update enrolment counts), all without the Enrollment context having any awareness of these downstream dependents.

The principal challenge in EDA is reliable event delivery. When a domain operation both modifies the database and publishes an event, these two actions must be atomic: either both succeed or neither takes effect. Performing the database write first, then publishing the event, creates a window in which the database write is visible but the event is not published—if the application crashes between the two operations, or if the message broker is unreachable, the system enters an inconsistent state. Conversely, publishing the event first and then writing the database risks acting upon an event whose associated state does not yet exist. This is known as the **dual-write problem**.

3.2.2. The Transactional Outbox Pattern

The transactional outbox pattern, popularised by Richardson [6] in the microservices community and independently described as the “application events” pattern in messaging literature, solves the dual-write problem by persisting events in the same database transaction as the domain mutation. Rather than publishing an event immediately to a message broker, the application inserts a row into an `outbox_events` table within the same database transaction that updates the domain aggregates. A separate background process—the **outbox publisher**—polls the outbox table for unprocessed events, publishes them to the message broker (or a message queue), and marks them as processed.

This design guarantees **at-least-once delivery**: an event may be published more than once if the publisher crashes after publishing but before marking the event as processed, but it will never be lost. Consumer applications use idempotency keys to achieve exactly-once semantics despite at-least-once transport.

3.2.3. NexaLearn’s Outbox Implementation

NexaLearn implements a variant of the transactional outbox pattern using PostgreSQL as the event store and a dedicated NestJS `OutboxProcessorService` as the publisher. The key components are:

- **OutboxEvent Entity.** A database table with columns: `id` (UUID primary key), `aggregate_type`, `aggregate_id`, `event_type`, `payload` (JSONB), `created_at`, `processed_at` (nullable), and `retry_count`.
- **OutboxPublisher Service.** A domain service that receives a domain event and inserts a corresponding outbox row. This service is invoked by domain services or use-case interactors after successfully modifying domain aggregates.
- **OutboxProcessorService.** A polling background worker (implemented as a NestJS `@Cron` scheduled task or a dedicated `@Processor` job queue consumer) that periodically queries `SELECT * FROM outbox_events WHERE processed_at IS NULL ORDER BY created_at ASC LIMIT n FOR UPDATE SKIP LOCKED`. The `FOR UPDATE SKIP LOCKED` clause ensures that multiple concurrent processor instances can work on different batches without contention—a critical property for horizontal scaling.

3.2.4. Idempotency and Exactly-Once Processing

Idempotency ensures that processing the same event multiple times produces the same outcome. NexaLearn’s **ProjectionEventLedger** table records which events have been applied to each read model. The ledger schema includes: `id` (UUID), `event_id` (foreign key to `outbox_events`), `projection_type` (e.g., `course-progress-view`), `handler_id` (the specific handler instance that processed the event), and `applied_at`. Before applying an event to a projection, the handler attempts `INSERT INTO projection_event_ledger (event_id, projection_type, handler_id) VALUES ($1, $2, $3) ON CONFLICT (event_id, projection_type) DO NOTHING`. If the insert succeeds, the event has not been

processed and the handler proceeds; if it returns zero rows (due to the conflict), the event was already processed and the handler skips it.

This approach is inspired by the idempotency-key pattern used by Stripe [7] for payment processing, adapted for internal event consumption. The combination of transactional outbox and idempotent projection handlers eliminates the dual-write problem without requiring distributed transactions, a two-phase commit coordinator, or an external message broker—a deliberate simplification discussed in Section 3.6.

3.2.5. CQRS-Lite Read Models

While NexaLearn does not implement full Command Query Responsibility Segregation (CQRS) [8] with separate write and read databases, it adopts a “CQRS-lite” pattern: several bounded contexts maintain specialised read models that are updated asynchronously by outbox event handlers. These read models are optimised for specific query patterns that would require expensive joins or recomputation across the write-optimised domain schema. Examples include:

- **CourseProgressView.** Denormalised per-learner progress summary showing total lessons, completed lessons, quiz scores, and streak data, updated reactively when `LessonCompleted`, `QuizAttemptSubmitted`, or `StreakUpdated` events are processed.
- **ResumeView.** A materialised row per learner recording the most recently accessed lesson, percentage through the video, and the last active timestamp, enabling efficient “resume where you left off” endpoints without scanning progress logs.
- **LessonAnalytics.** Aggregated view per lesson: total views, average completion rate, average quiz score, and discussion activity counts, computed from event streams rather than expensive OLAP queries on operational tables.
- **CourseSearchIndex.** A denormalised table containing course title, description, instructor name, category, average rating, and enrolment count, used by the search endpoint instead of joining across five write-optimised tables on every request.

These projections are stored in the same PostgreSQL instance but in a separate `read_models` schema, reinforcing the conceptual separation between write-optimised domain tables and read-optimised projection tables. The outbox processor guarantees that projections eventually converge to consistency with the write model, typically within milliseconds of the source event being committed. Mehta et al. [9] have demonstrated that this eventual-consistency approach scales to educational platforms processing millions of learner interaction events per day, with projection latency dominated by the outbox polling interval rather than database throughput.

3.3. Background on AI in E-Learning

3.3.1. OpenAI Whisper for Speech Recognition

OpenAI Whisper [10] is a general-purpose speech recognition model trained on 680,000 hours of multilingual and multitask supervised data collected from the web. Unlike prior automatic speech recognition (ASR) systems that relied on meticulously curated, single-domain datasets, Whisper leverages weak supervision at an unprecedented scale, resulting in robust performance across diverse acoustic environments, speaker accents, and recording qualities. The model architecture is a sequence-to-sequence Transformer encoder-decoder: audio is resampled to 16 kHz, converted to an 80-channel log-mel spectrogram, and encoded by the Transformer encoder; the decoder then autoregressively produces the transcript text, optionally including timestamps, language identification, and segment boundaries.

Whisper’s relevance to e-learning platforms is established by Panthel et al. [11], who evaluated Whisper transcriptions against human-produced subtitles across 120 university lecture recordings. They reported word error rates between 5.2% and 8.1% on English lectures with clear audio, rising to 12–15% on recordings with background noise, accented speech, or specialised technical terminology. Critically, the study concluded that Whisper’s output was “sufficient for automated downstream tasks such as keyword extraction, content indexing, and AI quiz generation [...] with manual correction recommended only for public-facing subtitles where near-perfect accuracy is required.”

NexaLearn integrates Whisper through a dedicated **AI Pipeline** bounded context. When a video asset completes transcoding (Section 3.4), the pipeline:

1. Downloads the MP4 rendition to a transient processing environment.
2. Extracts the audio track and segments it into 30-second chunks aligned with the Whisper context window.
3. Sends each chunk to the Whisper API or a self-hosted Whisper instance (configurable per deployment), requesting transcripts with word-level timestamps and segment boundaries.
4. Assembles the per-chunk transcripts into a monolithic transcript document, preserving the temporal mapping between transcript segments and video offset.
5. Stores the transcript as a JSON document associated with the `VideoAsset` entity, enabling downstream consumers to retrieve “the transcript segment starting at second 145” for fine-grained synchronisation.

3.3.2. LLM-Based Quiz Generation

The second AI capability of NexaLearn is automated quiz generation from lecture transcripts using large language models. Kumar et al. [12] demonstrated an end-to-end pipeline that inputs a lecture transcript and outputs a structured set of multiple-choice questions, true-or-false statements, and short-answer prompts, each with a correct answer, plausible distractors, and a justification. Their evaluation across 50 university lecture transcripts showed that

GPT-4-generated questions were rated by domain experts as acceptable or better in 73% of cases, with the primary failure mode being questions that tested recall of peripheral transcript details rather than conceptual understanding.

NexaLearn's AI pipeline generalises this approach. Upon receiving a transcript (either generated by Whisper or uploaded manually by the instructor), the pipeline constructs a prompt instructing the LLM to produce structured JSON adhering to a predefined schema:

```

1 {
2   "questions": [
3     {
4       "stem": "What is the primary invariant enforced by ...?",
5       "type": "multiple-choice",
6       "options": [
7         { "id": "A", "text": "...", "isCorrect": false },
8         { "id": "B", "text": "...", "isCorrect": true },
9         { "id": "C", "text": "...", "isCorrect": false },
10        { "id": "D", "text": "...", "isCorrect": false }
11      ],
12      "justification": "..."
13    }
14  ]
15 }
```

Listing 3.1: AI-generated quiz JSON schema

The generated quiz payload is not published directly to learners. Instead, NexaLearn applies a **human-in-the-loop review** workflow: the quiz is stored with a `READY_FOR_REVIEW` status, and the instructor is notified (via the Notifications context) to review, edit, approve, or reject each proposed question. Only upon instructor approval does the quiz status transition to `PUBLISHED`, making it available to enrolled learners. This review window ensures that the instructor retains pedagogical control while benefiting from automated question drafting—reducing the time required to create a quiz from a 60-minute lecture from approximately two hours of manual authoring to fifteen minutes of review and refinement.

3.3.3. HMAC-Signed Callback Authentication

The AI pipeline operates asynchronously: NexaLearn submits a transcription or quiz-generation request, the external AI service processes it (potentially taking seconds to minutes), and the service calls back a NexaLearn endpoint with the results. To authenticate these callbacks, NexaLearn implements an HMAC-signature scheme inspired by Stripe webhook signatures [7]. Each callback request includes an `X-Pipeline-Signature` header with format `t=<unix-timestamp>,v1=<hex-digest>`. The recipient—the AI Pipeline controller—recomputes the HMAC-SHA256 digest of the request body using the shared secret, verifies that the timestamp is within a configurable tolerance window (default 300

seconds), and rejects requests whose signature does not match or whose timestamp is stale. This prevents replay attacks and ensures that only callbacks from the genuine AI pipeline are accepted.

The HMAC callback pattern, combined with the outbox event processing described in Section 3.2, guarantees that AI-generated content is durably stored even if the application crashes during callback handling. The callback handler inserts the generated transcript or quiz into the database within a transaction that also creates an outbox event; the outbox processor subsequently notifies the instructor and updates related projections.

3.4. Background on Video Streaming Technology

3.4.1. HTTP Live Streaming (HLS)

HTTP Live Streaming (HLS), standardised as RFC 8216 [13], is the dominant adaptive-bitrate streaming protocol used by major video platforms including YouTube, Twitch, and Netflix. HLS works by segmenting a source video into short (typically 2–10 second) chunks encoded at multiple bitrates and resolutions. A **manifest file** (the `.m3u8` playlist) enumerates the available renditions and their segment URLs. The client—a web browser or mobile app—dynamically selects the highest bitrate segment that its current network conditions can sustain, switching to a lower bitrate when bandwidth drops and back to a higher bitrate when bandwidth recovers. This adaptive behaviour eliminates buffering interruptions while maximising playback quality.

The HLS specification defines a hierarchical playlist structure. A *master playlist* lists all variant streams (e.g., 360p, 720p, 1080p) and auxiliary media such as subtitles and alternate audio tracks. Each variant stream has its own *media playlist* containing the sequence of segment URLs and their durations. The protocol supports encryption via AES-128 or SAMPLE-AES, allowing content access control through key delivery authorisation.

3.4.2. Mux Video Platform

Mux [14] is a video API platform that abstracts the complexity of video ingestion, transcoding, storage, and delivery. NexaLearn integrates with Mux through the **Media** bounded context, which manages the full lifecycle of video assets. Mux provides:

- **Direct Uploads.** The client uploads a video file directly to Mux’s storage via a *direct upload endpoint*, obtaining a `upload_id` and a playback ID. The upload is initiated by NexaLearn’s backend, which requests a signed upload URL from Mux and returns it to the client; the client then PUTs the file directly to Mux, avoiding double-hop through the NexaLearn server.
- **Transcoding Pipeline.** Upon receiving a complete upload, Mux transcodes the source video into multiple HLS renditions (typically 360p, 480p, 720p, and 1080p) with configurable video codec (H.264), audio codec (AAC), and segment duration.
- **Webhook Notifications.** Mux sends webhooks to NexaLearn for lifecycle events:

`video.upload.asset_created`, `video.asset.ready`, `video.asset.errorred`, `video.asset.deleted`. These webhooks drive the video asset state machine.

- **Playback IDs and Signed URLs.** Each transcoded asset receives a playback ID. For gated content, Mux supports *signed playback URLs* using JSON Web Tokens (JWT) [15]. The NexaLearn backend generates a JWT containing the playback ID, an expiration time, and optional access restrictions (e.g., referrer whitelist, IP allowlist), signed with a Mux-specific signing key. The client embeds this signed JWT in the playback URL; Mux’s delivery infrastructure verifies the signature and expiration before serving segments.

3.4.3. Video Asset Lifecycle in NexaLearn

The Media bounded context implements a state machine for each `VideoAsset` entity. The states are:

AWAITING_UPLOAD. Initial state. A `Video` or `Lesson` entity has been created with an association to a new video, but the instructor has not yet uploaded the file. No Mux resource exists.

UPLOADED. The instructor has initiated an upload and Mux has confirmed receipt of the source file. The state transitions from `AWAITING_UPLOAD` to `UPLOADED`.

TRANSCODING. Mux sends `video.upload.asset_created`, indicating that the source file is being transcoded into HLS renditions. The NexaLearn backend creates a `GenerationJob` entity for the AI pipeline (Section 3.3) in `QUEUED` state, awaiting the completed transcript.

READY. Mux sends `video.asset.ready`, indicating that all HLS renditions are available. NexaLearn updates the `VideoAsset` with the playback ID and signed URL expiration policy. The AI pipeline receives a `VideoReady` event via the outbox, triggering the transcript-generation workflow.

FAILED. Transcoding failed. NexaLearn logs the Mux error code and message, notifies the instructor, and makes the upload available for retry.

3.4.4. MP4 Renditions for the AI Pipeline

While HLS is the delivery format for learner playback, the Whisper transcription model (Section 3.3) requires a contiguous audio or video file, not segmented HLS chunks. Mux provides a static MP4 rendition for each asset, accessible through a separate signed URL with an extended expiration window—sufficient for the AI pipeline to download the file, extract the audio track using FFmpeg, and feed it to Whisper. This dual-delivery strategy—HLS for learners, MP4 for automated processing—is transparent to the instructor, who uploads a single source file.

The interplay between video lifecycle and the AI pipeline exemplifies NexaLearn’s event-driven design: a single external event (`video.asset.ready`) propagates through the outbox

to trigger transcript generation, quiz creation, instructor notification, and search index updates, all without explicit orchestration code in the Media context. An asset may also transition directly to FAILED from UPLOADED or TRANSCODING if an error occurs at any stage.

3.5. Comparative Study of Previous Work

Several academic and industrial projects have explored architectural patterns for e-learning backends, event-driven reliability, and AI-assisted content generation. This section positions NexaLearn within the landscape of prior work by analysing four representative studies and identifying the gaps that NexaLearn’s combined approach addresses.

3.5.1. Dossena et al. (2022) — DDD in E-Learning Platforms

Dossena, Fugini, and Peroni [3] present a case study of applying DDD strategic and tactical patterns to the design of a university LMS. Their work decomposes the e-learning domain into six bounded contexts—Course Management, Enrollment, Assessment, Gradebook, Communication, and Reporting—and documents the context map with upstream/downstream relationships. The authors report that the bounded context decomposition reduced inter-module coupling by 40% compared to a previous monolithic version, and that the ubiquitous language improved communication between developers and faculty stakeholders during requirements gathering.

Limitations. The project was constrained to a single-semester pilot involving 15 courses and 300 students, limiting generalisability to production-scale deployments. The architecture relies on a relational database with no event-driven patterns, meaning that cross-context synchronisation is performed through direct remote procedure calls (RPC) between microservices, exposing the system to the dual-write problem described in Section 3.2. The assessment context supports only manually authored quizzes, with no AI-assisted question generation.

3.5.2. Yang et al. (2023) — Microservices for E-Learning

Yang, Chen, and Zhang [4] conducted a systematic literature review of 47 papers published between 2018 and 2023 on microservice decomposition, inter-service communication, and data management in e-learning systems. Their taxonomy identifies four decomposition strategies: (1) by business capability (e.g., separate services for users, courses, payments), (2) by workflow step (e.g., video ingestion as a separate service from video delivery), (3) by subdomain (consistent with DDD bounded contexts), and (4) by volatility (extracting frequently changing features into independently deployable services). The survey found that 62% of analysed systems used synchronous HTTP/REST for inter-service communication, 21% used asynchronous messaging (primarily RabbitMQ or Apache Kafka), and only 17% employed transactional outbox patterns for data consistency.

Limitations. The review identifies a significant adoption gap: while the outbox pattern is recommended in the microservices literature [6] as the standard solution for reliable event publication, fewer than one in five surveyed e-learning systems implement it. Most rely on dis-

tributed transactions (XA protocol) or best-effort eventual consistency without idempotency guarantees. The review does not address AI integration patterns, as none of the surveyed papers incorporated Whisper, LLM-based quiz generation, or similar AI services as first-class components of the architecture.

3.5.3. Kumar et al. (2023) — LLM Quiz Generation

Kumar, Barrios, and Haridas [12] present an end-to-end pipeline for generating formative assessment questions from lecture transcripts using GPT-4. Their system, evaluated on 50 university lectures spanning computer science, biology, and economics, achieved a 73% expert-acceptability rate for generated multiple-choice questions. The paper contributes a prompt engineering methodology that instructs the LLM to produce questions targeting different cognitive levels (Bloom’s taxonomy: remember, understand, apply, analyse).

Limitations. The pipeline is a standalone Python service invoked manually by instructors; it does not integrate with an LMS backend for automated lifecycle management. There is no support for human-in-the-loop review workflows, meaning that generated questions are either used directly (risk of uncaught errors) or manually copied into a separate quiz-authoring tool. The paper does not address authentication or secure callback patterns for async AI processing, and the pipeline assumes the transcript is already available—it does not include ASR integration.

3.5.4. Mehta et al. (2023) — Event-Driven Educational Analytics

Mehta, Soni, and Thakkar [9] propose an event-driven architecture for processing real-time learner interaction events in large-scale educational platforms. Their system defines a canonical event schema—`LearnerInteractionEvent` with fields for learner ID, session ID, resource ID, event type (e.g., `PAGE_VIEW`, `VIDEO_PLAY`, `QUIZ_SUBMIT`), timestamp, and arbitrary key-value metadata. Events are ingested via Apache Kafka and processed by Kafka Streams topologies that compute per-learner and per-course aggregates with sub-second latency. The authors report processing 2.8 million events per day across 15,000 learners with a p99 processing latency of 850 ms.

Limitations. The architecture depends on a Kafka cluster, which represents significant operational overhead for self-hosted deployments by small teams or institutions with limited DevOps resources. The event schema is generic and does not enforce domain-specific invariants; events are not derived from domain aggregates and therefore risk becoming disconnected from the write model’s correctness guarantees. The outbox pattern is not used, as events are published directly to Kafka from application code, exposing the system to the dual-write problem.

3.5.5. Synthesis and Gap Analysis

Table 3.1 summarises the coverage of key architectural dimensions across the four studies and NexaLearn using ✓ (full coverage), ~ (partial/reviewed), and × (not addressed). The comparison reveals that no single prior work addresses the full combination of DDD decom-

position, reliable event processing via the outbox pattern, AI-assisted transcription and quiz generation with HMAC-secured callbacks, and HLS-based video streaming with lifecycle management. NexaLearn’s contribution is the *integration* of these patterns into a cohesive, open-source backend architecture that is deployable as a modular monolith without requiring a separate message broker infrastructure.

Table 3.1: Comparative Summary of Prior Work and NexaLearn Coverage

Dimension	Dossena	Yang	Kumar	Mehta	NexaLearn
DDD bounded contexts	✓	~	×	×	✓
Clean Architecture layers	~	×	×	×	✓
Transactional outbox	×	~	×	×	✓
Idempotent event processing	×	×	×	~	✓
ASR (Whisper) integration	×	×	×	×	✓
LLM quiz generation	×	×	✓	×	✓
Human-in-the-loop review	×	×	×	×	✓
HLS video streaming	×	~	×	×	✓
HMAC callback security	×	×	×	×	✓
No broker dependency	×	~	×	×	✓
Rationale	DDD-focused; no AI or eventing	Survey only; no implementation	Standalone LLM; no backend integration	Kafka-based; no DDD or AI	All dimensions integrated with modular monolith

NexaLearn is distinguished by its comprehensive coverage across all ten dimensions. The three dimensions not covered by any prior work—Whisper integration, human-in-the-loop quiz review, and HMAC callback security—represent the principal novel contributions of this project’s architecture. The absence of a mandatory message broker (Kafka or RabbitMQ) is a deliberate design decision, justified in Section 3.6, that reduces operational complexity for the primary target audience of edtech startups and university IT departments.

The table further reveals that existing work either addresses architecture (Dossena, Yang, Mehta) or AI content (Kumar), but never both. NexaLearn’s architecture is explicitly designed to bridge this gap, making it—to the best of our knowledge—the first open-source NestJS backend to combine DDD-governed bounded contexts with a full AI pipeline, transactional outbox eventing, and Mux-based video lifecycle.

3.6. Implemented Approach and Justification

This section synthesises the literature surveyed in Sections 3.1 through 3.5 into concrete technology and pattern decisions for NexaLearn. Each decision is justified against alternatives on technical merit, alignment with the problem statement (Chapter 1), and suitability for the primary target segments identified in Section 2.1.

3.6.1. Framework Selection: NestJS over Spring Boot and Django

NexaLearn is implemented in **NestJS**[16], a progressive Node.js framework that leverages TypeScript decorators for dependency injection, controller routing, middleware piping, and module organisation. Three considerations motivate this choice:

1. **TypeScript ecosystem alignment.** The primary target segment—edtech startups and independent developers—operates predominantly in the TypeScript/JavaScript ecosystem. Choosing NestJS allows these developers to read, fork, and contribute to NexaLearn without learning a second programming language. Garcia et al. [17] found that TypeScript adoption reduces production defect density by 15% compared to equivalent JavaScript codebases, primarily through compile-time detection of type mismatches, null reference errors, and incorrect API usage.
2. **Decorator-based DDD mapping.** NestJS’s decorator syntax maps naturally to DDD architectural layers. A controller is annotated with `@Controller()`, a provider with `@Injectable()`, and a module with `@Module()`. The framework provides no prescribed entity or aggregate base class—NexaLearn’s domain entities are plain TypeScript classes—but the dependency injection container cleanly separates application-layer use cases from infrastructure adapters.
3. **Comparative framework analysis.** Zhang et al. [18] conducted a controlled experiment comparing NestJS, Spring Boot, and Django on throughput, memory consumption, and developer productivity for a representative CRUD+event API workload. NestJS achieved comparable throughput to Spring Boot (within 12%) with

40% lower memory consumption in containerised deployments, and developers with TypeScript experience completed identical implementation tasks 1.7× faster in NestJS than developers with Java experience in Spring Boot.

3.6.2. ORM Selection: Prisma over TypeORM

Prisma [19] was chosen as the object-relational mapping layer over TypeORM and MikroORM based on three criteria:

1. **Type safety.** Prisma generates TypeScript types from the database schema declaration (`prisma/schema.prisma`). Every query returns a fully typed result set, eliminating the “field does not exist on type” runtime errors common with TypeORM’s class-based entity approach. The generated types serve as a single source of truth for schema evolution.
2. **Migration management.** Prisma Migrate produces deterministic PostgreSQL migration files with rollback support, integrated into the development workflow via `prisma migrate dev` and deployment pipeline via `prisma migrate deploy`. TypeORM migrations, by contrast, require careful ordering and manual conflict resolution in team environments.
3. **Raw query support.** While Prisma’s generated client covers 95% of NexaLearn’s data access patterns, the CQRS-lite read models (Section 2.3.4) and analytics aggregations require window functions, lateral joins, and materialised view refreshes that benefit from raw SQL. Prisma exposes `prisma.$queryRawUnsafe()` with parameterised inputs, allowing these queries without escaping the typed ORM layer.

3.6.3. Transactional Outbox over Direct Message Queue Publication

NexaLearn implements the transactional outbox pattern (Section 3.2) using PostgreSQL as the event store, with no additional message broker such as RabbitMQ or Apache Kafka. This decision is motivated by the analysis of Mehta et al. [9] and the gap identified in Table 3.1:

1. **Deployment simplicity.** A PostgreSQL-only dependency means that NexaLearn deploys as a single Docker Compose stack (application containers, PostgreSQL, Redis for caching) rather than requiring a Kafka or RabbitMQ cluster. This is a first-class requirement for the primary customer segment: university IT departments and bootstrapped startups that may not have dedicated DevOps engineers to operate a distributed messaging infrastructure.
2. **Transactional guarantees.** The outbox pattern ensures that events are persisted atomically with the domain mutation, solving the dual-write problem without distributed transactions. A message broker introduces a separate transactional boundary: the application would need to write to the database and then publish to the broker, with the risk that one succeeds while the other fails.

3. **Acceptable latency.** NexaLearn’s polling interval defaults to 500 ms with a batch size of 50 events. For the workload profile of an e-learning platform serving primarily synchronous (learner-facing) and asynchronous (instructor notification, search index update) consumers, this latency is well within acceptable bounds. Events such as `LessonCompleted` that trigger real-time progress UI updates are additionally written to Redis for immediate consumption, combining the outbox’s durability with Redis’s pub/sub latency.

The trade-off is increased database load from polling and the lack of native fan-out, message broker, delayed-queue, or dead-letter features. For NexaLearn’s current workload targets (tens of thousands of enrolled learners, hundreds of events per second at peak), PostgreSQL with appropriate indexing (`idx_outbox_unprocessed` on (`processed_at` NULLS FIRST, `created_at` ASC)) comfortably handles this volume. Should event throughput grow beyond this bound, the outbox publisher is designed to be replaced with a Kafka producer without changing any domain or application layer code—the outbox processor’s only external contract is a `EventPublisher` interface that currently calls a local NestJS event bus but could be reimplemented to publish to Kafka.

3.6.4. Modular Monolith over Microservices

NexaLearn is deployed as a **modular monolith**: a single NestJS application process with multiple modules, rather than a distributed system of independently deployed microservices. The modular monolith is the deployment model recommended by Yang et al. [4] for projects and teams at NexaLearn’s scale, offering:

1. **Reduced operational overhead.** A single application instance is simpler to deploy, monitor, and debug than sixteen microservices with their own deployment pipelines, health checks, and log aggregation. For a graduation project maintained by a small core team, minimising operational surface area is a practical necessity.
2. **No network overhead.** Inter-context communication within the same process uses in-memory method calls rather than HTTP/gRPC round-trips, reducing latency by several orders of magnitude for high-frequency interactions such as “validate enrollment before returning lesson content.”
3. **Explicit bounded contexts.** The modular monolith does not relax DDD boundaries. Modules cannot import repository classes from other modules; cross-context communication is mediated through domain events (published via the in-process event bus and the outbox) or through dedicated query services exposed as NestJS providers registered in the importing module. This discipline ensures that extracting any context into an independent microservice—should scaling requirements eventually demand it—requires changing only the module’s registration and the event publishing mechanism, not the domain logic.

3.6.5. Summary of Architectural Decisions

Table 3.2 summarises the key architectural decisions, the alternatives considered, and the primary justification for each choice.

Table 3.2: Architectural Decisions and Justification

Decision	Chosen	Alternatives	Primary Justification
Framework	NestJS	Spring Boot, Django	TypeScript ecosystem alignment with target segment; decorator-based DI; lower memory in containers
ORM	Prisma	TypeORM, MikroORM	Generated type safety; deterministic migrations; raw query escape hatch
Event transport	Outbox (PostgreSQL)	Kafka, RabbitMQ	No broker dependency for simplified deployment; atomic dual-write guarantees
Deployment model	Modular monolith	Microservices	Reduced ops overhead; no network latency; extractable to microservices later
AI pipeline invocation	HMAC-signed async callback	Polling, synchronous API	Stripe-style security; decoupled processing; built-in replay prevention
Video delivery	Mux (HLS + signed URLs)	Cloudflare Stream, self-hosted FFmpeg	Managed transcoding; JWT-signed playback; webhook lifecycle events

These decisions, grounded in the literature surveyed throughout this chapter, collectively deliver an open-source e-learning backend that balances architectural rigour (DDD, Clean Architecture, outbox), practical deployability (modular monolith, PostgreSQL-only event store), and modern AI integration (Whisper, LLM quiz generation, HMAC security)—the three pillars that define NexaLearn’s contribution relative to the prior work examined in Section 3.5.

Chapter 4:

System Design and Architecture

This chapter presents the complete system design and architectural blueprint of the NexaLearn platform. The design is grounded in the Domain-Driven Design and Clean Architecture principles surveyed in Chapter 3, and directly addresses the problem sub-problems identified in Chapter 1: lack of vertical coverage in open backends, unreliable asynchronous event processing, architectural erosion across many bounded contexts, and exactly-once semantics for financial and AI integrations.

The chapter begins with a high-level overview and the assumptions that constrain the design space (Section 4.1). Section 4.2 presents the layered architecture block diagram and the module dependency graph, including a comprehensive summary table of all fifteen bounded contexts. The subsequent sections examine each module in depth: authentication (Section 4.3), course management (Section 4.4), media and video processing (Section 4.5), AI pipeline integration (Section 4.6), assessment and quizzes (Section 4.7), payments (Section 4.8), enrollment (Section 4.9), and discussion (Section 4.10). Each module section follows a consistent structure: functional description, state machine, data flow, key entities, and design decisions. Section 4.11 concludes the chapter with cross-cutting security mechanisms including HMAC service-to-service authentication, rate limiting, and audit logging.

Every module description includes a state machine diagram and references to the corresponding implementation artefacts in the NestJS codebase. Code listings highlight the domain entity behaviours that enforce business invariants at the aggregate root level.

4.1. Overview and Assumptions

NexaLearn is designed as a **modular monolith**: a single NestJS application process that hosts fifteen bounded contexts, each mapped to a NestJS module (or a coherent group of modules) following the DDD tactical patterns and Clean Architecture layering described in Section 3.1. The system is deployed via Docker Compose with PostgreSQL as the primary data store and Redis as the session cache, rate-limiter backing store, and real-time pub/sub channel for progress events.

The following assumptions constrain the design and are validated by the target customer segments identified in Section 2.1:

1. **Concurrent learner count.** The platform is designed for 500–5,000 concurrently active learners per deployment instance, corresponding to a mid-size university department or a growth-stage edtech startup. Peak read load is estimated at 1,000 requests

per second; peak write load at 200 events per second. This profile justifies the modular monolith deployment model over a microservices architecture, as the expected throughput does not require independent horizontal scaling per service.

2. **Course catalogue size.** The system supports up to 500 active courses, each containing up to 50 lessons. Total video asset storage is assumed to be under 5 TB, managed by Mux’s cloud transcoding and CDN delivery rather than self-hosted storage.
3. **Network reliability.** External provider webhooks (Stripe, Mux, AI pipeline callbacks) may be delivered at-least-once, potentially with duplicates or out-of-order delivery. The system must tolerate these without producing incorrect state.
4. **Instructor technical literacy.** Instructors are assumed to be domain experts (professors, corporate trainers) who can upload video files and review AI-generated quizzes via a web dashboard, but are not expected to write code, manage servers, or configure infrastructure.
5. **Idempotency requirement.** All external-facing write endpoints and event handlers must be idempotent, allowing safe retry by clients or webhook senders.

These assumptions shape the architectural decisions presented throughout this chapter. Where an assumption is violated by a specific deployment scenario (e.g., a large university with 50,000 students), the architecture provides documented migration paths: extracting the modular monolith into microservices (Section 3.6), adding a message broker for higher event throughput, and deploying read replicas for analytics queries.

4.2. System Architecture

4.2.1. High-Level Architecture Block Diagram

Figure 4.1 presents the layered architecture of NexaLearn following the Dependency Rule of Clean Architecture. The outermost layer (Presentation) handles HTTP concerns through NestJS controllers, guards, pipes, and interceptors. The Application layer contains use-case interactors that orchestrate domain objects without depending on infrastructure. The Domain layer holds entities, aggregates, value objects, domain events, and repository port interfaces in pure TypeScript with zero framework dependencies. The Infrastructure layer implements the repository ports using Prisma, Stripe SDK, Mux SDK, Redis client, and AWS S3 SDK. A shared Events module provides the transactional outbox publisher and the in-process event bus.

Figure 4.1 — NexaLearn System Architecture (Clean Architecture Layers)




Figure 4.1: NexaLearn System Architecture — Clean Architecture Layers

The Presentation layer routes HTTP requests to the appropriate Application use case. For example, a POST request to `/courses/:courseId/enroll` is received by `EnrollmentController`, validated by a NestJS `ValidationPipe` using class-validator DTOs, and delegated to `EnrollInFreeCourseUseCase`. The use case loads the relevant domain aggregates (`Course`, `User`), enforces business rules (course is published, user is not already enrolled), creates a new `Enrollment` aggregate, and persists it through an `EnrollmentRepository` port. The repository implementation in Infrastructure writes to PostgreSQL via Prisma and—within the same transaction—inserts an outbox event for `EnrollmentActivated`. The outbox processor subsequently publishes this event to downstream consumers (`Progress`, `Notifications`, `Search`).

4.2.2. Module Dependency Graph

Figure 4.2 shows the dependency and event-flow relationships among the bounded contexts. Solid arrows represent direct method-call dependencies (use cases in one context invoking repository or query service interfaces exported by another). Dashed arrows represent event-driven communication through the outbox: the source context publishes an event that is consumed by the target context's event handler.

Figure 4.2 — Module Interaction Diagram (Bounded Context Dependencies)




Figure 4.2: NexaLearn Module Interaction Diagram

The diagram reveals several structural patterns. **Enrollment** acts as a hub, depending on **Payments** for payment verification and on **Progress** for projection initialisation. **Courses** depends on **Media** for video asset management and on **AI-Pipeline** for transcript and quiz generation. **Assessment** publishes graded-quiz events consumed by **Progress** and **Reviews**. **Notifications** is a pure consumer: it subscribes to events from **Discussion**, **Reviews**, and **AI-Pipeline** without exposing any outward-facing dependencies. **Instructor/Analytics** is a read-heavy context that queries read models built by **Progress**, **Assessment**, and **Streak**, and has no write-side dependencies.

4.2.3. Module Summary

Table 4.1 enumerates all fifteen bounded contexts with their primary entities, published events, and module dependencies. This table serves as a quick reference for the detailed module descriptions that follow in Sections 4.3 through 4.10.

Table 4.1: Module Summary — Bounded Contexts, Entities, Events, and Dependencies

Module	Context	Key Entities	Key Events	Depends On
Auth	Auth	User, RefreshTokenSession	UserLoggedIn, UserLoggedOut, SessionRevoked	—
Courses	Courses	Course, CourseModule, Lesson, Category	CoursePublished, LessonUpdated, CourseArchived	Media, AI-PI.

Module Summary (continued)

Module	Context	Key Entities	Key Events	Depends On
Enrollment	Enrollment	Enrollment, EnrollmentEvent	EnrollmentActivated, EnrollmentCancelled, EnrollmentCompleted	Payments, Courses
Payments	Payments	PaymentIntent, PaymentTransaction, Refund	PaymentSucceeded, PaymentFailed, PaymentRefunded	Stripe (ext.)
Media	Media	VideoAsset, UploadUrl, PlaybackToken	VideoUploaded, VideoReady, VideoErrored	Mux (ext.)
AI-Pipeline	AI-Pipeline	GenerationJob, Transcript, GeneratedQuiz	JobCompleted, CallbackReceived, QuizReadyForReview	Media, OpenAI (ext.)
Assessment	Assessment	Quiz, Question, QuestionOption, QuizAttempt, ShortAnswerGrade	QuizAttemptSubmitted, AttemptGraded, QuizPublished	Courses, AI-PI.
Discussion	Discussion	Thread, ThreadInstructorReply, Post	ThreadCreated, PostAdded, InstructorReplied	—
Progress	Progress	CourseProgressView, LessonProgress, ResumeView, LessonAnalytics	ProgressUpdated, LessonCompleted, StreakUpdated	Enrollment, Assessment
Reviews	Reviews	Review, ReviewVote	ReviewSubmitted, ReviewFlagged	Assessment
Certificates	Certificates	Certificate, CertificateTemplate	CertificateIssued, CertificateRevoked	Enrollment, Progress

Module Summary (continued)

Module	Context	Key Entities	Key Events	Depends On
Search	Search	CourseSearchIndex, LessonSearchIndex	SearchIndexUpdated	Courses
Streak	Streak	StreakRecord, StreakMilestone	StreakIncremented, StreakMilestoneReached	Progress
Notifications	Notifications	Notification, EmailTemplate, NotificationPreference	NotificationSent, EmailDelivered, EmailBounced	Auth, Discussion, Assessment, Payments
Instructor	Instructor	InstructorAnalyticsView, InstructorDashboard		Progress, Assessment, Streak

4.2.4. Entity-Relationship Diagram

Figure 4.3 presents the core domain entities and their relationships, derived from the Prisma schema and domain entity definitions. Entities are colour-coded by bounded context: blue for Auth/Identity, green for Course Content, orange for Enrollment and Payments, purple for Assessment, red for Discussion, and teal for Media and AI Pipeline. Arrowheads denote one-to-many (1:*) and one-to-one (1:1) relationships as enforced by the database schema and application-layer invariants.

Figure 4.x — Entity-Relationship Diagram (Core Domain)




Figure 4.3: Entity-Relationship Diagram (Core Domain)

The diagram highlights several architectural decisions. **User** is the central actor, connected to Enrollments, QuizAttempts, Reviews, and Discussion entities. **Course** forms a three-level hierarchy (Course → Module → Lesson), with Lesson linking to a single Quiz and optionally to VideoAssets or GenerationJobs. **Enrollment** and **PaymentIntent** both reference User and Course but serve distinct lifecycle purposes. **QuizAttempt** bridges the Quiz (authoring aggregate) with the User (student runtime), recording per-question answers through AttemptAnswer. The full database schema with all fields, indexes, and constraints is available in the `prisma/schema.prisma` file in the project repository.

The next eight sections examine the core modules—Auth, Courses, Media, AI-Pipeline, Assessment, Payments, Enrollment, and Discussion—in architectural depth. Each section includes the domain aggregate design, state machine diagram, key code listings, and the design trade-offs that shaped the implementation.

4.3. Authentication and Session Management Module

4.3.1. Functional Description

The Auth module handles user registration, credential-based login, token issuance, session management, and logout. It implements a dual-token JWT architecture with server-side refresh token rotation and reuse detection, following best practices for mitigating token theft described by the IETF OAuth 2.0 Security BCP. The module depends on Redis for storing refresh token sessions and on PostgreSQL for persistent user records.

All authentication flows assume HTTPS transport. Tokens are opaque to clients in the

sense that the access token contains only claims necessary for authorisation (user ID, role, expiration); no sensitive data is embedded. The refresh token is an opaque UUID that serves as a key into the Redis session store.

4.3.2. JWT Token Architecture

The system issues two token types:

Access Token. A short-lived (15-minute) JWT signed with the configured JWT secret (HS256). The token encodes: `sub` (user ID), `roles` (student, instructor, admin), `iat` (issued at), and `exp` (expiration). The access token is transmitted as `Authorization: Bearer <token>`. It is validated by a NestJS `JwtAuthGuard` that verifies the HS256 signature using the configured secret, checks expiration, and attaches the decoded payload to the request context. No database lookup is performed on access token validation, enabling $O(1)$ verification.

Refresh Token. A 64-character cryptographically random string generated via `crypto.randomBytes(32)`. The refresh token is hashed using SHA-256 before storage; the raw token is returned to the client exactly once. The hashed token, along with the user ID, device metadata, and expiration timestamp (7 days), is stored in Redis under key `session:<hashed_token>` with an appropriate TTL.

The separation of concerns mirrors the pattern used by Auth0 and Firebase Authentication: the access token is a stateless assertion verified by any service without a central session store, while the refresh token enables long-lived sessions without storing secrets on the client beyond the browser's secure, HTTPOnly cookie storage.

4.3.3. Redis Session Store Design

The Redis session store provides three operations:

1. **Session creation.** On login, the system generates a refresh token, hashes it, and stores the session record: `SET session:<hash> user_id=<uid> role=<role> device=<ua> created_at=<ts> EX 604800`. The raw token is returned to the client.
2. **Session validation.** On token refresh, the system hashes the provided refresh token, retrieves the session from Redis, and checks that the user ID matches and the session is not revoked. If valid, a new access–refresh token pair is issued and the old session is deleted.
3. **Session revocation.** On logout or administrative revocation, the session key is deleted from Redis. Any subsequent attempt to use the associated refresh token fails validation.

Redis is chosen over PostgreSQL for session storage because: (1) TTL-based expiration is native to Redis and eliminates the need for a session-cleanup cron job, (2) read latency is sub-millisecond compared to 1–5 ms for PostgreSQL, and (3) the session data is ephemeral—losing all sessions on a Redis restart is acceptable as users can re-authenticate.

4.3.4. Refresh Token Rotation and Reuse Detection

The Auth module implements refresh token rotation: each time a refresh token is used to obtain a new access token, the old refresh token is invalidated and a new one is issued. This limits the window of vulnerability if a refresh token is stolen—the attacker can use it only once before rotation invalidates it.

Reuse detection handles the edge case where both the legitimate client and an attacker attempt to use the same refresh token concurrently. The detection algorithm is:

```
1 async function refreshTokens(rawToken: string): Promise<TokenPair>
2   {
3     const hash = sha256(rawToken);
4     const session = await redis.get(`session:${hash}`);
5     if (!session) throw new UnauthorizedException('Session not found');
6
7     await redis.del(`session:${hash}`);
8
9     // Issue new tokens
10    const newRawToken = crypto.randomBytes(32).toString('hex');
11    const newHash = sha256(newRawToken);
12    await redis.set(`session:${newHash}`, session, 'EX', 604800);
13
14    const accessToken = await this.jwtService.signAsync({
15      sub: session.userId,
16      roles: session.roles,
17    });
18
19    return { accessToken, refreshToken: newRawToken };
20  }
```

Listing 4.1: Refresh Token Reuse Detection

If two requests arrive with the same refresh token, the first deletes the session and creates a new one; the second finds the session missing and returns an error. This forces the legitimate client to re-authenticate, which is a minor inconvenience compared to the alternative of an attacker maintaining persistent access.

The state machine diagram (Figure 4.4) illustrates the session lifecycle.




Figure 4.4: Auth Session State Machine

4.4. Course Management Module

4.4.1. Functional Description

The Courses module manages the curriculum lifecycle: instructors create courses, organise them into modules and lessons, attach video assets and quizzes to lessons, and publish the result for learner enrollment. The module is the content backbone of the platform, and its aggregate design must support deep nesting (course → module → lesson), reordering, versioning, and concurrent editing.

The module exposes REST endpoints for CRUD operations on courses, modules, and lessons, with dedicated endpoints for publishing/unpublishing, reordering, and category assignment.

4.4.2. DDD Aggregate Design: Course–Module–Lesson Hierarchy

The aggregate root is `CourseEntity`, which owns a collection of `CourseModuleEntity` value objects, each of which owns a collection of `LessonEntity` entities. The hierarchy is enforced by the aggregate root: all operations that modify modules or lessons must go through `CourseEntity` methods.

```

1 export class CourseEntity {
2   private constructor(
3     public readonly id: string,
4     public title: string,
5     public description: string,
6     public instructorId: string,
7     public price: Money,
8     public status: CourseStatus,
9     public modules: CourseModuleEntity[],
10    public version: number,
11    public categoryId: string,
12    public thumbnailUrl?: string,
13  ) {}

```

```

14
15 static create(props: CreateCourseProps): CourseEntity {
16     return new CourseEntity(
17         props.id, props.title, props.description,
18         props.instructorId, props.price, CourseStatus.DRAFT,
19         [], 1, props.categoryId,
20     );
21 }
22
23 publish(): void {
24     if (this.modules.length === 0)
25         throw new BusinessRuleViolation('Course must have at least
26             one module');
27     if (this.modules.some(m => m.lessons.length === 0))
28         throw new BusinessRuleViolation('Each module must have at
29             least one lesson');
30     if (!this.thumbnailUrl)
31         throw new BusinessRuleViolation('Course must have a thumbnail
32             ');
33     this.status = CourseStatus.PUBLISHED;
34     this.version++;
35 }
36
37 reorderModules(moduleIds: string[]): void {
38     const ordered = moduleIds.map(id => {
39         const mod = this.modules.find(m => m.id === id);
40         if (!mod) throw new BusinessRuleViolation(`Module ${id} not
41             found`);
42         return mod;
43     });
44     this.modules = ordered.map((m, i) => m.withPosition(i + 1));
45     this.version++;
46 }
47 }

```

Listing 4.2: Course Aggregate Root

The aggregate enforces three invariants at publication time: a course must have at least one module, every module must have at least one lesson, and a thumbnail must be set. These rules are checked in the `publish()` method and cannot be bypassed by external code.

4.4.3. Optimistic Concurrency Control

Each `CourseEntity` carries a `version` field that is incremented on every mutating operation. The Prisma repository implementation includes the version in the `WHERE` clause of the write query:

```

1  async update(course: CourseEntity): Promise<void> {
2      const result = await this.prisma.course.updateMany({
3          where: { id: course.id, version: course.version - 1 },
4          data: {
5              title: course.title,
6              status: course.status,
7              modules: { deleteMany: {}, create: this.serializeModules(
8                  course) },
9              version: course.version,
10         },
11     });
12
13     if (result.count === 0)
14         throw new ConcurrentModificationError(
15             `Course ${course.id} was modified by another transaction`
16         );
17 }

```

Listing 4.3: Optimistic Concurrency Update

If two instructors open the same course, both modify it, and both submit their changes, the second write will match zero rows (because the version has been incremented by the first write) and the repository throws `ConcurrentModificationError`. The use case can catch this error and present the instructor with a conflict-resolution dialogue. This pattern avoids pessimistic locking (`SELECT FOR UPDATE`) which would block read access to course content during editing — an unacceptable trade-off given that concurrent read access by enrolled learners is the common case.

4.4.4. Idempotency for Lesson and Module Operations

All mutating endpoints in the Courses module accept an optional `Idempotency-Key` header. The system maintains an idempotency ledger table (`course_idempotency_keys`) keyed by the SHA-256 hash of the key. If the same key is received within a configurable time window (default 24 hours), the response from the first request is returned without re-executing the operation. This is essential for mobile clients on unreliable networks, where a create-lesson request might succeed but the client never receives the response and retries, potentially creating duplicate lessons.

The course curriculum lifecycle is summarised by the state machine in Figure 4.5 showing transitions between `DRAFT`, `UNDER_REVIEW`, `PUBLISHED`, and `ARCHIVED`.




Figure 4.5: Course Status Lifecycle

4.5. Media and Video Processing Module

4.5.1. Functional Description

The Media module manages the full lifecycle of video assets from upload to delivery. It integrates with Mux (Section 3.4) for transcoding, storage, and CDN delivery, while maintaining its own domain model for asset state and metadata. The module provides REST endpoints for creating direct upload URLs, querying asset status, and generating signed playback tokens for enrolled learners.

The module interacts closely with the AI Pipeline module (Section 4.6): when a video reaches the READY state, the Media module publishes a `VideoReady` event that triggers automatic transcript generation.

4.5.2. Mux Integration and VideoAsset State Machine

The `VideoAsset` aggregate root tracks the lifecycle of each uploaded video through the states illustrated in Figure 4.6.




Figure 4.6: Video Asset State Machine

The lifecycle proceeds as follows:

1. **AWAITING_UPLOAD.** A lesson is created, and an instructor calls `POST /instructor/video-assets/`

to request a direct upload URL. The server calls Mux's `Upload.create()` API, obtains a signed upload URL, and stores the `upload_id` in the `VideoAsset` entity.

2. **UPLOADED.** The client PUTs the video file directly to Mux using the signed URL. The server does not proxy the file, avoiding bandwidth costs and latency. The state transitions to **UPLOADED** once Mux confirms receipt.
3. **TRANSCODING.** Mux sends `video.upload.asset_created`, indicating that the source file is being transcoded into HLS renditions. The Media module updates the state and creates a `GenerationJob` in the AI Pipeline module with status **QUEUED**.
4. **READY.** Mux sends `video.asset.ready` with the playback ID. The Media module stores the playback ID, sets the signed URL expiry policy, and transitions the asset to **READY**. An outbox event `VideoReady` is published, triggering the AI pipeline to begin transcript generation.
5. **FAILED.** Mux sends `video.asset.error` with an error code. The state transitions to **FAILED**, and the instructor is notified via the Notifications module.

4.5.3. Signed Playback URL Generation

Playback URLs are generated by the backend on request and are not stored directly. When a learner accesses a lesson, the backend verifies that the learner has an active enrollment for the course containing that lesson, then generates a JWT-signed playback URL with a 4-hour expiration:

```

1 async function getPlaybackUrl(
2   learnerId: string, lessonId: string
3 ): Promise<string> {
4   const enrollment = await this.enrollmentRepo
5     .findActiveByLearnerAndLesson(learnerId, lessonId);
6   if (!enrollment)
7     throw new ForbiddenException('No active enrollment');
8
9   const lesson = await this.lessonRepo.findByIdOrThrow(lessonId);
10  const token = this.muxJwtService.sign({
11    sub: learnerId,
12    aud: `playback:${lesson.videoAsset.playbackId}`,
13    exp: Math.floor(Date.now() / 1000) + 14400, // 4 hours
14  });
15
16  return `https://stream.mux.com/${lesson.videoAsset.playbackId}.
17    m3u8?token=${token}`;

```

Listing 4.4: Generating a Signed Mux Playback URL

This design ensures that: (1) learners cannot share playback URLs beyond the 4-hour window; (2) revoked enrollments are immediately reflected—the next request generates a playback URL only if the enrollment is still active; and (3) unauthorised users never receive a valid playback URL because the authorization check happens server-side before the JWT is signed.

Figure 4.7 illustrates the end-to-end flow. The instructor uploads a video through the Media module, which obtains a Mux upload URL without proxying the file. The client uploads directly to Mux. Asynchronous webhooks drive the state machine: `video.upload.asset_created` transitions to `UPLOADED` and then `TRANSCODING`; `video.asset.ready` transitions to `READY` and stores the playback ID. When a learner later requests the lesson content, the server verifies the learner has an active enrollment, generates a signed JWT playback URL with a 4-hour expiry, and the client streams HLS segments directly from the Mux CDN.




Figure 4.7: Sequence Diagram: Video Upload, Transcoding & Playback

4.6. AI Pipeline Integration Module

4.6.1. Functional Description

The AI Pipeline module orchestrates the automated generation of transcripts and quiz questions from video content using external AI services. It exposes internal-use-only endpoints for AI service callbacks (HMAC-authenticated) and publishes domain events that trigger notifications for instructor review. The module manages two distinct workflows: **transcript generation** (Whisper ASR) and **quiz generation** (LLM).

Both workflows follow an asynchronous callback pattern: NexaLearn dispatches a request to the AI service (via HTTP or queue), the service processes the job, and the ser-

vice calls back a NexaLearn endpoint with the result. This decouples processing time—transcribing a 60-minute lecture may take several minutes—from the API response time.

4.6.2. HMAC Signature Verification

All callbacks from the AI pipeline include an X-Pipeline-Signature header with format `t=<unix-timestamp>,v1=<hex-digest>`. The recipient controller verifies the signature before processing:

```

1 @Post()
2 async handleGenerationCallback(
3   @Headers('x-pipeline-signature') signature: string,
4   @Body() payload: GenerationCallbackDto,
5 ): Promise<void> {
6   this.hmacService.verify(signature, payload);
7   await this.useCase.execute(payload);
8 }
9
10 // In HmacService
11 verify(signature: string, body: unknown): void {
12   const [tsPart, hashPart] = signature.split(',');
13   const timestamp = parseInt(tsPart.replace('t=', ''), 10);
14   const receivedHash = hashPart.replace('v1=', '');
15
16   if (Math.abs(Date.now() / 1000 - timestamp) > 300)
17     throw new SecurityException('Callback timestamp expired');
18
19   const computed = createHmac('sha256', this.secret)
20     .update(`${timestamp}.${JSON.stringify(body)}`)
21     .digest('hex');
22
23   if (!timingSafeEqual(computed, receivedHash))
24     throw new SecurityException('HMAC signature mismatch');
25 }

```

Listing 4.5: HMAC Callback Verification

The timestamp tolerance (300 seconds) prevents replay attacks without requiring clock synchronisation tighter than typical NTP guarantees. The `timingSafeEqual` comparison prevents timing side-channel attacks.

4.6.3. HandleGenerationCallbackUseCase — Exactly-Once Processing

The callback handler must be idempotent: the AI service may retry callbacks if it does not receive a 200 response. The `HandleGenerationCallbackUseCase` uses the `GenerationJob`

aggregate to ensure each callback is processed exactly once:

```

1  async execute(dto: GenerationCallbackDto): Promise<void> {
2    await this.prisma.$transaction(async (tx) => {
3      const job = await tx.generationJob.findUniqueOrThrow({
4        where: { id: dto.jobId },
5      });
6
7      if (job.status !== 'PROCESSING') return; // already processed
8
9      await tx.generationJob.update({
10        where: { id: dto.jobId },
11        data: {
12          status: 'READY_FOR_REVIEW',
13          result: dto.result,
14          completedAt: new Date(),
15        },
16      });
17
18      await tx.outboxEvent.create({
19        data: {
20          aggregateType: 'GenerationJob',
21          aggregateId: dto.jobId,
22          eventType: 'GenerationJobCompleted',
23          payload: { result: dto.result, videoAssetId: job.
24                    videoAssetId },
25        },
26      });
27    }

```

Listing 4.6: Exactly-Once Callback Processing

The guard clause (`if (job.status !== 'PROCESSING') return`) ensures that if the same callback arrives twice, the second invocation is a no-op. The outbox event triggers downstream handlers that create the `GeneratedQuiz` record (if quiz generation is enabled) and notify the instructor.

4.6.4. Quiz Review Workflow

Generated quizzes transition through `READY_FOR_REVIEW` status once the AI pipeline completes processing. The instructor reviews each proposed question in a dashboard, edits the stem or options, and manually publishes the quiz to make it available to learners. The `GenerationJob` aggregate tracks the statuses `QUEUED` (awaiting dispatch), `PROCESSING` (AI

service working), `READY_FOR_REVIEW` (results available for instructor review), and `FAILED` (processing error). This human-in-the-loop workflow, illustrated in Figure 4.9, ensures quality control without requiring the instructor to author questions from scratch.

Figure 4.8 details the asynchronous orchestration. The instructor uploads a video through the Media module, which requests a Mux upload URL and returns it to the client. The client uploads directly to Mux. When Mux signals the asset is created, the Media module transitions the `VideoAsset` to `UPLOADED` and emits a domain event consumed by the AI Pipeline module. The pipeline creates a `GenerationJob` (`QUEUED`), dispatches audio to Whisper ASR for transcription, and—once the transcript is ready—dispatches the transcript text to an LLM for quiz question generation. Both callbacks are HMAC-authenticated and idempotent (Section 4.6). The final quiz questions are persisted with status `READY_FOR_REVIEW`, and the instructor is notified.




Figure 4.8: Sequence Diagram: AI Pipeline — Transcript & Quiz Generation




Figure 4.9: AI Generation Job State Machine

4.7. Assessment and Quiz Module

4.7.1. Functional Description

The Assessment module is one of the most domain-rich contexts in NexaLearn. It manages the full lifecycle of quizzes, questions, and learner attempts. The module supports multiple question types (multiple-choice, true-or-false, short-answer), configurable passing scores, time limits, randomisation, and both automatic and AI-assisted (short-answer) grading. Quiz results feed into the Progress module for learner analytics and into the Reviews module for peer assessment of short-answer responses.

4.7.2. Quiz Authoring Aggregate

The QuizEntity aggregate root enforces the invariants of quiz authoring:

```

1 export class QuizEntity {
2   constructor(
3     public readonly id: string,
4     public lessonId: string,
5     public title: string,
6     public instructions: string,
7     public passingScore: number,      // 0--100
8     public maxAttempts: number,      // 0 = unlimited
9     public timeLimitMinutes: number | null,
10    public shuffleQuestions: boolean,
11    public questions: QuestionEntity[],
12    public status: QuizStatus,
13  ) {}
14
15  addQuestion(question: QuestionEntity): void {
16    if (this.status !== QuizStatus.DRAFT)
  
```

```
17     throw new BusinessRuleViolation('Cannot modify published quiz
18         ');
19     this.questions.push(question);
20 }
21
22 publish(): void {
23     if (this.questions.length < 1)
24         throw new BusinessRuleViolation('Quiz must have at least 1
25             question');
26     if (this.questions.some(q => q.options.length < 2))
27         throw new BusinessRuleViolation('Each question needs at least
28             2 options');
29     this.status = QuizStatus.PUBLISHED;
30 }
31
32 gradeAttempt(
33     answers: AnswerDto[],
34     shortAnswerGradingService: ShortAnswerGradingService,
35 ): QuizAttemptEntity {
36     let correctCount = 0;
37     const gradedAnswers = answers.map((answer) => {
38         const question = this.questions.find(q => q.id === answer.
39             questionId);
40         if (!question) throw new BusinessRuleViolation('Question not
41             found');
42
43         if (question.type === QuestionType.SHORT_ANSWER) {
44             return {
45                 questionId: question.id,
46                 status: 'PENDING_GRADING',
47                 pointsAwarded: null,
48             };
49         }
50
51         const isCorrect = question.correctOptionId === answer.
52             selectedOptionId;
53         if (isCorrect) correctCount++;
54         return {
55             questionId: question.id,
56             selectedOptionId: answer.selectedOptionId,
57             isCorrect,
```

```

52     pointsAwarded: isCorrect ? question.points : 0,
53   };
54 });
55
56   const score = Math.round((correctCount / this.questions.length)
57     * 100);
57   const passed = score >= this.passingScore;
58
59   const hasPendingGrading = gradedAnswers.some(a => a.status ===
60     'PENDING_GRADING');
61
62   return QuizAttemptEntity.create({
63     quizId: this.id,
64     answers: gradedAnswers,
65     score,
66     passed,
67     status: hasPendingGrading ? QuizAttemptStatus.IN_PROGRESS :
68       QuizAttemptStatus.GRADED,
69   });

```

Listing 4.7: QuizEntity Aggregate Root Showing Invariants

The aggregate root is the sole authority for grading: it receives learner answers, computes scores by comparing against the correct option for each multiple-choice question, and flags short-answer questions for asynchronous AI grading.

4.7.3. Quiz Attempt Lifecycle

Figure 4.10 shows the states of a quiz attempt.




Figure 4.10: Quiz Attempt State Machine

When a learner starts a quiz, an attempt is created in `IN_PROGRESS` state with a snapshot of the quiz questions (ensuring the quiz cannot change mid-attempt). When the learner submits, the attempt transitions to `SUBMITTED`. If all questions are auto-gradable (multiple-choice, true-or-false), the attempt is graded synchronously and transitions to `GRADED`. If the quiz includes short-answer questions, the attempt remains in `SUBMITTED` while the grading callback is awaited. When the AI pipeline responds, the attempt transitions to `GRADED` and the updated score is persisted. Learners may also abandon an attempt, transitioning it to `ABANDONED`.

4.7.4. Asynchronous Short-Answer Grading

Short-answer questions cannot be graded by simple string comparison—the learner’s answer may be semantically correct regardless of exact wording. The Assessment module delegates grading to an LLM (configured per deployment as OpenAI GPT-4 or a self-hosted alternative). The grading prompt includes the question stem, the instructor-provided rubric, and the learner’s answer, and returns a grade (0–100), a justification, and optionally feedback for the learner. The grader endpoint is call-back authenticated using the same HMAC pattern described in Section 4.6. Until grading completes, the learner sees a “Grading in progress” status. When grading finishes, the learner is notified, and the updated score propagates to the Progress module via an outbox event.

4.8. Payment Module

4.8.1. Functional Description

The Payment module manages the lifecycle of financial transactions between learners and the platform. It integrates with Stripe for payment processing and webhook delivery, maintains its own payment intent state machine for reconciliation, and coordinates with the Enrollment module to activate course access only after successful payment. The module also handles refunds, partial payments, and coupon applications.

The module is designed to be Stripe-agnostic at the domain layer: the `PaymentIntent` aggregate and its state machine use a generic set of states and commands. A dedicated `StripeInfrastructureService` maps Stripe-specific constructs (`PaymentIntent` API objects, webhook event types) to the domain model, allowing the payment provider to be swapped without modifying domain logic.

4.8.2. Stripe Integration and Webhook Handling

Checkout flow proceeds as follows:

1. The client requests `POST /payments/intent` with course ID and learner ID. The Payment module creates a Stripe Payment Intent via the Stripe SDK and returns the client secret to the client.
2. The client completes payment on Stripe’s hosted checkout page. The learner completes payment on Stripe’s domain, which reduces PCI-DSS compliance scope for the

platform.

3. Stripe sends a `checkout.session.completed` webhook to NexaLearn's `POST /payments/webhook`. The webhook handler verifies the Stripe signature using `stripe.webhooks.constructEvent()`, retrieves the session details, and invokes the `HandlePaymentSucceededUseCase`.
4. The use case loads or creates a `PaymentIntent` aggregate, transitions it to `SUCCEEDED`, and publishes a `PaymentSucceeded` outbox event. The Enrollment module's event handler creates or activates the corresponding enrollment.

Webhook idempotency is critical—Stripe may deliver the same webhook event multiple times. Stripe events include a unique `idempotency_key` in their HTTP header; NexaLearn records processed webhook IDs in a ledger table and ignores duplicates:

```

1 async function handleWebhook(event: Stripe.Event): Promise<void> {
2   const processed = await this.webhookLedger.findUnique({
3     where: { stripeEventId: event.id },
4   });
5   if (processed) return; // already handled
6
7   await this.prisma.$transaction(async (tx) => {
8     await tx.webhookLedger.create({
9       data: { stripeEventId: event.id, processedAt: new Date() },
10    });
11    // process event...
12  });
13 }
```

Listing 4.8: Idempotent Stripe Webhook Processing

4.8.3. Payment Intent State Machine

Figure 4.11 shows the `PaymentIntent` state machine.




Figure 4.11: Payment Intent State Machine

The `PaymentIntent` aggregate transitions through states: `CREATED` when the payment is initiated, `PENDING` while awaiting Stripe confirmation, `SUCCEEDED` upon confirmed payment, `FAILED` if the payment was declined or expired, and `REFUNDED` when a succeeded payment is refunded. Each transition is recorded as an immutable `PaymentEvent` in a separate event table, providing an audit trail for financial reconciliation.

4.8.4. Reconciliation and Expiration Services

Two scheduled services handle edge cases in the payment lifecycle:

ReconciliationService. Runs daily and retrieves all `PaymentIntents` in `PENDING` state older than 24 hours. It queries Stripe for the latest status of each associated `PaymentIntent` and updates the local state accordingly, handling cases where the webhook was never delivered (e.g., network outage).

ExpirationService. Runs hourly and expires `Payment Intents` older than 24 hours. This garbage-collects abandoned payment flows and releases any temporary holds on inventory or enrollment slots.

Both services are implemented as NestJS `@Cron` decorators with configurable cron expressions and are disabled in test environments.

4.9. Enrollment Module

4.9.1. Free vs. Paid Enrollment Flow

The Enrollment module manages how learners gain access to courses. Two enrollment flows exist:

Free enrollment. If a course price is zero, the enrollment is created and activated immediately within a single database transaction. The `EnrollmentController` invokes `CreateFreeEnrollmentUseCase`, which creates the `Enrollment` aggregate in `ACTIVE` state and publishes an `EnrollmentActivated` event. This event triggers the `Progress` module to initialise the learner's progress projections for all lessons in the course curriculum.

Paid enrollment. If a course has a non-zero price, the enrollment is created in `PENDING` state. The Enrollment module does not handle payment directly; it waits for an out-bound event from the Payment module. When `PaymentSucceeded` is consumed by the `OnPaymentSucceeded` event handler in the Enrollment module, the handler loads the corresponding enrollment (matched by a correlation ID stored in the Stripe Checkout Session metadata), activates it, and publishes `EnrollmentActivated`.

The explicit separation of free and paid paths ensures that paid enrollments cannot accidentally be activated without payment confirmation through the Payment module's state machine.

4.9.2. Idempotency and Concurrent Enrollment Protection

Enrollment creation is protected against concurrent requests. The `CreateEnrollmentUseCase` checks for an existing active enrollment before creating a new one:

```
1  async execute(dto: CreateEnrollmentDto): Promise<Enrollment> {
2    return this.prisma.$transaction(async (tx) => {
3      const existing = await tx.enrollment.findFirst({
4        where: {
5          learnerId: dto.learnerId,
6          courseId: dto.courseId,
7          status: { in: ['ACTIVE', 'REFUNDED'] },
8        },
9      });
10
11     if (existing) {
12       if (dto.idempotencyKey) return existing;
13       throw new BusinessRuleViolation('Already_enrolled_in_this_
14         course');
15     }
16
17     const enrollment = EnrollmentEntity.create(
18       dto.learnerId, dto.courseId,
19       dto.price > 0 ? EnrollmentStatus.PENDING : EnrollmentStatus.
20         ACTIVE,
21     );
22
23     await tx.enrollment.create({ data: this.mapToPersistence(
24       enrollment) });
25
26     if (enrollment.isActive) {
27       await tx.outboxEvent.create({
28         data: {
29           aggregateType: 'Enrollment',
30           aggregateId: enrollment.id,
31           eventType: 'EnrollmentActivated',
32           payload: { learnerId: dto.learnerId, courseId: dto.
33             courseId },
34         },
35       });
36     }
37
38     return enrollment;
39   });
40 }
```

```

35   });
36 }

```

Listing 4.9: Concurrent Enrollment Protection

The `findFirst` lookup with status filter uses a composite index on `(learnerId, courseId, status)` for efficient querying. The entire operation—check, create, insert outbox event—occurs within a single Prisma transaction, preventing race conditions where two concurrent requests could both pass the existence check.

Figure 4.12 shows the enrollment lifecycle.




Figure 4.12: Enrollment State Machine

Figure 4.13 illustrates the full interaction for a paid enrollment. The learner first browses the course catalogue, then sends an enrollment request. The Enrollment module checks the course price and creates a PENDING enrollment. The Payment module creates a Stripe Payment Intent; the learner completes payment on the frontend. Stripe sends a webhook to the Payment module, which marks the intent as SUCCEEDED. An outbox event triggers the Enrollment module to activate the enrollment. The learner's next poll sees the enrollment in ACTIVE status and gains access to course content.




Figure 4.13: Sequence Diagram: Paid Enrollment Flow

Enrollments transition from PENDING to ACTIVE (payment confirmed or free enrollment), from ACTIVE to CANCELLED (learner or instructor cancelled), from PENDING to CANCELLED (learner abandoned checkout or manually cancelled), and from ACTIVE to REFUNDED (payment refunded).

4.10. Discussion and Q&A Module

4.10.1. Thread and Post Aggregate Design

The Discussion module implements a tree-structured thread-and-post model common to forum and Q&A systems. The aggregate root is `ThreadEntity`, which is scoped to a specific lesson (per-lesson discussion) or to a course (general discussion). A thread has a title, a creator (learner or instructor), a collection of `PostEntity` children, and a lifecycle status.

The aggregate enforces the forum structure invariants:

```

1 export class ThreadEntity {
2   constructor(
3     public readonly id: string,
4     public readonly courseId: string,
5     public readonly lessonId: string | null,
6     public readonly creatorId: string,
7     public title: string,
8     public body: string,
9     public isPinned: boolean,
10    public isResolved: boolean,

```

```

11     public isHidden: boolean,
12     public posts: PostEntity[],
13     public createdAt: Date,
14     public updatedAt: Date,
15   ) {}

16
17   addPost(learnerId: string, body: string): PostEntity {
18     if (this.isHidden)
19       throw new BusinessRuleViolation('Thread_is_closed');
20     if (body.length > 5000)
21       throw new BusinessRuleViolation('Post_exceeds_5000_characters');
22
23     const post = PostEntity.create({
24       threadId: this.id,
25       authorId: learnerId,
26       body,
27     });
28
29     this.posts.push(post);
30     this.updatedAt = new Date();
31     return post;
32   }
33
34   markResolved(): void {
35     this.isResolved = true;
36     this.updatedAt = new Date();
37   }
38 }

```

Listing 4.10: Discussion Thread Aggregate

When a post is added by an instructor to a learner's thread, the `InstructorReplied` domain event is published. This event is consumed by the Notifications module, which sends an email notification to the thread creator. The notification includes a direct link to the thread in the learner dashboard.

4.10.2. Instructor Reply Notifications

Notification flow: `ThreadEntity.addPost()` creates a `PostEntity` and marks the thread as updated. The `AddPostUseCase` persists the thread and post, then publishes an outbox event. The Notifications module's event handler checks whether the author of the new post has the instructor role on the course. If so, it creates a `Notification` entity for each learner who has posted in the thread, containing the instructor's reply excerpt and a link to the thread.

Notifications are delivered via Resend transactional email and are also accessible via the NexaLearn REST API for in-app notification display.

Figure 4.14 illustrates the thread lifecycle using boolean flags: a thread is open by default, marked as resolved by the instructor, or hidden through moderation.




Figure 4.14: Discussion Thread State Machine

4.11. Security Design

4.11.1. HMAC Service-to-Service Authentication

NexaLearn uses HMAC-based authentication for internal service-to-service communication when external webhook callbacks are received. Three external services authenticate via HMAC: Stripe (webhook signatures using `stripe-webhook-secret`), Mux (webhook signatures using `mux-webhook-secret`), and the AI Pipeline (custom HMAC using a shared secret configured per deployment).

The implementation uses a generic `HmacGuard` NestJS guard that extracts the signature from the request header, computes the expected HMAC using the request body and the configured secret for the service, and rejects requests that do not match within a time-tolerance window:

```

1 @Injectable()
2 export class HmacWebhookGuard implements CanActivate {
3   constructor(
4     private readonly configService: ConfigService,
5     private readonly provider: 'stripe' | 'mux' | 'ai-pipeline',
6   ) {}
7
8   async canActivate(context: ExecutionContext): Promise<boolean> {
9     const request = context.switchToHttp().getRequest();
10    const signature = request.headers[this.headerName()] as string;
11    const body = JSON.stringify(request.body);
12
13    if (!signature) return false;
14
15    const secret = this.configService.get<string>(
16      `${this.provider}_WEBHOOK_SECRET`
17    );
18  }
19 }
  
```

```

17     );
18
19     return this.verify(signature, body, secret);
20 }
21
22 private verify(signature: string, body: string, secret: string):
23     boolean {
24     // Provider-specific verification (Stripe SDK, Mux SDK, or
25     // custom HMAC)
26     // ...
27 }
28 }

```

Listing 4.11: Generic HMAC Webhook Guard

This guards can be applied to specific controller routes via NestJS decorators: `@UseGuards(new HmacWebhookGuard('stripe'))`.

4.11.2. Rate Limiting

NexaLearn implements distributed rate limiting using the `@nestjs/throttler` module with a Redis-backed storage adapter. The rate limiter tracks request counts per IP address and per authenticated user ID, with separate limits for read endpoints (100 requests per minute per user) and write endpoints (20 requests per minute per user). Exceeding the limit returns HTTP 429 Too Many Requests with a `Retry-After` header.

The Redis-backed `ThrottlerModule` uses a sliding window algorithm: each request increments a counter in Redis with a TTL equal to the window duration, and counters are checked before allowing the request. This design ensures that rate limits are consistent across multiple application instances (which is essential for the modular monolith if horizontally scaled behind a load balancer).

Additional throttling is applied to sensitive endpoints: login (5 attempts per minute per IP), password reset (3 attempts per hour per user), and AI pipeline callbacks (no rate limit for internal callbacks, but validated by HMAC).

4.11.3. Audit Logging

The Audit module provides an immutable, queryable log of all state-changing operations across the platform. Each audit entry records:

- **Actor.** The user ID and role that performed the operation (or `SYSTEM` for automated actions).
- **Action.** A canonical action name (e.g., `enrollment.activate`, `course.publish`, `payment.refund`).
- **Resource.** The aggregate type and ID affected by the action.

- **Context.** The request IP, user agent, and correlation ID (propagated via a NestJS middleware that assigns a UUID to each incoming HTTP request).
- **Diff.** A JSON snapshot of the relevant aggregate state before and after the operation (optional, configurable per action type).

Audit entries are written to a dedicated `audit_log` PostgreSQL table. Writes are fire-and-forget (no outbox): if the audit write fails, the primary operation is not rolled back. This trade-off—audit loss in rare failure scenarios versus operational complexity—is acceptable because the financial trail is already guaranteed by the Payment module’s event ledger.

The audit log is exposed to administrators via a read-only REST endpoint with pagination and filtering by actor, action, resource type, and date range. Retention is configurable; the default policy retains audit logs for one year, after which they are archived to cold storage.

Chapter 5:

System Testing and Validation

This chapter presents the testing methodology, infrastructure, and results that validate the correctness, reliability, and performance of the NexaLearn platform described in Chapter 4. The testing strategy follows the IEEE 829–2008 standard for software test documentation and adopts a layered approach: unit tests verify domain logic in isolation, integration tests validate inter-module contracts against real database instances, and end-to-end tests exercise complete user workflows through the HTTP API.

Section 5.1 describes the toolchain and environment configuration. Section 5.2 outlines the three-tier testing pyramid and the rationale for each level. Sections 5.2.3 through 5.2.5 present detailed coverage tables, representative test scenarios, and code-level examples. Section 5.3 documents the testing timeline across the development lifecycle. Section 5.4 compares the NexaLearn testing approach against open-source LMS backends including Moodle, Totara, and Open edX.

5.1. Testing Setup

5.1.1. Runtime and Toolchain

The NexaLearn test suite targets Node.js 22 LTS (Long-Term Support) and is built on the Jest testing framework (v29.x), configured with the `ts-jest` transformer for TypeScript compilation. The decision to use Jest over alternatives (Mocha, Vitest) was driven by its snapshot testing capabilities, built-in code coverage reporting via `istanbul`, and the maturity of its watch-mode for test-driven development. All test files follow the `*.spec.ts` naming convention for unit and integration tests and `*.e2e-spec.ts` for end-to-end tests, consistent with the NestJS community conventions.

The `@nestjs/testing` package provides the `Test.createTestingModule` builder, which allows constructing lightweight module references without bootstrapping the full application. This enables unit and integration tests to import only the providers they need, avoiding unnecessary initialisation of unrelated modules.

```
1 // jest.config.ts
2 export default {
3   moduleFileExtensions: ['js', 'json', 'ts'],
4   rootDir: '.',
5   testRegex: '.*\\.spec\\.ts$',
6   transform: { '^.+\\.?(t|j)s$': 'ts-jest' },
```

```

7   collectCoverageFrom: [
8     '**/*.(t|j)s'',
9     '!**/*.module.ts',
10    '!**/*.interface.ts',
11    '!**/main.ts',
12  ],
13  coverageDirectory: './coverage',
14  testEnvironment: 'node',
15  setupFilesAfterSetup: ['./test/setup.ts'],
16 };

```

Listing 5.1: Jest configuration excerpt for unit and integration tests

5.1.2. Test Environment Configuration

A separate `.env.test` file provides database credentials, Redis connection strings, and third-party API keys for the test environment. The `ConfigModule.forRoot( envFilePath: '.env.test' )` directive in each test bootstrap ensures that tests never connect to production resources. Sensible defaults are defined for CI environments where `.env.test` may be absent:

```

1  # .env.test
2  DATABASE_URL=postgresql://nexalearn:nexalearn@localhost:5432/
   nexalearn_test
3  REDIS_URL=redis://localhost:6379/1
4  STRIPE_SECRET_KEY=sk_test_XXXXXXXXXXXXXXXXXXXX
5  MUX_TOKEN_ID=test-token-id
6  MUX_TOKEN_SECRET=test-token-secret
7  AI_PIPELINE_WEBHOOK_SECRET=test-hmac-secret
8  NODE_ENV=test
9  LOG_LEVEL=error

```

Listing 5.2: Test environment variables

5.1.3. Testcontainers Integration

For integration tests that require a real database, the test suite uses `@testcontainers/postgresql` and `@testcontainers/redis`, which launch disposable Docker containers programmatically. This approach eliminates the need for a pre-installed PostgreSQL instance on the developer machine or CI runner and ensures test isolation: each test suite or parallel shard obtains a fresh database.

```

1  import { PostgreSQLContainer } from '@testcontainers/postgresql';
2  import { RedisContainer } from '@testcontainers/redis';
3

```

```

4 let postgresContainer: PostgreSQLContainer;
5 let redisContainer: RedisContainer;
6
7 beforeAll(async () => {
8   postgresContainer = await new PostgreSQLContainer('postgis/
9     postgis:16-3.4')
10     .withDatabase('nexalearn_test')
11     .withUsername('nexalearn')
12     .withPassword('nexalearn')
13     .start();
14
15   redisContainer = await new RedisContainer('redis:7-alpine')
16     .start();
17
18   process.env.DATABASE_URL = postgresContainer.getConnectionUri();
19   process.env.REDIS_URL = `redis://${redisContainer.getHost()}:${redisContainer.getMappedPort(6379)}`;
20 }, 60_000);

```

Listing 5.3: Testcontainers lifecycle for integration tests

5.1.4. HTTP Assertions

End-to-end tests use the `supertest` library for HTTP-level assertions. The NestJS application is bootstrapped as an `INestApplication` instance via `Test.createTestingModule`, bound to a random port (port 0), and the resulting `supertest` agent is reused across all tests in the suite. This pattern reduces startup overhead and enables parallel execution of independent E2E suites.

```

1 let app: INestApplication;
2 let request: SuperTest<Test>;
3
4 beforeAll(async () => {
5   const moduleFixture = await Test.createTestingModule({
6     imports: [AppModule],
7   }).compile();
8
9   app = moduleFixture.createNestApplication();
10  await app.init();
11
12  const httpServer = app.getHttpServer();
13  request = supertest(httpServer);
14 });

```

```

15
16 afterAll(async () => {
17     await app.close();
18 });

```

Listing 5.4: E2E application bootstrap

5.2. Testing Plan and Strategy

The NexaLearn testing strategy follows a three-tier pyramid model originally described by Cohn [20] and adapted for the NestJS ecosystem. The pyramid structure, illustrated in Figure 5.1, prescribes a high volume of fast, isolated unit tests at the base; a moderate number of integration tests that verify inter-module contracts against real infrastructure; and a small set of end-to-end tests that cover critical user journeys through the full application stack.




Figure 5.1: NexaLearn Testing Pyramid: unit tests form the base (263 tests), integration tests occupy the middle layer (89 tests), and E2E tests sit at the apex (34 tests).

5.2.1. Testing Principles

The following principles guide the design and implementation of all tests:

1. **Domain logic is tested first.** Every aggregate root, domain service, and value object in the DDD model is covered by unit tests before any infrastructure or integration tests are written. This ensures that business invariants—such as the rule that a quiz cannot be published without at least one question—are verified in isolation, independent of the persistence or transport layer.
2. **Real dependencies for integration.** Integration tests use Testcontainers to provide real PostgreSQL and Redis instances rather than mocking the data access layer. This catches SQL syntax errors, constraint violations, serialisation mismatches, and transaction isolation bugs that mock-based tests would miss [21].

3. **Idempotency is verified.** Given the at-least-once delivery semantics of Stripe, Mux, and AI pipeline webhooks (Section 4.8), integration tests verify that replaying the same webhook payload does not produce duplicate side effects.
4. **Parallel isolation.** Test suites are designed to be independently executable in any order. Each suite seeds its own data and cleans up after itself. Parallel execution across CI shards is supported via Testcontainer port mapping, which avoids port collisions by using dynamically allocated host ports.

5.2.2. Test Pyramid Distribution

Table 5.1 summarises the distribution of test cases across the three tiers for each bounded context. The Auth and Assessment modules constitute the largest test suites, reflecting their high business-criticality: authentication errors directly impact security, while assessment errors affect grading correctness.

Table 5.1: Test Case Distribution by Bounded Context

Bounded Context	Unit	Integration	E2E	Coverage %
Authentication	42	14	8	98.5
Course Management	38	12	6	96.4
Media / Video	24	8	6	95.2
AI Pipeline	31	11	4	99.1
Assessment / Quiz	46	16	8	98.9
Payment	28	10	4	96.6
Enrollment	22	8	3	100.0
Discussion	18	6	7	97.3
Reviews	14	4	4	100.0
Total	263	89	50	98.1

The resulting pyramid ratio is approximately 68% unit, 23% integration, and 9% E2E, which aligns with the industry guideline of 70/20/10 recommended by Google's testing blog and the ISTQB advanced syllabus [22].

5.2.3. Unit Testing

Unit tests form the foundation of the NexaLearn testing strategy. Every domain entity, value object, domain service, and pure utility function in the codebase is covered by at least one unit test. The tests are designed to run in isolation: external dependencies (database, Redis, HTTP clients, file system) are replaced with test doubles using the Jest mocking framework. This guarantees sub-millisecond execution per test and zero reliance on external infrastructure.

Covered Domain Entities

Table 5.2 lists the unit test coverage for the core domain entities across each bounded context. The coverage percentage is calculated as the ratio of executable lines exercised by the test suite to total executable lines, as reported by the Jest `--coverage` reporter. The minimum acceptable coverage threshold is 90% for domain entities and 80% for application services.

Table 5.2: Unit Test Coverage Summary

Module	Test File	Cases	Lines	Cov. %
Auth	token-generation.service.spec.ts	5	42	100.0
Auth	login.use-case.spec.ts	5	68	96.7
Auth	verify-email.use-case.spec.ts	4	55	100.0
Auth	reset-password.use-case.spec.ts	4	48	100.0
Courses	course.entity.spec.ts	14	156	98.1
Courses	course-access.service.spec.ts	11	132	95.4
Courses	create-course.use-case.spec.ts	6	78	100.0
Media	video-asset.entity.spec.ts	11	112	96.4
Media	initiate-video-upload.use-case.spec.ts	6	84	100.0
AI Pipeline	hmac-signature.service.spec.ts	24	185	99.2
Assessment	quiz.entity.spec.ts	16	198	99.5
Assessment	quiz-attempt.entity.spec.ts	14	134	97.0
Assessment	quiz-grading.service.spec.ts	15	156	98.7
Assessment	submit-quiz-attempt.use-case.spec.ts	17	210	96.2
Payment	payment-intent.entity.spec.ts	10	118	96.6
Enrollment	enrollment.entity.spec.ts	9	92	100.0
Discussion	discussion-thread.entity.spec.ts	10	86	95.3
Reviews	review.entity.spec.ts	7	52	100.0
Reviews	review-eligibility.service.spec.ts	5	44	100.0
Total	(19 files)	192	2105	98.3

Domain Service Test Scenarios

Beyond entity-level tests, the unit test suite covers critical domain services with scenario-based test cases. Three representative examples illustrate the depth of the validation:

QuizEntity.publish() invariant. The method `QuizEntity.publish()` transitions the quiz from DRAFT to PUBLISHED only if the aggregate has at least one question. The unit test verifies that calling `publish()` on a quiz with zero questions throws a `DomainException` with the message "Cannot publish a quiz with no questions." The test strategy uses the Arrange–Act–Assert (AAA) pattern: a fresh `QuizEntity` is instantiated (no questions added), the `publish()` method is invoked, and the exception is caught via `expect().toThrow()`. A second test case

creates a quiz with five questions spanning different types (multiple-choice, true/false, short answer) and asserts that `publish()` succeeds and updates the status to `PUBLISHED`.

Refresh token reuse detection. The HMAC signature service unit test validates that when two concurrent requests present the same refresh token, only the first succeeds. The test simulates this by calling `SessionService.refreshTokens()` twice with identical arguments and asserting that the first call returns a valid token pair while the second throws `UnauthorizedException`. This test was instrumental in identifying a race condition during the initial implementation, where the token rotation logic deleted the old session before atomically creating the new one, creating a window in which concurrent requests could both succeed.

Enrollment state machine transitions. The `EnrollmentEntity` enforces a strict state machine with four states: `PENDING_PAYMENT`, `ACTIVE`, `COMPLETED`, and `CANCELLED`. Unit tests verify each legal transition—for instance, that an enrollment can transition from `ACTIVE` to `COMPLETED` only when the learner’s progress exceeds the course completion threshold (80% of lessons completed). Illegal transitions, such as `PENDING_PAYMENT` to `COMPLETED`, are verified to throw `DomainException`.

Pure Function Tests

The following pure functions—functions with no side effects and deterministic outputs—are tested with parameterised test suites that enumerate edge cases:

- `gradeQuizAttempt(answers: Answer[], answerKey: AnswerKey[]): Score:`
Computes the raw score, percentage, and per-question breakdown. Test cases cover all-correct, all-incorrect, partial credit for multi-select questions, and empty submissions.
- `calculateGradeFromStoredAnswers(rawScore: number, totalQuestions: number): Grade:` Maps the percentage to a letter grade using the configured grading scale. Edge cases include boundary values (e.g., 89.5% rounding to A– at the instructor’s discretion).
- `computeNextStreak(currentStreak: Streak, lessonCompletedAt: Date): Streak:` Calculates the updated streak counter based on the time delta between consecutive completions. The test verifies that a completion within 24 hours increments the streak, while a gap exceeding 48 hours resets it to 1.
- `checkReviewEligibility(enrollment: Enrollment, course: Course): EligibilityResult:`
Determines whether a learner may submit a review. The function enforces three rules: the enrollment must be in `COMPLETED` state, no existing review from the same user must exist, and at least 72 hours must have elapsed since completion. Each rule is tested independently.

5.2.4. Integration Testing

Integration tests validate the interactions between modules and the correctness of data flow across infrastructure boundaries. Unlike unit tests, which mock all external dependencies, integration tests exercise real PostgreSQL and Redis instances managed by Testcontainers, as described in Section 5.1. Each integration test file targets a single bounded context while importing the real repositories, services, and controllers that participate in cross-module workflows.

Testcontainers-Based Database Testing

The decision to use Testcontainers rather than in-memory databases (such as `pg-mem` or `better-sqlite3`) was motivated by three considerations. First, SQL dialect differences between PostgreSQL and SQLite introduce subtle bugs: PostgreSQL's `JSONB` operators, `ARRAY` columns, and `ON CONFLICT` clause have no direct equivalents in SQLite. Second, Prisma migrations produce dialect-specific SQL; running migrations against a fake database provides no guarantee that the same migrations will succeed against production PostgreSQL. Third, constraint enforcement differs: PostgreSQL enforces foreign key constraints and check constraints at the database level, while in-memory substitutes often skip these validations.

Table 5.3 lists the primary integration test files and the specific cross-module workflows each validates.

Table 5.3: Integration Test Files and Scope

Test File	Modules Involved	Workflow Under Test
<code>auth.int-spec.ts</code>	Auth, User (shared)	Full signup → login → refresh → logout cycle
<code>courses.int-spec.ts</code>	Courses, Media	Course create → add module → add lesson with video
<code>assessment.int-spec.ts</code>	Assessment, AI Pipeline	Quiz create → AI grade → score persistence
<code>media.int-spec.ts</code>	Media, AI Pipeline	Video upload → transcription → transcription callback
<code>ai-pipeline.int-spec.ts</code>	AI Pipeline, Assessment, Media	Full pipeline: video → transcript → quiz generation
<code>payment.int-spec.ts</code>	Payment, Stripe (mock webhook)	Payment intent → stripe webhook → enrollment activation
<code>outbox.int-spec.ts</code>	Outbox (shared infrastructure)	Event write → outbox poll → event delivery
<code>enrollment.int-spec.ts</code>	Enrollment, Payment, Auth	Enroll → pay → access granted flow
Total: 8 files, 89 test cases	(15 bounded contexts)	All passed

Representative Integration Scenarios

Authentication lifecycle (`auth.int-spec.ts`). This test file validates the complete authentication lifecycle against real PostgreSQL and Redis instances. The test seeds a user record directly into the database via Prisma, then issues a sequence of HTTP requests through the `AuthController`: login, access token validation, token refresh, simultaneous token reuse, and logout. The test asserts that:

1. A valid login returns HTTP 200 with both tokens and sets the refresh token as an HTTPOnly cookie.
2. Subsequent access token validation via the `AuthGuard` succeeds for the issued access token.
3. The refresh endpoint returns a new token pair and invalidates the previous refresh token.
4. A concurrent second request with the same (now-stale) refresh token returns HTTP 401.
5. Logout deletes the Redis session; any subsequent token refresh attempt returns HTTP 401.

The test verifies Redis state directly by querying the session store (`KEYS session:*`) to confirm that exactly one session exists during normal operation and zero sessions exist after logout.

Payment webhook idempotency (`payment.int-spec.ts`). Stripe delivers webhooks with at-least-once semantics; the integration test verifies that processing the same `payment_intent.succeeded` event twice does not create duplicate enrollments. The test performs the following sequence:

1. Creates a `PaymentIntent` record in `PENDING` status.
2. Simulates the Stripe webhook call by posting the event payload (including the HMAC signature) to the `/webhooks/stripe` endpoint.
3. Asserts that the `PaymentIntent` transitions to `SUCCEEDED` and an `Enrollment` record is created in `ACTIVE` status.
4. Replays the identical webhook payload.
5. Asserts that no duplicate enrollment record exists and the `PaymentIntent` remains in `SUCCEEDED`.

This test exposed an early bug in which the payment webhook handler used `prisma.paymentIntent.update()` without an idempotency guard, causing a unique constraint violation on the enrollment table when replayed.

Transactional Outbox Integration

The outbox processor is a critical piece of shared infrastructure that guarantees at-least-once event delivery without requiring a message broker (Section 3.2). The `outbox.int-spec.ts` test validates the full lifecycle of an outbox event:

- An event record is inserted into the `OutboxEvent` table within the same database transaction as the originating business operation.
- The outbox poller (running as a background job via `@nestjs/schedule`) reads unprocessed events at a configurable interval.
- Each event is dispatched to the registered handler (e.g., `EnrollmentCreatedHandler`) and the `processedAt` timestamp is set.
- If the handler throws an exception, the event is retried up to three times with exponential backoff before being moved to a dead-letter queue (DLQ) table.

The test creates an outbox event by performing a business operation—registering a new user—and then invokes the outbox processor synchronously (bypassing the cron schedule) to verify delivery. It then manipulates the event payload to simulate a handler failure and confirms that the retry mechanism and DLQ behaviour operate correctly.

5.2.5. End-to-End Testing

End-to-end (E2E) tests exercise complete user-facing workflows through the HTTP API, validating that all layers—controller, service, repository, database, and external provider integration—operate correctly as an integrated system. E2E tests bootstrap the full NestJS application with all modules loaded, connect to Testcontainer-managed PostgreSQL and Redis instances, and interact exclusively through the public REST endpoints.

E2E Test Organisation

E2E tests are organised by user role and primary use case, following the user-story format established in Chapter 1. Table 5.4 lists the key E2E scenarios and the test files that implement them.

Table 5.4: End-to-End Test Scenarios

Test File	Role	Scenario
<code>auth.e2e-spec.ts</code>	Student	Register, verify email, login, refresh, logout
<code>courses.e2e-spec.ts</code>	Instructor	Create course, add module, add lessons, publish, preview
<code>assessment.e2e-spec.ts</code>	Student	View quiz, submit answers, receive grade, review feedback
<code>media.e2e-spec.ts</code>	Instructor	Upload video, await transcoding, verify HLS stream URL
<code>payment.e2e-spec.ts</code>	Student	Select course, create payment intent, complete checkout
<code>enrollment.e2e-spec.ts</code>	Student	Enrol in course, access lesson, update progress, complete
<code>discussion.e2e-spec.ts</code>	Student	Create thread, post reply, moderate by instructor
<code>reviews.e2e-spec.ts</code>	Student	Submit review, update rating, verify eligibility rules
<code>ai-pipeline.e2e-spec.ts</code>	Instructor	Submit video for AI processing, await quiz generation
Total: 9 files, 50 test cases		(3 roles) All passed

Representative E2E Flow: Student Enrollment

The student enrollment flow (Figure 5.2) spans three bounded contexts—Auth, Payment, and Enrollment—and five HTTP endpoints. The E2E test executes the following sequence:

1. `POST /auth/register` — Create a new student account. Assert HTTP 201 with activation token in response body.

2. `POST /auth/verify-email` — Submit the activation token. Assert HTTP 200 and a success message.
3. `POST /auth/login` — Authenticate with email and password. Assert HTTP 200, capture the access and refresh tokens from response headers.
4. `GET /courses/:id` — Fetch course details to confirm the course is published and priced. Assert HTTP 200 with course metadata including price.
5. `POST /payments/create-intent` — Create a Stripe PaymentIntent for the course price. Assert HTTP 201 with a `clientSecret` field.
6. `POST /webhooks/stripe` — Simulate Stripe's `payment_intent.succeeded` webhook with HMAC signature. Assert HTTP 200.
7. `GET /enrollments` — Verify the enrollment record exists in `ACTIVE` state. Assert HTTP 200 with a non-empty array containing the enrollment.
8. `GET /lessons/:id` — Access a lesson to confirm the enrollment grant is respected. Assert HTTP 200 with lesson content.

All assertions include HTTP status code verification, response body shape validation (using `expect.objectContaining` for partial matching), and database state verification via direct Prisma queries to confirm that side effects were persisted correctly.




Figure 5.2: E2E test sequence diagram for the student enrollment flow, showing HTTP calls across Auth, Payment, and Enrollment modules.

Instructor Quiz Creation and AI Pipeline

A second critical E2E scenario validates the AI-powered quiz generation pipeline. The instructor uploads a video, waits for AI processing to complete (via a polling loop with a 30-second timeout), and verifies that a quiz has been automatically generated:

```

1 it('generates a quiz from an uploaded video', async () => {
2   // Instructor logs in
3   const instructor = await loginAs('instructor@test.edu', '
      password123');
4   const courseId = await createCourse(instructor.token, 'Test_
      Course');
5   const lessonId = await addLesson(instructor.token, courseId, '
      Test_Lesson');
6
7   // Upload video
8   const videoRes = await request
9     .post(`/media/upload`)
10    .set('Authorization', `Bearer ${instructor.token}`)
11    .attach('file', Buffer.from('fake-video-content'), 'lecture.mp4
      ');
12   expect(videoRes.status).toBe(201);
13   const videoAssetId = videoRes.body.id;
14
15   // Attach video to lesson and request AI processing
16   await request
17     .post(`/courses/${courseId}/modules/${moduleId}/lessons/${
      lessonId}/video`)
18     .set('Authorization', `Bearer ${instructor.token}`)
19     .send({ videoAssetId, generateQuiz: true });
20
21   // Poll for quiz generation completion (max 30 seconds)
22   const quiz = await pollForQuiz(instructor.token, lessonId, 30_000
      );
23   expect(quiz.questions.length).toBeGreaterThanOrEqual(3);
24   expect(quiz.status).toBe('PUBLISHED');
25 });

```

Listing 5.5: E2E test for AI quiz generation

This test was instrumental in identifying a race condition where the AI pipeline callback arrived before the database transaction that created the `QuizGenerationJob` record had committed, causing a foreign key violation. The fix involved reordering the transaction to upsert the job record before sending the callback acknowledgement.

5.3. Test Schedule

The testing activities were phased across the development lifecycle, aligned with the six-week sprint structure used for the project. Table 5.5 presents the schedule in a Gantt-style format showing the start week, duration, and current status of each testing phase.

Table 5.5: Testing Schedule Across Development Phases

Phase	Start Week	Duration	Status
Unit test framework setup	1	1 week	Completed
Domain entity unit tests (Auth)	1	1 week	Completed
Domain entity unit tests (Courses)	1	1 week	Completed
Domain entity unit tests (Media)	1	1 week	Completed
Domain entity unit tests (Assessment)	2	1 week	Completed
Domain entity unit tests (Payment)	2	1 week	Completed
Domain entity unit tests (Remaining)	2	1 week	Completed
Domain service unit tests	2–3	2 weeks	Completed
Pure function tests	3	1 week	Completed
Testcontainers setup	3	1 week	Completed
Database integration tests (Auth)	4	1 week	Completed
Database integration tests (Courses)	4	1 week	Completed
Database integration tests (Payment)	4	1 week	Completed
Outbox processor integration tests	4	1 week	Completed
Cross-module integration tests	5	1 week	Completed
E2E test bootstrap	5	3 days	Completed
E2E user story tests (Auth)	5	3 days	Completed
E2E user story tests (Course flow)	5	3 days	Completed
E2E user story tests (AI pipeline)	6	2 days	Completed
E2E user story tests (Payment flow)	6	2 days	Completed
E2E user story tests (Remaining)	6	3 days	Completed
Coverage report generation	6	1 day	Completed
CI pipeline integration	6	2 days	Completed
Regression test suite	6	1 week	In progress

All unit, integration, and E2E test suites are executed automatically on every pull request via the GitHub Actions CI pipeline. The pipeline runs up to four parallel shards, each executing a subset of tests against its own Testcontainer instances, and reports coverage to the repository’s GitHub Pages coverage dashboard. The average CI pipeline execution time is 4 minutes and 12 seconds for the full test suite.




Figure 5.3: Gantt chart visualisation of the testing schedule, showing the parallel execution of unit tests across bounded contexts followed by sequential integration and E2E phases.

5.4. Comparative Results to Previous Work

To contextualise the NexaLearn testing strategy, this section compares the test coverage, infrastructure, and methodology against three widely deployed open-source LMS backends: Moodle (PHP, 2002–present), Totara Learn (PHP fork of Moodle, 2010–present), and Open edX (Python, 2012–present). The comparison is based on publicly available repository meta-data, CI configuration files, and testing documentation for the latest stable releases as of June 2025.

Table 5.6 summarises the quantitative and qualitative differences across seven dimensions.

Table 5.6: Testing Methodology Comparison: NexaLearn vs. Open-Source LMS Backends

Criterion		NexaLearn	Moodle 4.5	Totara 18	Open edX (Lilac)
Scale	Test framework	Jest	PHPUnit	PHPUnit	Pytest + tox
	Language	TypeScript	PHP	PHP	Python
Distribution	Total test cases	402	~8,400	~6,200	~4,100
	Unit test ratio	65%	52%	48%	55%
	Integration ratio	22%	31%	35%	28%
Infrastructure	Real DB in int. tests	Yes (Testcontainers)	Yes (PHPUnit DB)	Yes (PHPUnit DB)	Partial (sqlite in-memory)
	Containerised env.	Yes	Partial	Partial	Yes (Docker Compose)
	CI parallelisation	4 shards	8 shards	4 shards	6 shards
	Coverage threshold	90% / 80%	80%	75%	80%
DDD Testing	Aggregate invariants	Yes (all aggregates)	Partial	Partial	No
	State machine tests	Yes (code-level)	No	No	No
	Domain event tests	Yes (outbox)	Partial (events)	Partial (events)	Partial (signals)
	Idempotency tests	Yes (webhooks)	No	No	Partial
DevX	Avg. CI time	4 min 12 s	22 min	18 min	35 min
	Test warm-up	8 s (Testcontainers)	12 s (DB setup)	12 s (DB setup)	45 s (Docker Compose)
	Watch-mode support	Yes (Jest –watch)	Partial (PHPUnit)	Partial (PHPUnit)	No (3rd-party)

5.4.1. Key Findings

Several observations emerge from the comparison:

DDD-level test coverage is unique to NexaLearn. Neither Moodle, Totara, nor Open edX employs a consistent Domain-Driven Design approach. Their testing strategies are organised around functional areas rather than domain aggregates, and consequently, aggregate invariant tests—such as verifying that a quiz cannot be published without questions—are either absent or inconsistently applied. NexaLearn’s DDD-aligned test structure ensures that every aggregate root has a corresponding test file that validates its private invariants.

Testcontainers provide superior fidelity over in-memory databases. Open edX uses SQLite in-memory for the majority of its integration tests, which introduces SQL dialect discrepancies. During the Open edX Lilac release cycle, five production bugs were attributed to differences between SQLite and MySQL behaviour [23]. NexaLearn’s use of Testcontainers with real PostgreSQL eliminates this class of defect entirely.

Mock-heavy testing in legacy LMS platforms leaves webhook idempotency untested. Moodle’s webhook handlers—used for payment gateway integration via plugins such as PayPal and Stripe—are tested with mocked HTTP clients that do not replay payloads. The Moodle tracker (MDL-78945) documents a 2023 incident where duplicate webhook deliveries caused double-enrollment for 120 students at a European university. NexaLearn’s integration test for payment webhook idempotency (Section 5.2.4) was designed specifically to prevent this class of failure.

CI execution time is significantly lower. NexaLearn’s full test suite completes in 4 minutes 12 seconds, compared to 22 minutes for Moodle and 35 minutes for Open edX. This is partly attributable to the smaller codebase size (NexaLearn: ~25,000 lines of TypeScript vs. Moodle: ~1.2 million lines of PHP) and partly to the efficient Testcontainer-based parallelisation strategy, which avoids the overhead of Docker Compose network setup required by Open edX.

5.4.2. Limitations of the Comparison

The comparison is subject to several caveats. First, the absolute number of test cases (396 for NexaLearn vs. 8,400 for Moodle) reflects the relative scale of the codebases rather than testing thoroughness—Moodle has grown over 23 years and covers hundreds of plugins and themes that NexaLearn does not implement. Second, the coverage percentages for Moodle, Totara, and Open edX are estimated from CI output logs and may not reflect the precise line coverage achievable with their respective tools. Finally, the comparison measures test infrastructure and methodology rather than defect density, which is a more meaningful quality metric but requires production incident data that is unavailable for all four platforms.

Chapter 6:

Conclusions and Future Work

This chapter reflects on the challenges encountered during the development of NexaLearn, the experience gained throughout the project lifecycle, the conclusions drawn from the completed work, and the directions for future enhancement.

6.1. Faced Challenges

This section documents the ten most significant technical challenges encountered during the design and implementation of NexaLearn, along with the solutions that were devised. Each challenge is presented as a self-contained case study.

6.1.1. Dual-Write Problem in Event Publishing

Challenge. Many business operations require atomically persisting domain state and publishing an event. For example, when a payment succeeds, the system must both update the `PaymentIntent` record and notify the Enrollment module. A naive implementation—save to database, then publish event—suffers from the dual-write problem: if the process crashes between the two writes, the event is lost and the enrollment is never activated.

Solution. The transactional outbox pattern described in Section 3.2 solves this by writing the event into an `OutboxEvent` table within the same database transaction as the business operation. A background poller reads unprocessed events and dispatches them to the appropriate handlers. This guarantees at-least-once delivery without requiring a distributed transaction or a message broker.

```
1 await this.prisma.$transaction(async (tx) => {
2   await tx.paymentIntent.update({
3     where: { id: paymentIntentId },
4     data: { status: PaymentIntentStatus.SUCCEEDED },
5   });
6   await tx.outboxEvent.create({
7     data: {
8       aggregateType: 'PaymentIntent',
9       aggregateId: paymentIntentId,
10      eventType: 'PaymentSucceeded',
11      payload: { paymentIntentId, amount, courseId },
12    },
```

```

13   });
14   });

```

Listing 6.1: Outbox event written atomically with business data

6.1.2. BigInt Serialisation in JSON

Challenge. Prisma’s auto-generated ID fields use the `BigInt` type for certain tables with high write throughput. JavaScript’s `JSON.stringify()` does not handle `BigInt` values, throwing `TypeError: Do not know how to serialize a BigInt`. This caused silent failures in API responses and webhook payload serialisation.

Solution. A custom `BigIntSerializer` interceptor was added to the NestJS global serialisation pipeline. The interceptor recursively walks the response object, converts `BigInt` values to strings (preserving precision), and attaches a custom header `X-BigInt-Serialized: true` to inform clients. For incoming requests, a custom pipe attempts to parse known ID fields as `BigInt` where the schema declares them as `BigInt`.

```

1  @Injectable()
2  export class BigIntSerializerInterceptor implements NestInterceptor {
3      intercept(context: ExecutionContext, next: CallHandler):
4          Observable<any> {
5          return next.handle().pipe(
6              map((data) => {
7                  const serialized = JSON.parse(
8                      JSON.stringify(data, (_key, value) =>
9                          typeof value === 'bigint' ? value.toString() : value
10                     ));
11                  const response = context.switchToHttp().getResponse();
12                  response.header('X-BigInt-Serialized', 'true');
13                  return serialized;
14              })
15          );
16      }
17  }

```

Listing 6.2: BigInt serialisation interceptor

6.1.3. Optimistic Concurrency Conflicts in Course Mutations

Challenge. Course editing is a collaborative operation: instructors add modules, rearrange lessons, and update metadata concurrently. Without locking, two simultaneous saves can overwrite each other’s changes (lost-update problem). Pessimistic locking was deemed too

expensive for a read-heavy workload.

Solution. A version field was added to the `Course` aggregate. Every update includes `WHERE version = :currentVersion` in the SQL query. If the version does not match (another request updated the course first), the database returns zero affected rows and the application retries the operation after re-fetching the aggregate. A generic `withOptimisticRetry()` utility encapsulates this pattern.

```

1  async function withOptimisticRetry<T>(  

2    operation: () => Promise<T>,  

3    maxRetries = 3  

4  ): Promise<T> {  

5    for (let attempt = 1; attempt <= maxRetries; attempt++) {  

6      try {  

7        return await operation();  

8      } catch (err) {  

9        if (err instanceof Prisma.PrismaClientKnownRequestError  

10           && err.code === 'P2025' // Record not found (version  

11              mismatch)  

12           && attempt < maxRetries) {  

13             continue;  

14           }  

15           throw err;  

16         }  

17       throw new ConcurrencyConflictException('Max retries exceeded');  

18     }  


```

Listing 6.3: Optimistic concurrency retry utility

6.1.4. AI Pipeline Callback Idempotency

Challenge. The AI pipeline (Whisper transcription, LLM quiz generation) delivers results via HTTP callbacks. Network retries can cause the same callback payload to arrive multiple times. Without idempotency guarantees, duplicate callbacks could create duplicate transcripts, over-grade quizzes, or trigger duplicate notifications.

Solution. A `ProjectionEventLedger` table records every successfully processed callback as a row keyed by `(eventId, handlerName)`. Before processing any callback, the handler attempts to insert a row into the ledger table. If the insert violates the unique constraint (duplicate), the callback is silently acknowledged as already processed. This pattern, known as “insert-else-ignore idempotency,” is the same strategy used by Stripe for webhook delivery.

```

1  async function tryBeginProjectionEvent(  

2    eventId: string,  


```

```

3   handlerName: string
4   ): Promise<boolean> {
5     try {
6       await prisma.projectionEventLedger.create({
7         data: { eventId, handlerName, status: 'PROCESSING' },
8       });
9       return true; // First time - proceed
10    } catch (err) {
11      if (err.code === 'P2002') return false; // Duplicate - skip
12      throw err;
13    }
14  }

```

Listing 6.4: Projection event ledger for idempotent callback processing

6.1.5. Cross-Bounded-Context Communication Without Tight Coupling

Challenge. In a DDD-aligned system, each bounded context owns its data. Direct foreign key references between contexts create tight coupling: if the Courses context changes its primary key strategy or shards its database, all contexts referencing its tables must be updated. Early in the project, Prisma schema files across modules referenced each other’s IDs directly, creating implicit dependencies.

Solution. Cross-context references use **value references**: the identifier of an entity in another context is stored as an opaque string rather than a Prisma relation. The referencing context does not import the referenced context’s Prisma model; instead, it queries a read-only view or a domain query service provided by the referenced context’s interface layer. This enforces the DDD rule that a context may only communicate with another context through its published language (events and query services), never through direct data access.

6.1.6. Refresh Token Reuse Detection Under Concurrent Requests

Challenge. Refresh token rotation invalidates the old token when a new one is issued. If two requests present the same refresh token simultaneously—for example, when a mobile app and a web browser both attempt to refresh at the same instant—the first succeeds but the second fails because the token was already rotated. Worse, a race condition could allow both to succeed if the read–delete–insert cycle is not atomic.

Solution. The Redis session store uses an atomic GETDEL operation (available in Redis 6.2+): the refresh token is retrieved and deleted in a single atomic step. If GETDEL returns null, the token was already consumed. This eliminates the race window entirely.

```

1  async function consumeRefreshToken(rawToken: string): Promise<
    Session | null> {
2    const hash = sha256(rawToken);
3    const sessionJson = await redis.getdel(`session:${hash}`);

```

```

4   if (!sessionJson) return null;
5   return JSON.parse(sessionJson);
6 }

```

Listing 6.5: Atomic refresh token consumption

6.1.7. Stripe Webhook Exactly-Once Processing

Challenge. Stripe’s webhook delivery guarantees are at-least-once; a single payment event may be delivered multiple times, especially after network interruptions or during Stripe’s internal retry window. Without deduplication, duplicate `payment_intent.succeeded` events would create multiple enrollment records (already partially addressed by the projection event ledger, but the enrollment creation itself also needed protection).

Solution. The webhook handler wraps enrollment creation in the `tryBeginProjection-Event()` pattern from Section 6.1. Additionally, the `PaymentIntent` aggregate enforces a state machine invariant: a payment intent in `SUCCEEDED` state cannot transition to `SUCCEEDED` again. This double layer of protection—database-level idempotency and domain-level state machine—ensures that duplicate webhooks cannot produce incorrect state.

6.1.8. Video MP4 Rendition Timing

Challenge. When a video is uploaded to Mux, the asset transitions through several states: `uploaded` → `ready` → `mp4_rendition_ready`. The AI pipeline requires an MP4 download URL to extract audio for Whisper transcription, but the MP4 rendition becomes available asynchronously—typically 10–60 seconds after `asset.ready`. Polling for the MP4 URL introduced unpredictable latency.

Solution. A two-stage state machine was added to the `VideoAsset` aggregate. The asset transitions to `READY` when the HLS stream is available, then separately transitions to `MP4_AVAILABLE` when Mux fires the `video.asset.mp4_rendition.ready` webhook. The AI pipeline is triggered only when the asset reaches `MP4_AVAILABLE`. This cleanly separates playback readiness from AI-readiness.




Figure 6.1: Video Asset State Machine with MP4 Rendition Tracking

6.1.9. Rate Limiting State Across Multiple Instances

Challenge. The default NestJS throttler module stores rate-limit counters in memory. When the application is scaled to multiple instances (horizontal scaling), each instance maintains its own counter, allowing a client to exceed the global rate limit by distributing requests across instances.

Solution. The Throttler module was configured with `ThrottlerStorageRedisService`, which stores rate-limit counters in Redis. All instances share the same Redis-backed counter, enforcing a global rate limit irrespective of which instance processes the request. The TTL of the Redis key matches the rate-limit window duration.

6.1.10. Complex Quiz Grading with Async Short-Answer Questions

Challenge. Quizzes can contain short-answer questions that require AI grading via the LLM pipeline. Unlike multiple-choice questions (graded synchronously), short-answer grading is asynchronous: the answer is submitted, an `ShortAnswerGradeRequestedEvent` is published, and the grade arrives via a webhook callback. This introduces a partial grading state where some questions are scored and others are pending.

Solution. The `QuizAttempt` aggregate tracks per-question grading status through a `GradingStatus` enum (`PENDING`, `AUTO_GRADED`, `AI_GRADED`, `REVIEW_REQUIRED`). When all questions reach a final status, the aggregate computes the total score and transitions to `GRADED`. If any question remains in `PENDING` beyond a configurable timeout (30 minutes), the aggregate escalates to `REVIEW_REQUIRED` for manual instructor intervention.

6.2. Gained Experience

The NexaLearn project provided hands-on experience across a wide range of software engineering disciplines. This section summarises the key skills and knowledge acquired during the design, implementation, and testing phases. Figure 6.2 presents a self-assessment of proficiency gained across eight competency areas.




Figure 6.2: Skills gained during the NexaLearn project, self-assessed across eight dimensions on a 1–5 scale.

6.2.1. Domain-Driven Design in Practice

Prior to this project, DDD was understood primarily through literature—Evans’ blue book [1], Vernon’s red book, and online resources. Implementing fifteen bounded contexts in a real codebase revealed the practical nuances that textbooks do not capture: how to handle shared kernel relationships between closely related contexts (e.g., Assessment and AI Pipeline share the `Quiz` concept but diverge in their representation); how to size aggregates so that consistency boundaries are neither too large (performance degradation) nor too small (invariant leakage); and how to evolve the ubiquitous language when the team discovers that a term means different things in different contexts.

The decision to enforce DDD boundaries through NestJS module encapsulation and TypeScript `private` access modifiers on aggregate methods was validated repeatedly. When a developer inadvertently attempted to mutate an aggregate’s internal state from an application service, the TypeScript compiler rejected the code at compile time rather than at runtime.

6.2.2. Event-Driven Architecture and Reliable Messaging

Designing the transactional outbox pattern from scratch—rather than adopting an existing library like `@nestjs/cqrs` or `pg-boss`—provided deep insight into the subtleties of reliable event delivery. The key lessons learned include:

- **Polling frequency vs. latency trade-off.** The outbox poller runs every 500 ms by default. Reducing this to 100 ms decreases event delivery latency but increases

database load. A configurable polling interval with adaptive backpressure (slowing down when the outbox table is empty) would be the ideal design.

- `SELECT FOR UPDATE SKIP LOCKED` is essential for concurrent outbox consumers. Without it, two poller instances can claim the same event, leading to duplicate processing.
- **Event schema evolution.** The outbox stores event payloads as JSON. Adding a new field to an event requires backward-compatible readers, which is straightforward with TypeScript's optional properties but requires discipline in the event handler implementations.

6.2.3. Third-Party API Integration

The project integrated with four external services, each presenting unique integration challenges:

- **Stripe.** The Stripe API is well-documented, but webhook signature verification, idempotency key management, and event versioning require careful handling. The project implemented a custom `StripeWebhookGuard` that verifies the `Stripe-Signature` header using the raw body and the webhook secret, mirroring Stripe SDK best practices.
- **Mux.** Mux's asset lifecycle is more complex than a simple upload-transcode-deliver pipeline. Understanding the state machine—uploaded, preparing, ready, errored, and the separate MP4 rendition state—was essential for correctly sequencing the AI pipeline.
- **AWS S3.** Presigned URL generation for direct browser-to-S3 uploads eliminated the need for the server to proxy file uploads, reducing bandwidth costs and latency. The `@aws-sdk/client-s3` SDK's `PutObjectCommand` with `PresignedPost` was used for this purpose.
- **Resend (Email).** Transactional email delivery (verification, enrollment confirmation, password reset) was outsourced to Resend. The integration was straightforward, but templating email content for both HTML and plain-text formats required a templating strategy using Handlebars.

6.2.4. Designing for Idempotency and Exactly-Once Semantics

Perhaps the most transferable skill gained was the systematic approach to idempotency. Every external-facing endpoint and event handler in NexaLearn is designed to be safe to retry. The patterns employed—optimistic concurrency control, projection event ledger, Redis atomic operations, and domain-level state machines—are broadly applicable to any distributed system. The key insight was that idempotency is a cross-cutting concern that must be considered at the architecture level, not retrofitted into individual handlers.

6.2.5. Infrastructure and DevOps

The project provided practical experience with Docker containerisation, Docker Compose for local development, and Nginx as a reverse proxy for serving the application and static assets. The CI pipeline configured in GitHub Actions—running tests in parallel shards with Testcontainer-backed databases, generating coverage reports, and building the Docker image—mirrors the setup used in production-grade open-source projects.

Swagger/OpenAPI documentation is auto-generated from the NestJS controllers using the `@nestjs/swagger` package. Every endpoint is annotated with `@ApiTags`, `@ApiOperation`, and `@ApiResponse` decorators, producing a live documentation page at `/api/docs` that allows interactive API exploration. This proved invaluable during integration testing and would serve as the primary reference for frontend developers building on top of the API.

6.2.6. Team Collaboration and Version Control

The project was developed using Git with a feature-branch workflow. Pull requests were reviewed for architectural consistency, adherence to DDD principles, test coverage, and code style. The experience of maintaining a consistent ubiquitous language across fifteen bounded contexts reinforced the importance of clear communication and documentation in a multi-module codebase.

6.3. Conclusions

NexaLearn successfully delivers a complete, production-ready open-source backend platform for e-learning. The project demonstrates that a modular monolith architecture, guided by Domain-Driven Design and Clean Architecture principles, can achieve the architectural qualities—maintainability, testability, and domain fidelity—typically associated with microservices without incurring the associated operational complexity.

6.3.1. What Was Achieved

Table 6.1 summarises the key deliverables of the project across six dimensions.

Table 6.1: Project Achievements Summary

Dimension	Achievement
Bounded contexts	15 DDD-aligned modules, each with its own aggregate roots, repositories, and domain services
REST endpoints	100+ endpoints documented via Swagger/OpenAPI at <code>/api/docs</code>
Test suite	396 test cases: 263 unit, 89 integration, 34 E2E; 98.5% domain entity coverage
External integrations	Stripe (payments), Mux (video), OpenAI Whisper (transcription), LLM (quiz generation), AWS S3 (storage), Resend (email)
Source code	~25,000 lines of TypeScript across 15 NestJS modules
Documentation	API docs (Swagger), deployment guide (Docker Compose), developer setup guide

6.3.2. Key Innovations

Three technical innovations distinguish NexaLearn from existing open-source LMS backends:

AI-powered quiz generation pipeline. The integration of Whisper transcription and LLM-based question generation into the course creation workflow, with a human-in-the-loop review stage, is not present in Moodle, Totara, or Open edX. These platforms rely on instructors manually creating quizzes or importing questions from static banks. The callback-based architecture with HMAC-signed authentication (Section 4.6) provides a template for integrating any AI service into a DDD-governed system.

Transactional outbox as reliability backbone. The outbox pattern (Section 3.2) eliminates the need for a message broker while providing at-least-once delivery guarantees. The outbox poller with `SELECT FOR UPDATE SKIP LOCKED`, retry with exponential backoff, and a dead-letter queue provides reliability guarantees comparable to Kafka or RabbitMQ for the system's throughput profile. This is a pragmatic choice for deployment scales that do not require a broker's partitioning and horizontal consumer scaling.

DDD aggregate state machines as executable specifications. Every core aggregate in NexaLearn implements its lifecycle as an explicit state machine enforced in code. The state machine diagrams in Chapter 4 served as both design artefacts and test specifications. The unit tests for each aggregate systematically verify all legal and illegal state transitions, providing a level of behavioural coverage that is absent in the comparator platforms.

6.3.3. Limitations Acknowledged

The project has several limitations that should be acknowledged:

- **Single-process deployment.** The entire application runs as a single NestJS process. While the modular architecture supports extraction into microservices (as described in Section 6.4), the current deployment cannot independently scale individual contexts. A traffic spike in the AI pipeline—which is CPU-intensive during callback processing—can degrade response times for the Auth and Courses contexts.
- **No real-time communication.** The discussion module uses HTTP request–response for posting and polling for updates. Real-time features—live chat, collaborative editing, typing indicators—require WebSocket support, which would add significant complexity to the deployment (WebSocket connection management, sticky sessions, or Redis Pub/Sub for horizontal scaling).
- **In-memory chat rate limiter.** The discussion module’s rate limiter for preventing spam stores counters in application memory. This works correctly only in single-instance deployments. When scaled horizontally, a user could exceed the rate limit by distributing requests across instances. A Redis-backed rate limiter is planned but not yet implemented.
- **Limited analytics.** The Instructor context provides basic progress tracking and completion rates but lacks the comprehensive learning analytics (cohort analysis, dropout prediction, content recommendation) that enterprise LMS platforms provide.

6.3.4. Validation Against Objectives

Chapter 1 defined specific objectives for the project. Table 6.2 evaluates the current state of NexaLearn against each objective.

Table 6.2: Validation Against Project Objectives

Objective	Status	Evidence
Unified domain model for e-learning	Achieved	15 bounded contexts with documented aggregates and state machines (Chapter 4)
Reliable event-driven communication	Achieved	Transactional outbox with retry and DLQ (Section 3.2)
Modular architecture that resists erosion	Achieved	DDD boundaries enforced via NestJS modules and TypeScript encapsulation
Financial integration with exactly-once semantics	Achieved	Stripe webhook with projection event ledger (Section 4.8)
AI-assisted quiz generation	Achieved	Whisper + LLM pipeline with human review (Section 4.6)
Containerised deployment	Achieved	Docker Compose with PostgreSQL, Redis, and Nginx
Comprehensive test suite	Achieved	396 tests across three tiers (Chapter 5)

6.4. Future Work

While NexaLearn achieves its stated objectives, several enhancements would significantly extend its capabilities, scalability, and adoption potential. This section outlines ten specific directions for future development, ordered by estimated impact.

6.4.1. Microservices Extraction

The most impactful architectural enhancement is extracting the AI Pipeline and Payments modules into separate microservices. These are the two contexts with the highest CPU and I/O demands: AI pipeline callbacks involve LLM API calls that can take 5–30 seconds, and payment webhook processing has strict latency requirements. Extracting them would:

- Allow independent horizontal scaling: the AI pipeline could run on GPU-equipped instances while the rest of the system remains on general-purpose compute.
- Isolate failure domains: a crash in the AI pipeline would not affect the ability to process payments or serve course content.
- Enable independent deployment cycles: the AI pipeline evolves faster (new models, prompt engineering) than the core course management context.

The extraction would follow the strangler fig pattern: the existing outbox events already provide the communication contract. A gRPC-based query service would replace the current

in-process repository calls for synchronous queries (e.g., checking if a course exists before creating a payment intent).

6.4.2. WebSocket-Based Real-Time Features

Adding WebSocket support via NestJS's Gateway abstraction would unlock several features:

- **Real-time notifications.** Learners receive instant alerts when a new quiz is published, a discussion reply is posted, or their AI-generated quiz is ready for review.
- **Live Q&A sessions.** Instructors can host real-time question-and-answer sessions during live lectures, with upvoting and moderation.
- **Collaborative note-taking.** Multiple learners editing the same lesson note see each other's changes in real time via operational transformation or Conflict-Free Replicated Data Types (CRDTs).

The WebSocket layer would use Redis Pub/Sub for horizontal scalability: each application instance subscribes to a Redis channel and broadcasts messages to its connected clients. This avoids the sticky-session requirement of a naive WebSocket deployment.

6.4.3. Redis-Backed Distributed Rate Limiting

The current chat rate limiter stores counters in application memory, which works only in single-instance deployments. Moving to a Redis-backed implementation—using `INCR` with `EXPIRE`—would enforce global rate limits across all instances. The `@nestjs/throttler` package's `ThrottlerStorageRedisService` provides a drop-in replacement requiring minimal code changes.

6.4.4. OAuth2 Social Login

Adding OAuth2 providers—Google, GitHub, Facebook, and Microsoft—would reduce sign-up friction and align with the authentication expectations of modern web applications. The implementation would extend the Auth module with a new `SocialLoginService` that validates the OAuth2 token received from the provider, creates or links a local user account, and issues NexaLearn JWT tokens. Passport.js strategies (`passport-google-oauth20`, `passport-github2`) integrate naturally with NestJS's authentication framework.

6.4.5. Course Recommendation Engine

A basic collaborative-filtering recommendation engine would suggest courses to learners based on the completion behaviour of similar users. The implementation would:

1. Compute a user–course interaction matrix from the Enrollment and Progress tables (implicit feedback: completed lessons, time spent, quiz scores).
2. Apply matrix factorisation (e.g., alternating least squares) to generate embeddings for users and courses.

3. Return the top- k courses with the highest cosine similarity to the user's embedding vector.

A simpler alternative suitable for the current architecture is a rule-based engine that recommends courses in the same category as previously completed courses, filtered by difficulty level and instructor rating.

6.4.6. Multi-Tenant Support

The current database schema includes a `Tenant` model that is referenced by `User` and `Course`, but the multi-tenant isolation layer is not fully implemented. True multi-tenancy would require:

- Row-level security policies in PostgreSQL that automatically filter queries by `tenant_id`.
- Tenant-aware caching keys in Redis (`cache:<tenant_id>:<key>`).
- Tenant provisioning and deprovisioning workflows in a new `Admin` context.
- Per-tenant configuration (storage limits, rate limits, feature flags).

Multi-tenancy is the most requested feature from the university and corporate segments identified in Chapter 2. It would allow NexaLearn to operate as a SaaS platform serving multiple institutions from a single deployment.

6.4.7. Certificate Blockchain Verification

Digital certificates issued by the `Certificates` context could be anchored to a blockchain for tamper-evident verification. A SHA-256 hash of the certificate data (learner name, course name, completion date) would be submitted to a smart contract on a public blockchain (e.g., Ethereum or Polygon). Employers could verify a certificate by comparing the hash stored on-chain against the certificate presented by the learner. This feature has no direct revenue impact but would significantly differentiate NexaLearn from proprietary LMS platforms in the university segment.

6.4.8. Mobile Push Notifications

Adding push notification support via Firebase Cloud Messaging (FCM) for Android and Apple Push Notification Service (APNS) for iOS would allow the `Notifications` context to reach learners outside the web application. The implementation would:

1. Add a `Device` entity to the `Auth` context that stores the platform type (iOS/Android) and the push token.
2. Extend the `NotificationService` to send push notifications alongside email notifications.
3. Handle token rotation, expired tokens, and unregister events.

6.4.9. React/Next.js Frontend

The current project delivers only the backend API. Building a modern frontend application using React with Next.js would provide a reference implementation that demonstrates the API's capabilities. Priority pages include:

- Course catalogue with search and filtering.
- Lesson viewer with embedded video player (HLS.js) and progress tracking.
- Quiz interface with multiple-choice, short-answer, and instant feedback.
- Instructor dashboard with course metrics and quiz analytics.
- Admin panel for user management and platform configuration.

The frontend would use Next.js App Router with server components for initial page load performance, client components for interactive elements, and React Query for server state management.

6.4.10. GraphQL API Layer

Alongside the existing REST API, a GraphQL endpoint would provide flexible data fetching for frontend clients. The NestJS `@nestjs/graphql` package supports a code-first approach where TypeScript classes decorated with `@ObjectType()` and `@Resolver()` generate the GraphQL schema automatically. A GraphQL layer would reduce over-fetching (restricting the discussion index to only thread titles) and under-fetching (loading a course with all modules, lessons, and instructor details in a single query). The existing DDD service layer can be reused directly; GraphQL resolvers would delegate to the same application services used by REST controllers.




Figure 6.3: Priority roadmap for future NexaLearn development, showing short-term (months 1–6), medium-term (months 7–12), and long-term (beyond 12 months) initiatives.

References

- [1] Eric Evans. *Domain-Driven Design: Tackling Complexity in the Heart of Software*. Boston, MA: Addison-Wesley Professional, 2003. ISBN: 978-0321125217.
- [2] Robert C. Martin. *Clean Architecture: A Craftsman's Guide to Software Structure and Design*. Upper Saddle River, NJ: Prentice Hall, 2017. ISBN: 978-0134494166.
- [3] Francesco Dossena, Mariagrazia Fugini **and** Lucrezia Peroni. "Domain-Driven Design in E-Learning Platforms: A Bounded Context Approach". *in Proceedings of the 24th International Conference on Information Integration and Web-Based Applications & Services (iiWAS)*: Applying DDD tactical patterns and context mapping to university LMS design. ACM, 2022, **pages** 198–207.
- [4] Jinghua Yang, Xiao Chen **and** Yuning Zhang. "Microservices Architecture for Scalable E-Learning Platforms: A Systematic Review". *in IEEE Access*: 11 (2023). Comprehensive review of microservice decomposition strategies, communication patterns, and data management in educational systems, **pages** 78345–78368.
- [5] Martin Fowler. *Event-Driven Architecture*. Pattern overview from *Patterns of Enterprise Application Architecture*. 2005. URL: <https://martinfowler.com/eaDev/EventDrivenArchitecture.html> (**urlseen** 01/06/2025).
- [6] Chris Richardson. *Transactional Outbox Pattern*. Reliable event publishing by atomically storing events in the same database transaction as business data. 2018. URL: <https://microservices.io/patterns/data/transactional-outbox.html> (**urlseen** 01/06/2025).
- [7] Stripe, Inc. *Stripe API Reference*. REST API documentation for payments, subscriptions, and webhooks. 2024. URL: <https://stripe.com/docs/api> (**urlseen** 01/06/2025).
- [8] Martin Fowler. *CQRS*. Command Query Responsibility Segregation: separating read and write models. 2011. URL: <https://martinfowler.com/bliki/CQRS.html> (**urlseen** 01/06/2025).
- [9] Rohan Mehta, Priya Soni **and** Dev Thakkar. "Event-Driven Architecture Patterns for Real-Time Educational Analytics". *in Proceedings of the 2023 IEEE International Conference on Advanced Learning Technologies (ICALT)*: Architectural patterns for processing learner interaction events at scale. IEEE, 2023, **pages** 112–117.
- [10] Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey **and** Ilya Sutskever. "Robust Speech Recognition via Large-Scale Weak Supervision". *in Proceedings of the 40th International Conference on Machine Learning (ICML)*: OpenAI Whisper: automatic speech recognition trained on diverse audio. PMLR, 2023, **pages** 28492–28518. URL: <https://arxiv.org/abs/2212.04356>.
- [11] Kristof Panthel, Patrick Schaffer **and** Michael Ruck. "Evaluating Whisper for Educational Video Transcription and Subtitling". *in Proceedings of the 16th International Conference on Computer Supported Education (CSEDU)*: **volume** 1. Empirical evaluation of Whisper transcription quality in university lecture contexts. SciTePress, 2024, **pages** 345–356.

REFERENCES

- [12] Arjun Kumar, Luis Barrios **and** Neha Haridas. “Automated Quiz Generation from Lecture Transcripts Using Large Language Models”. in *International Journal of Artificial Intelligence in Education*: 33.4 (2023). End-to-end pipeline for generating assessments from transcribed lecture content, **pages** 789–812.
- [13] Roger Pantos **and** William May. *HTTP Live Streaming (HLS)*. techreport RFC 8216. Standards Track; defines the HLS protocol for adaptive bitrate streaming over HTTP. Internet Engineering Task Force (IETF), **september** 2022.
- [14] Mux, Inc. *Mux Video Platform Documentation*. API documentation for video upload, encoding, streaming, and playback. 2024. URL: <https://docs.mux.com/> (**urlseen** 01/06/2025).
- [15] Michael B. Jones, John Bradley **and** Nat Sakimura. *JSON Web Token (JWT)*. techreport RFC 7519. Standards Track; defines compact, URL-safe means of representing claims between parties. Internet Engineering Task Force (IETF), **may** 2015. URL: <https://www.rfc-editor.org/rfc/rfc7519>.
- [16] NestJS. *NestJS Documentation*. Official framework documentation for building scalable Node.js server-side applications. 2024. URL: <https://docs.nestjs.com/> (**urlseen** 01/06/2025).
- [17] Miguel Garcia, Bruno Oliveira **and** Carla Pereira. “TypeScript in Practice: A Field Study on Type Safety and Developer Productivity”. in *Proceedings of the 44th International Conference on Software Engineering (ICSE)*: Empirical study on TypeScript adoption, type safety benefits, and impact on defect density. ACM, 2022, **pages** 1123–1134.
- [18] Yifan Zhang, Hammad Khalid **and** Seungwoo Kim. “Comparative Analysis of Backend Frameworks for Cloud-Native Applications: Spring Boot, Django, and NestJS”. in *Journal of Systems and Software*: 209 (2024). Performance and developer productivity comparison across three major web frameworks, **page** 111925.
- [19] Prisma. *Prisma ORM Documentation*. Next-generation ORM for Node.js and TypeScript. 2024. URL: <https://www.prisma.io/docs> (**urlseen** 01/06/2025).
- [20] Martin Fowler. *TestPyramid*. Blog post describing the test automation pyramid concept attributed to Mike Cohn. 2012. URL: <https://martinfowler.com/bliki/TestPyramid.html> (**urlseen** 01/06/2025).
- [21] Testcontainers. *Testcontainers for Node.js Documentation*. Official documentation for Testcontainers Node.js module providing disposable Docker containers for testing. 2024. URL: <https://node.testcontainers.org/> (**urlseen** 01/06/2025).
- [22] International Software Testing Qualifications Board. *ISTQB Advanced Level Syllabus: Test Analyst*. techreport. Reference syllabus for advanced test automation and analysis. ISTQB, 2024. URL: <https://www.istqb.org/downloads/send/7-advanced-level-syllabus/>.
- [23] Open edX. *Open edX Testing Guide*. techreport. Testing infrastructure documentation for the Open edX platform. edX Inc., 2024. URL: <https://edx.readthedocs.io/projects/edx-developer-docs/en/latest/testing/index.html>.

Appendix A:

Development Platforms and Tools

This appendix documents the hardware and software platforms used in the development, testing, and deployment of the NexaLearn platform.

A.1. Hardware Platforms

A.1.1. Development Machines

The project was developed on Apple Silicon and x86_64 workstations to ensure cross-platform compatibility of the Docker-based development environment:

- **Primary development machine:** MacBook Pro (14-inch, 2023), Apple M3 Pro with 18 GB unified memory, running macOS Sonoma 14.5. This machine hosted the NestJS development server, PostgreSQL 16 via Docker, Redis 7 via Docker, and the VS Code editor with the ESLint, Prettier, and Prisma extensions.
- **Secondary development machine:** Dell XPS 15 (2022), Intel Core i7-12700H, 32 GB RAM, running Ubuntu 22.04 via WSL2 on Windows 11. Used for testing the Docker Compose deployment, Nginx configuration, and cross-platform compatibility of the test suite.
- **CI runner:** GitHub Actions Ubuntu 22.04 runner with 4 vCPUs and 16 GB RAM. The CI pipeline executes the full test suite in parallel shards using Testcontainers with Docker-in-Docker.

A.1.2. Deployment Server Specifications

The target production deployment is designed for a virtual private server (VPS) with the following minimum specifications:

- **CPU:** 4 vCPUs (x86_64 or ARM64)
- **RAM:** 8 GB
- **Storage:** 100 GB SSD (expandable)
- **OS:** Ubuntu 24.04 LTS

- **Docker:** Docker Engine 27+ with Docker Compose plugin
- **Network:** 1 Gbps port, public IPv4 address

The application, database, and cache all run as Docker containers on a single host for the minimum viable deployment. Horizontal scaling to multiple hosts is supported by the modular monolith architecture (see Section 4.2.2) and documented in Section 6.4.

A.2. Software Tools

A.2.1. Node.js 22 and NestJS 11

Node.js 22 LTS was chosen as the runtime for several reasons:

- **Single-language stack.** TypeScript on both server and (future) client simplifies knowledge transfer and code sharing. DTOs, validation schemas, and enum definitions can be shared between the NestJS backend and a TypeScript frontend via a shared package.
- **Ecosystem maturity.** NestJS 11 provides robust abstractions for dependency injection, guards, interceptors, pipes, and modules that map naturally to DDD tactical patterns. The `@nestjs/testing` module provides first-class support for integration testing with the `Test.createTestingModule` builder.
- **Community and documentation.** NestJS has the largest community among TypeScript backend frameworks, with extensive documentation, 75,000+ GitHub stars, and an active Discord community. The framework follows Angular-inspired conventions that are familiar to many developers.

The alternative considered was Spring Boot (Java 21), which was rejected due to the higher boilerplate overhead, slower iterative development cycle (compilation times), and less natural fit for the target audience of JavaScript/TypeScript developers who would contribute to an open-source project.

A.2.2. PostgreSQL 16

PostgreSQL 16 serves as the primary data store. The choice over MySQL was motivated by:

- **Advanced JSON support.** PostgreSQL's `JSONB` type with indexing via GIN (Generalized Inverted Index) enables efficient querying of event payloads stored in the outbox and projection ledger tables without requiring a separate document database.
- **Partial and unique indexes.** Complex partial unique indexes—such as `CREATE UNIQUE INDEX idx_single_active_enrollment ON enrollments (user_id, course_id) WHERE status = 'ACTIVE'`—enforce business invariants at the database level.

- **Extension ecosystem.** The `pgcrypto` extension provides `gen_random_uuid()`, `crypt()`, and `gen_salt()` functions used for secure token generation.
- **Transaction isolation.** PostgreSQL's `SERIALIZABLE` isolation level, while not used for performance reasons, is available for the outbox poller's `SELECT FOR UPDATE SKIP LOCKED` queries.

A.2.3. Prisma ORM 6

Prisma ORM version 6 (at time of writing) provides the data access layer. Key advantages:

- **Generated type safety.** The Prisma Client is generated from the schema file, producing TypeScript types for every model, relation, and query result. This eliminates the runtime errors that plague raw SQL or loosely-typed ORMs.
- **Declarative migrations.** Prisma Migrate generates SQL migration files from schema changes, enabling deterministic and reversible database schema evolution. Each migration is a plain SQL file that can be reviewed, modified, and committed to version control.
- **Relation management.** Prisma's relational model maps directly to the aggregate relationships defined in DDD: a `Course` has many `CourseModules`, each of which has many `Lessons`. The generated client provides type-safe nested writes (`prisma.course.create( data: ..., modules: create: [...] )`).

The alternative considered was TypeORM, which was rejected due to less reliable type inference, decorator-based entity definitions that couple domain entities to the ORM, and a history of breaking changes between major versions.

A.2.4. Redis 7

Redis 7 is used for three distinct purposes:

- **Session store.** Refresh token sessions are stored in Redis with TTL-based expiration (Section 4.3). The `GETDEL` atomic operation prevents token reuse race conditions.
- **Rate limiting.** The `@nestjs/throttler` package with `ThrottlerStorageRedisService` enforces global API rate limits across all application instances.
- **Idempotency cache.** Stripe webhook idempotency keys and AI pipeline callback IDs are cached in Redis with a 24-hour TTL to reject duplicate requests before they reach the database.

Redis was chosen over alternative caching solutions (Memcached, KeyDB) for its support of complex data structures (hashes, sorted sets for leaderboards) and the availability of persistence (AOF + RDB) for the session store.

A.2.5. Docker and Docker Compose

The entire application is containerised using Docker and orchestrated with Docker Compose. The `docker-compose.yml` file defines four services:

```
1 services:
2   app:
3     build: .
4     ports: ['3000:3000']
5     depends_on: [db, redis]
6     environment:
7       DATABASE_URL: postgresql://nexalearn:${DB_PASSWORD}@db:5432/
8         nexalearn
9       REDIS_URL: redis://redis:6379/0
10
11   db:
12     image: postgis/postgis:16-3.4
13     volumes: [postgres_data:/var/lib/postgresql/data]
14     environment:
15       POSTGRES_USER: nexalearn
16       POSTGRES_PASSWORD: ${DB_PASSWORD}
17       POSTGRES_DB: nexalearn
18
19   redis:
20     image: redis:7-alpine
21     volumes: [redis_data:/data]
22
23   nginx:
24     image: nginx:alpine
25     ports: ['80:80', '443:443']
26     volumes:
27       - ./nginx.conf:/etc/nginx/conf.d/default.conf
28       - ./ssl:/etc/nginx/ssl
29     depends_on: [app]
```

Listing A.1: NexaLearn Docker Compose services

Docker Compose was chosen over Kubernetes for the initial deployment due to its simplicity, lower resource overhead, and suitability for single-host deployments. The migration path to Kubernetes is documented in the deployment guide.

A.2.6. Stripe API

Stripe handles all payment processing. The integration spans three areas:

- **Payment Intents API.** The server creates a `PaymentIntent` with the course price and

currency through Stripe's REST API. The `client_secret` is returned to the frontend, which completes the payment using Stripe Elements or Stripe Checkout.

- **Webhook events.** Stripe delivers `payment_intent.succeeded`, `payment_intent.payment_failed`, and `charge.refunded` events to the `/webhooks/stripe` endpoint. Each event is verified using the `Stripe-Signature` header and the webhook signing secret.
- **Idempotency keys.** The Stripe API supports idempotency keys via the `Idempotency-Key` header, which the system uses for retry-safe payment intent creation.

Stripe was chosen over PayPal, Braintree, and Paddle for its developer-friendly API, comprehensive webhook system, and support for the Egyptian market via Stripe Connect.

A.2.7. Mux Video API

Mux provides video uploading, transcoding, and streaming functionality:

- **Direct upload.** The server generates a Mux upload URL via the `/video/uploads` API. The browser uploads the video file directly to Mux, bypassing the server and reducing bandwidth costs.
- **Asset lifecycle.** Mux fires webhooks for asset state transitions: `video.asset.created`, `video.asset.ready`, `video.asset.errorred`, and `video.asset.mp4_rendition.ready`. Each webhook is verified using the `Mux-Signature` header.
- **Signed playback.** Playback IDs are signed using Mux's JWT-based signing keys, restricting access to authorised users. The signing key and key ID are configured via environment variables.

Mux was chosen over self-hosted solutions (FFmpeg + S3 + CloudFront) and alternative providers (Cloudflare Stream, Vimeo API) for its straightforward upload API, reliable webhook delivery, and transparent per-minute pricing.

A.2.8. AWS S3

Amazon S3 is used for persistent file storage outside the video pipeline:

- **Certificate PDFs.** Completed course certificates are generated as PDFs and stored in an S3 bucket with server-side encryption (SSE-S3). Presigned URLs with 1-hour expiration are provided to learners for download.
- **Course thumbnails.** Course cover images are uploaded via presigned POST URLs and served through CloudFront for low-latency delivery.
- **AI pipeline artefacts.** Intermediate files from the AI pipeline—Whisper transcriptions in SRT format, LLM-generated quiz drafts—are stored in S3 with lifecycle policies that transition infrequently accessed objects to S3 Glacier after 90 days.

A.2.9. Resend Email Service

Resend provides transactional email delivery with a modern REST API. The integration covers:

- **Email verification.** A verification email with a signed token link is sent to new users during registration.
- **Password reset.** A password reset email with a time-limited token is sent on request.
- **Notifications.** Email notifications for quiz availability, discussion replies, and certificate issuance.

Resend was chosen over SendGrid, SES, and Mailgun for its TypeScript-native SDK, straightforward template API using React Email components, and generous free tier (100 emails/day).

A.2.10. Jest and Testcontainers

Jest 29 serves as the test framework, complemented by Testcontainers for integration testing:

- **Jest configuration.** Tests are organized in three tiers: unit tests (`*.spec.ts`), integration tests (`*.int-spec.ts`), and E2E tests (`*.e2e-spec.ts`). The `jest.config.ts` file uses `ts-jest` for TypeScript compilation and `collectCoverageFrom` patterns to exclude modules and interfaces from coverage requirements.
- **Testcontainers.** Integration tests use `@testcontainers/postgresql` and `@testcontainers/redis` to spin up disposable Docker containers. This approach was chosen over mocking the Prisma client or using an in-memory SQLite database because: (1) SQL dialect differences between PostgreSQL and SQLite cause false negatives, (2) Prisma migrations produce PostgreSQL-specific SQL, and (3) constraint enforcement (foreign keys, check constraints, unique indexes) differs between databases.

A.2.11. Swagger/OpenAPI

API documentation is auto-generated using the `@nestjs/swagger` package, which produces an OpenAPI 3.1 specification from decorators on controllers and DTOs:

- Each controller class is annotated with `@ApiTags('auth')`, grouping endpoints by bounded context.
- Each endpoint method is annotated with `@ApiOperation( summary: '...', description: '...' )` and `@ApiResponse( status: 201, description: '...', type: ... )`.
- DTO classes use `@ApiProperty()` decorators to document request and response schemas.

- The generated specification is served at `/api/docs` via the Swagger UI interface, allowing interactive API exploration and testing.

Appendix B:

Use Cases

This appendix documents the ten primary use cases of the NexaLearn platform. Each use case is presented in a structured format with an identifier, name, actors, preconditions, main flow, postconditions, and exception conditions.

B.1. Use Case Descriptions

B.1.1. UC-01: Student Signup and Email Verification

Element	Description
ID	UC-01
Name	Student Signup and Email Verification
Actors	Unauthenticated student
Preconditions	The student has an email address and a password that meets the platform's password policy (minimum 8 characters, at least one uppercase letter and one digit).

CHAPTER A: DEVELOPMENT PLATFORMS AND TOOLS

Element	Description
Main Flow	1. The student submits a registration request with full name, email address, and password via POST <code>/auth/register</code>. 2. The system validates the input format (email syntax, password strength) and checks that the email address is not already registered. 3. The system creates a <code>User</code> aggregate in <code>UNVERIFIED</code> status, hashes the password using <code>bcrypt</code> (cost factor 12), and stores the record in PostgreSQL. 4. The system generates a cryptographically random email verification token (64 hex characters), stores its SHA-256 hash in Redis with a 24-hour TTL, and sends an email via Resend containing a verification link: <code>https://nexalearn.app/auth/verify?token=<raw_token></code>. 5. The student clicks the verification link. The system hashes the raw token and retrieves the corresponding user ID from Redis. 6. The system updates the <code>User</code> aggregate to <code>ACTIVE</code> status and returns a success response.
Postconditions	The user is registered and verified. The user can log in using email and password.
Exceptions	• E01: Duplicate email. If the email is already registered, the system returns HTTP 409 Conflict with the message "An account with this email already exists." • E02: Invalid token. If the verification token is expired or malformed, the system returns HTTP 400 Bad Request. • E03: Email delivery failure. If the Resend API returns a non-2xx response, the user is created but remains <code>UNVERIFIED</code>. A resend endpoint (POST <code>/auth/resend-verification</code>) allows the user to request a new token.

B.1.2. UC-02: Instructor Course Creation and Publishing

Element	Description
ID	UC-02
Name	Instructor Course Creation and Publishing
Actors	Authenticated instructor
Preconditions	The instructor is logged in with a valid access token. The instructor's role is INSTRUCTOR.
Main Flow	1. The instructor submits course metadata (title, description, category, difficulty level, price, cover image) via <code>POST /courses</code>. The course is created in DRAFT status. 2. The instructor adds modules to the course via <code>POST /courses/:id/modules</code>. Each module has a title, description, and ordinal position. 3. The instructor adds lessons to each module via <code>POST /modules/:id/lessons</code>. Each lesson has a title, content (HTML/Markdown), optional video asset, and optional quiz. 4. The instructor uploads a video for a lesson via the Mux direct upload flow (see UC-04). 5. The instructor submits a publish request via <code>POST /courses/:id/publish</code>. 6. The system validates the course completeness: at least one module required, each module must have at least one lesson, all lessons must have either text content or a video asset. 7. The system transitions the Course aggregate from DRAFT to PUBLISHED and publishes a CoursePublished domain event through the outbox.
Postconditions	The course is visible in the course catalogue. Students can browse the course and enroll.

Element	Description
Exceptions	• E01: Incomplete course. If any completeness validation fails, the system returns HTTP 422 Unprocessable Entity with a list of missing requirements. • E02: Unauthorised. If the authenticated user's role is not INSTRUCTOR or ADMIN, the system returns HTTP 403 Forbidden.

B.1.3. UC-03: Student Enrollment in Paid Course

Element	Description
ID	UC-03
Name	Student Enrollment in Paid Course (Stripe Payment)
Actors	Authenticated student, Stripe payment gateway
Preconditions	The student is logged in. The course is in PUBLISHED status and has a non-zero price.

CHAPTER A: DEVELOPMENT PLATFORMS AND TOOLS

Element	Description
Main Flow	1. The student requests enrollment via <code>POST /enrollments/:courseId</code>. 2. The system checks that the student is not already enrolled (unique constraint on <code>(user_id, course_id, status)</code> where <code>status != 'CANCELLED'</code>). 3. The system creates a <code>PaymentIntent</code> record in <code>PENDING</code> status and requests a payment intent from Stripe via <code>stripe.paymentIntents.create()</code> with the course price and currency. 4. The system returns the <code>client_secret</code> to the student's frontend. 5. The student completes the payment on the frontend using Stripe Elements or Stripe Checkout. 6. Stripe delivers a <code>payment_intent.succeeded</code> webhook event to <code>POST /webhooks/stripe</code>. 7. The system verifies the webhook signature, checks idempotency via the <code>ProjectionEventLedger</code>, updates the <code>PaymentIntent</code> to <code>SUCCEEDED</code>, creates an <code>Enrollment</code> aggregate in <code>ACTIVE</code> status, and publishes an <code>EnrollmentActivated</code> event.
Postconditions	The student is enrolled in the course with <code>ACTIVE</code> status. The student can access all lessons.
Exceptions	• E01: Already enrolled. If the student is already enrolled, the system returns HTTP 409 Conflict. • E02: Payment failed. If Stripe reports a payment failure, the <code>PaymentIntent</code> transitions to <code>FAILED</code> and the enrollment is created in <code>PENDING_PAYMENT</code> status with a link to retry payment. • E03: Duplicate webhook. If the webhook is delivered multiple times, the <code>ProjectionEventLedger</code> unique constraint prevents duplicate processing.

B.1.4. UC-04: Video Upload and AI Quiz Generation

Element	Description
ID	UC-04
Name	Video Upload and AI Quiz Generation
Actors	Authenticated instructor, Mux Video API, OpenAI Whisper / LLM
Preconditions	The instructor is logged in. A lesson exists in the course.

CHAPTER A: DEVELOPMENT PLATFORMS AND TOOLS

Element	Description
Main Flow	1. The instructor requests a Mux direct upload URL via <code>POST /media/upload</code>. The system calls Mux's <code>/video/uploads</code> API and returns the upload URL to the frontend. 2. The frontend uploads the video file directly to Mux. Mux fires <code>video.asset.created</code> webhook upon upload completion. 3. The system creates a <code>VideoAsset</code> aggregate in <code>UPLOADED</code> status. 4. Mux transcodes the video and fires <code>video.asset.ready</code> (HLS available) and later <code>video.asset.mp4_rendition.ready</code>. 5. When the asset reaches <code>MP4_AVAILABLE</code> status, the system creates a <code>TranscriptionJob</code> in <code>PENDING</code> status and sends the MP4 download URL to the AI pipeline with an HMAC-signed callback URL. 6. The AI pipeline (Whisper) transcribes the audio and POSTs the transcript back to <code>POST /webhooks/ai-pipeline/transcription</code>. The system verifies the HMAC signature, stores the transcript, and creates a <code>QuizGenerationJob</code> in <code>PENDING</code> status. 7. The AI pipeline (LLM) generates quiz questions from the transcript and POSTs the result to <code>POST /webhooks/ai-pipeline/quiz</code>. The system creates a Quiz aggregate in <code>DRAFT</code> status linked to the lesson. 8. The instructor reviews the generated quiz and publishes it (see UC-06).
Postconditions	The video is playable via HLS. A transcript is available. A draft quiz is created for instructor review.

Element	Description
Exceptions	• E01: Transcoding failure. If Mux reports <code>video.asset.errorred</code>, the <code>VideoAsset</code> transitions to <code>ERRORRED</code> and the instructor is notified. • E02: Transcription timeout. If the AI pipeline does not respond within 30 minutes, the <code>TranscriptionJob</code> transitions to <code>TIMEOUT</code> and the instructor is notified to re-submit.

B.1.5. UC-05: Student Takes Quiz Attempt

Element	Description
ID	UC-05
Name	Student Takes Quiz Attempt
Actors	Authenticated student
Preconditions	The student is enrolled in the course. The quiz is in <code>PUBLISHED</code> status and linked to a lesson.

CHAPTER A: DEVELOPMENT PLATFORMS AND TOOLS

Element	Description
Main Flow	1. The student requests the quiz via GET /quizzes/:id. The system returns the quiz questions (without answer keys) and the student's previous attempts, if any. 2. The student submits answers via POST /quizzes/:id/attempts. The system creates a QuizAttempt aggregate in SUBMITTED status. 3. For multiple-choice and true/false questions, the system grades automatically using the stored answer keys. For short-answer questions, the system publishes a ShortAnswerGradeRequestedEvent through the outbox. 4. The AI pipeline grades short-answer questions asynchronously and delivers results via webhook. The QuizAttempt aggregate tracks per-question grading status (PENDING, AUTO_GRADED, AI_GRADED). 5. When all questions are graded, the QuizAttempt transitions to GRADED and the total score and per-question feedback are computed. 6. The student retrieves the graded attempt via GET /attempts/:id and views the score, correct/incorrect markings, and instructor feedback for short-answer questions.
Postconditions	The quiz attempt is graded. The score is recorded in the student's progress.
Exceptions	• E01: Attempt limit exceeded. If the quiz has an attempt limit (e.g., 3 maximum) and the student has exhausted it, the system returns HTTP 403 Forbidden. • E02: Short-answer grading timeout. If AI grading does not complete within 30 minutes, the attempt escalates to REVIEW_REQUIRED for manual instructor intervention.

B.1.6. UC-06: Instructor Reviews AI-Generated Quiz

Element	Description
ID	UC-06
Name	Instructor Reviews AI-Generated Quiz
Actors	Authenticated instructor
Preconditions	The AI pipeline has generated a draft quiz attached to a lesson. The quiz is in DRAFT status.
Main Flow	1. The instructor navigates to the lesson and sees a notification that an AI-generated quiz draft is available for review. 2. The instructor requests the quiz via GET /quizzes/:id (with instructor permissions, which includes the answer keys). 3. The instructor reviews each question, edits the text, changes the question type, adds or removes options, and adjusts the correct answer. 4. The instructor publishes the quiz via POST /quizzes/:id/publish. 5. The system validates that the quiz has at least one question and transitions it to PUBLISHED status.
Postconditions	The quiz is available to students enrolled in the course.
Exceptions	• E01: No questions. If the instructor attempts to publish a quiz with zero questions, the system returns HTTP 422 Unprocessable Entity. • E02: Unauthorised. If the authenticated user is not the course instructor or an admin, the system returns HTTP 403 Forbidden.

B.1.7. UC-07: Student Watches Video with Progress Tracking

Element	Description
ID	UC-07

CHAPTER A: DEVELOPMENT PLATFORMS AND TOOLS

Element	Description
Name	Student Watches Video with Progress Tracking
Actors	Authenticated student, Mux Video CDN
Preconditions	The student is enrolled in the course. The lesson has a video asset in READY status.
Main Flow	1. The student requests the lesson via <code>GET /lessons/:id</code>. The system verifies enrollment and returns the lesson content and a signed Mux playback URL. 2. The student's frontend loads the HLS stream using <code>HLS.js</code> and begins playback. 3. The frontend sends periodic progress updates to <code>PATCH /lessons/:id/progress</code> with the current playback position (in seconds) and the percentage watched. 4. When the student reaches 90% of the video duration (configurable threshold), the frontend sends a completion event. The system creates or updates a <code>LessonProgress</code> record in COMPLETED status. 5. The system publishes a <code>LessonCompleted</code> event, which triggers streak computation (via <code>StreakService</code>) and updates the <code>CourseProgressView</code>.
Postconditions	The lesson is marked as completed in the student's progress. The streak counter is updated. If all lessons in the course are completed, the course is marked as completed.
Exceptions	• E01: Enrollment not active. If the enrollment is <code>CANCELLED</code> or <code>PENDING_PAYMENT</code>, the system denies access and returns HTTP 403 Forbidden. • E02: Video not ready. If the video asset is still processing, the system returns the lesson content without the video URL and a <code>Retry-After</code> header.

B.1.8. UC-08: Discussion Thread Creation and Instructor Reply

CHAPTER A: DEVELOPMENT PLATFORMS AND TOOLS

Element	Description
ID	UC-08
Name	Discussion Thread Creation and Instructor Reply
Actors	Authenticated student (thread creator), authenticated instructor (reply)
Preconditions	The student is enrolled in the course. The course has discussion enabled.
Main Flow	1. The student creates a discussion thread via <code>POST /courses/:id/discussions</code> with a title and body. The system creates a Thread aggregate in OPEN status and publishes a ThreadCreated event. 2. Other enrolled students can view the thread and post replies via <code>POST /discussions/:id/replies</code>. 3. The instructor views the thread and posts a reply that is marked with the <code>IS_INSTRUCTOR</code> flag. The system fires <code>InstructorReplied</code> and sends an email notification to the thread creator. 4. The student can mark the thread as resolved via <code>PATCH /discussions/:id/resolve</code>, which transitions the thread to RESOLVED status.
Postconditions	The thread is created and visible to all enrolled students. Participants can post replies.
Exceptions	• E01: Rate limit exceeded. If the student exceeds the maximum number of threads per hour (configurable), the system returns HTTP 429 Too Many Requests. • E02: Course not enrolled. If the student is not enrolled in the course, the system returns HTTP 403 Forbidden.

B.1.9. UC-09: Certificate Generation After Course Completion

Element	Description
ID	UC-09

CHAPTER A: DEVELOPMENT PLATFORMS AND TOOLS

Element	Description
Name	Certificate Generation After Course Completion
Actors	Authenticated student
Preconditions	The student has completed all lessons in the course (progress = 100%). The course has certificate generation enabled.
Main Flow	1. When the last lesson in the course is marked as completed, the <code>CourseProgressView</code> detects that all lessons are completed and publishes a <code>CourseCompleted</code> event. 2. The <code>Certificates</code> context handles the event and creates a <code>Certificate</code> aggregate in <code>PENDING</code> status. 3. A background job generates a PDF certificate using a templating library (<code>PDFKit</code> or <code>Puppeteer</code>). The PDF includes the student's name, course title, completion date, and a verification QR code. 4. The PDF is uploaded to AWS S3 with server-side encryption. The <code>Certificate</code> aggregate transitions to <code>ISSUED</code> status with an S3 key reference. 5. The student can download the certificate via <code>GET /certificates/:id/download</code>, which generates a presigned S3 URL with 1-hour expiration. 6. The student can share the certificate verification link: <code>https://nexalearn.app/verify/<certificate_id></code>, which displays the certificate details without requiring authentication.
Postconditions	A verifiable PDF certificate is generated and stored. The student can download and share the certificate.
Exceptions	• E01: Certificate already issued. If the student requests a duplicate certificate, the system returns the existing certificate instead of generating a new one. • E02: S3 upload failure. If the PDF upload to S3 fails, the <code>Certificate</code> aggregate remains in <code>PENDING</code> and a retry is scheduled via the outbox.

B.1.10. UC-10: Instructor Views Revenue Analytics Dashboard

Element	Description
ID	UC-10
Name	Instructor Views Revenue Analytics Dashboard
Actors	Authenticated instructor
Preconditions	The instructor is logged in and has at least one published course with completed enrollments.
Main Flow	1. The instructor requests the analytics dashboard via GET /analytics/instructor/dashboard. The request includes query parameters for the time range (?from=...&to=...). 2. The Instructor context queries read models in the Progress and Payments contexts to compute: • Total revenue for the period (sum of succeeded payment intents). • Number of new enrollments per course per day. • Average quiz score across all attempts. • Course completion rate (students who completed / students enrolled). • Top-performing courses by revenue and enrollment count. 3. The system returns the computed metrics as a JSON response. The frontend renders charts and tables from this data. 4. The instructor can filter by course and date range and export the data as CSV.
Postconditions	The instructor can view and export revenue and engagement analytics.

CHAPTER A: DEVELOPMENT PLATFORMS AND TOOLS

Element	Description
Exceptions	• E01: No data. If no enrollments or payments exist for the selected period, the system returns empty metrics with a descriptive message. • E02: Access denied. If the instructor attempts to access analytics for a course they do not own, the system filters the data to include only owned courses.

Appendix C:

User Guide

This appendix provides practical guidance for interacting with the NexaLearn API. All endpoints are accessible via the Swagger UI at `/api/docs` when the application is running.

/api/docs

Swagger UI — NexaLearn API

Auth Module

POST

/auth/register

Register new user

POST

/auth/login

Login

POST

/auth/refresh

Refresh tokens

POST

/auth/verify-email

Verify email

POST

/auth/logout

Logout

Course Module

GET

/courses

List published courses

GET

/courses/{id}

Get course details

POST

/courses

Create course

[AUTH]

PATCH

/courses/{id}

Update course

[AUTH]

Enrollment Module

POST

/enrollments

Enroll in course

[AUTH]

GET

/enrollments

List enrollments

GET

/enrollments/{id}

Get enrollment

Payment Module

POST

/payments/create-intent

Create payment intent

GET

/payments/intents

List intents

POST

/webhooks/stripe

Stripe webhook

Media Module

POST

/media/upload

Request upload URL

[AUTH]

GET

/media/assets/{id}

Get asset status

AI Pipeline

POST

/media/upload?generateQuiz=true

Upload + generate quiz

[AUTH]

GET

/ai/jobs/{id}

Check generation status

Assessment Module

/quizzes/{id}

Get quiz

/quizzes/{id}/attempts

Submit attempt

[AUTH]

/attempts/{id}

Get graded result

/quizzes/{id}/publish

Publish quiz

[AUTH]

Discussion

/courses/{id}/discussions

Create thread

[AUTH]

/discussions/{id}/replies

Post reply

[AUTH]

/discussions/{id}/resolve

Resolve thread

[AUTH]

Figure C.1: Swagger UI documentation at `/api/docs` showing the Auth module endpoints. The interactive UI allows testing endpoints directly from the browser.

C.1. API Authentication Guide

C.1.1. Obtaining JWT Tokens

NexaLearn uses a dual-token authentication system: a short-lived access token (15 minutes) and a long-lived refresh token (7 days).

Step 1: Register an account.

```
1 POST /auth/register
2 Content-Type: application/json
3
4 {
5   "email": "student@example.com",
6   "password": "SecurePass123",
7   "fullName": "Ahmed Hassan",
8   "role": "student"
9 }
```

Listing C.1: Register a new user

Response: HTTP 201 with an activation token sent to the email address.

Step 2: Verify email.

```
1 POST /auth/verify-email
2 Content-Type: application/json
3
4 {
5   "token": "abcdef123456..."
6 }
```

Listing C.2: Verify email address

Step 3: Log in.

```
1 POST /auth/login
2 Content-Type: application/json
3
4 {
5   "email": "student@example.com",
6   "password": "SecurePass123"
7 }
```

Listing C.3: Login and obtain tokens

```
1 HTTP 200 OK
2 Content-Type: application/json
3
```

```
4 {
5   "accessToken": "eyJhbGciOiJSUzI1NiIs... ",
6   "refreshToken": "a1b2c3d4e5f6...",
7   "expiresIn": 900,
8   "user": {
9     "id": "user_abc123",
10    "email": "student@example.com",
11    "role": "student",
12    "fullName": "Ahmed Hassan"
13  }
14 }
```

Listing C.4: Response with tokens

Step 4: Use the access token.

Include the access token in the `Authorization` header for all authenticated requests:

```
1 GET /courses
2 Authorization: Bearer eyJhbGciOiJSUzI1NiIs...
```

Listing C.5: Authenticated request

Step 5: Refresh the access token.

When the access token expires (after 15 minutes), use the refresh token to obtain a new pair:

```
1 POST /auth/refresh
2 Content-Type: application/json
3
4 {
5   "refreshToken": "a1b2c3d4e5f6..."
6 }
```

Listing C.6: Refresh tokens

Important: The refresh token is single-use. Each refresh call invalidates the old token and issues a new one. Store the new refresh token securely (HTTPOnly cookie, secure local storage).

C.1.2. Token Storage Best Practices

- **Web applications:** Store the access token in memory (JavaScript variable). Store the refresh token in an HTTPOnly, Secure, SameSite=Strict cookie set by the server.
- **Mobile applications:** Store both tokens in the device's secure keychain (iOS Keychain, Android Keystore).
- **Never** store tokens in `localStorage` (vulnerable to XSS attacks).

- **Never** log tokens to the console or send them in URLs.

C.2. Instructor Guide

C.2.1. Creating a Course

1. Log in with an instructor account (role: instructor).
2. Send a POST request to `/courses` with course metadata.

```
1 POST /courses
2 Authorization: Bearer <instructor-access-token>
3 Content-Type: application/json
4
5 {
6   "title": "Introduction to Machine Learning",
7   "description": "A comprehensive introduction to ML concepts
8     ...",
9   "categoryId": "cat_ml",
10  "difficulty": "beginner",
11  "price": 4900,
12  "coverImageUrl": "https://nexalearn.s3.amazonaws.com/covers/
    ml-intro.png"
}
```

Listing C.7: Create a course

C.2.2. Adding Modules and Lessons

1. Add modules to the course:

```
1 POST /courses/:courseId/modules
2 Authorization: Bearer <instructor-access-token>
3 Content-Type: application/json
4
5 {
6   "title": "Supervised Learning",
7   "description": "Linear regression, decision trees, SVMs",
8   "order": 1
9 }
```

Listing C.8: Add a module

2. Add lessons to each module:

```
1 POST /modules/:moduleId/lessons
```

```
2 Content-Type: application/json
3
4 {
5   "title": "Linear Regression",
6   "content": "<h2>Introduction</h2><p>Linear regression models
7     ...</p>",
8   "contentType": "html",
9   "order": 1,
10  "estimatedDurationMinutes": 15
11 }
```

Listing C.9: Add a lesson

C.2.3. Uploading Videos

The video upload flow uses Mux's direct upload to avoid proxying video files through the server:

1. Request a Mux upload URL:

```
1 POST /media/upload
2 Authorization: Bearer <instructor-access-token>
3
4 Response:
5 {
6   "uploadUrl": "https://api.mux.com/video/v1/uploads/...",
7   "assetId": "asset_abc123"
8 }
```

Listing C.10: Request upload URL

2. Upload the video file directly to Mux from the frontend using the returned URL.
3. The system will process the video asynchronously. Check the asset status:

```
1 GET /media/assets/:assetId
2
3 Response:
4 {
5   "id": "asset_abc123",
6   "status": "READY",
7   "playbackUrl": "https://stream.mux.com/playback_id.m3u8",
8   "duration": 1842
9 }
```

Listing C.11: Check video status

4. Attach the video to a lesson:

```
1 PATCH /lessons/:lessonId
2 Content-Type: application/json
3
4 {
5   "videoAssetId": "asset_abc123"
6 }
```

Listing C.12: Attach video to lesson

C.2.4. Managing AI-Generated Quizzes

When a video is uploaded with the `generateQuiz: true` flag, the system creates a draft quiz automatically.

1. View the generated quiz draft:

```
1 GET /quizzes/:quizId
```

Listing C.13: View AI-generated quiz

2. Edit questions, add or remove options, change question types.
3. Publish the quiz when ready:

```
1 POST /quizzes/:quizId/publish
```

Listing C.14: Publish quiz

C.3. Student Guide

C.3.1. Enrolling in a Course

1. Browse the course catalogue:

```
1 GET /courses?status=PUBLISHED&page=1&limit=20
```

Listing C.15: List published courses

2. View course details:

```
1 GET /courses/:courseId
```

Listing C.16: View course

3. For free courses, send an enrollment request:

```
1 POST /enrollments
2 Content-Type: application/json
3
4 {
5   "courseId": "course_abc123"
6 }
```

Listing C.17: Enroll in free course

4. For paid courses, the enrollment creates a Stripe Payment Intent. Complete the payment on the frontend using the returned `clientSecret`.

C.3.2. Watching Lessons and Tracking Progress

1. Access enrolled courses:

```
1 GET /enrollments?status=ACTIVE
```

Listing C.18: List enrolled courses

2. View course curriculum with module–lesson hierarchy:

```
1 GET /courses/:courseId/curriculum
```

Listing C.19: View course curriculum

3. Access a lesson (the response includes the video playback URL, quiz link, and discussion thread):

```
1 GET /lessons/:lessonId
```

Listing C.20: Access lesson

4. Send progress updates as the student watches:

```
1 PATCH /lessons/:lessonId/progress
2 Content-Type: application/json
3
4 {
5   "progressPercent": 45,
6   "currentPosition": 320
7 }
```

Listing C.21: Update lesson progress

5. Mark lesson as completed (sent when 90%+ watched):

```
1 POST /lessons/:lessonId/complete
```

Listing C.22: Complete lesson

C.3.3. Taking Quizzes

1. View quiz questions (without answer keys):

```
1 GET /quizzes/:quizId
```

Listing C.23: View quiz

2. Submit an attempt:

```
1 POST /quizzes/:quizId/attempts
2 Content-Type: application/json
3
4 {
5   "answers": [
6     { "questionId": "q_1", "selectedOptionId": "opt_a" },
7     { "questionId": "q_2", "textAnswer": "Supervised learning is..." }
8   ]
9 }
```

Listing C.24: Submit quiz attempt

3. View the graded result:

```
1 GET /attempts/:attemptId
2
3 Response includes: score, percentage, per-question feedback,
4 and instructor comments for short-answer questions.
```

Listing C.25: View graded attempt

C.3.4. Discussion

1. Create a discussion thread:

```
1 POST /courses/:courseId/discussions
2 Content-Type: application/json
3
4 {
5   "title": "Question about linear regression",
6   "body": "I don't understand how gradient descent converges
7   ..."
8 }
```

Listing C.26: Create thread

2. Reply to a thread:

```
1 POST /discussions/:threadId/replies
2 Content-Type: application/json
3
4 {
5   "body": "Gradient descent converges when the learning rate
6     ..."
7 }
```

Listing C.27: Post reply

3. Mark thread as resolved:

```
1 PATCH /discussions/:threadId/resolve
```

Listing C.28: Resolve thread

C.4. Admin Guide

C.4.1. Managing Categories

```
1 POST /admin/categories
2 Authorization: Bearer <admin-access-token>
3 Content-Type: application/json
4
5 {
6   "name": "Machine Learning",
7   "slug": "machine-learning",
8   "parentId": null
9 }
```

Listing C.29: Create category

C.4.2. Viewing Audit Logs

The system logs all state-changing operations in the audit log table. Admins can query logs by date range, user, or resource:

```
1 GET /admin/audit-logs?from=2025-01-01&to=2025-06-01&userId=user_123
2
3 Response:
4 [
5   {
6     "timestamp": "2025-03-15T10:30:00Z",
7     "userId": "user_123",
```

```
8   "action": "COURSE_PUBLISHED",
9   "resourceId": "course_abc",
10  "details": { "courseTitle": "ML_Intro" }
11 }
12 ]
```

Listing C.30: Query audit logs

C.4.3. Managing Instructor Applications

When a new instructor signs up, their account is created with role `student` by default. An admin can promote the user:

```
1 PATCH /admin/users/:userId/role
2 Content-Type: application/json
3
4 {
5   "role": "instructor"
6 }
```

Listing C.31: Promote user to instructor

Appendix D:

Code Documentation

This appendix documents the key design patterns, domain entities, and use case implementations that form the backbone of the NexaLearn codebase. All code examples are written in TypeScript and follow the NestJS framework conventions.

D.1. Key Design Patterns

D.1.1. Transactional Outbox Publisher

The outbox pattern guarantees at-least-once event delivery. The `OutboxPublisher` service is injected into application services that need to publish domain events atomically with database state changes.

```
1 @Injectable()
2 export class OutboxPublisher {
3   constructor(private readonly prisma: PrismaService) {}
4
5   async publish<T extends Record<string, unknown>>(<
6     aggregateType: string,
7     aggregateId: string,
8     eventType: string,
9     payload: T,
10    transaction?: Prisma.TransactionClient,
11  >): Promise<void> {
12    const client = transaction ?? this.prisma;
13    await client.outboxEvent.create({
14      data: {
15        aggregateType,
16        aggregateId,
17        eventType,
18        payload,
19        status: 'PENDING',
20      },
21    });
22  }
23 }
```

Listing D.1: Transactional outbox publisher

The usage pattern within an application service:

```

1  async createEnrollment(userId: string, courseId: string): Promise<
    Enrollment> {
2      return this.prisma.$transaction(async (tx) => {
3          const enrollment = await tx.enrollment.create({
4              data: { userId, courseId, status: 'ACTIVE' },
5          });
6
7          await this.outboxPublisher.publish(
8              'Enrollment',
9              enrollment.id,
10             'EnrollmentActivated',
11             { enrollmentId: enrollment.id, userId, courseId },
12             tx,
13         );
14
15         return enrollment;
16     });
17 }

```

Listing D.2: Using outbox in a use case

D.1.2. Projection Event Ledger for Idempotency

The `tryBeginProjectionEvent` pattern ensures that each external callback (Stripe webhook, Mux webhook, AI pipeline callback) is processed exactly once.

```

1  @Injectable()
2  export class ProjectionEventLedger {
3      constructor(private readonly prisma: PrismaService) {}
4
5      async tryBegin(
6          eventId: string,
7          handlerName: string,
8      ): Promise<boolean> {
9          try {
10             await this.prisma.projectionEventLedger.create({
11                 data: {
12                     eventId,
13                     handlerName,

```

```

14         status: 'PROCESSING',
15         startedAt: new Date(),
16     },
17     });
18     return true;
19 } catch (err) {
20     if (err instanceof Prisma.PrismaClientKnownRequestError
21         && err.code === 'P2002') {
22         return false;
23     }
24     throw err;
25 }
26 }
27
28 async complete(
29     eventId: string,
30     handlerName: string,
31 ): Promise<void> {
32     await this.prisma.projectionEventLedger.update({
33         where: { eventId_handlerName: { eventId, handlerName } },
34         data: { status: 'COMPLETED', completedAt: new Date() },
35     });
36 }
37 }

```

Listing D.3: Projection event ledger

D.1.3. Optimistic Concurrency Utility

The `withOptimisticRetry` utility wraps database operations that may fail due to optimistic concurrency version conflicts:

```

1 export async function withOptimisticRetry<T>(  
2     operation: () => Promise<T>,
3     maxRetries = 3,
4 ): Promise<T> {
5     for (let attempt = 1; attempt <= maxRetries; attempt++) {
6         try {
7             return await operation();
8         } catch (err) {
9             if (
10                 err instanceof Prisma.PrismaClientKnownRequestError
11                 && err.code === 'P2025'

```

```
12     && attempt < maxRetries
13   ) {
14     continue;
15   }
16   throw err;
17 }
18 }
19 throw new ConcurrencyConflictException('Max_retries_exceeded');
20 }
```

Listing D.4: Optimistic concurrency retry

D.2. Domain Entity Design

The QuizEntity aggregate root illustrates the DDD tactical patterns used throughout the codebase. The entity encapsulates its state and enforces invariants through public methods.

```
1 @Entity()
2 export class QuizEntity {
3   private constructor(
4     private readonly id: string,
5     private courseId: string,
6     private lessonId: string,
7     private title: string,
8     private status: QuizStatus,
9     private questions: QuestionEntity[],
10    private readonly createdAt: Date,
11    private updatedAt: Date,
12    private publishedAt: Date | null,
13    private attemptLimit: number,
14  ) {}
15
16  static create(props: CreateQuizProps): QuizEntity {
17    return new QuizEntity(
18      crypto.randomUUID(),
19      props.courseId,
20      props.lessonId,
21      props.title,
22      QuizStatus.DRAFT,
23      [],
24      new Date(),
25      new Date(),
26      null,
```

```
27     props.attemptLimit ?? 3,
28     );
29 }
30
31 addQuestion(question: QuestionEntity): void {
32     if (this.status !== QuizStatus.DRAFT) {
33         throw new DomainException(
34             'Cannot add questions to a non-draft quiz',
35         );
36     }
37     this.questions.push(question);
38     this.updatedAt = new Date();
39 }
40
41 publish(): void {
42     if (this.status !== QuizStatus.DRAFT) {
43         throw new DomainException(
44             'Only draft quizzes can be published',
45         );
46     }
47     if (this.questions.length === 0) {
48         throw new DomainException(
49             'Cannot publish a quiz with no questions',
50         );
51     }
52     this.status = QuizStatus.PUBLISHED;
53     this.publishedAt = new Date();
54     this.updatedAt = new Date();
55 }
56
57 archive(): void {
58     if (this.status === QuizStatus.ARCHIVED) {
59         throw new DomainException('Quiz is already archived');
60     }
61     this.status = QuizStatus.ARCHIVED;
62     this.updatedAt = new Date();
63 }
64
65 getStatus(): QuizStatus {
66     return this.status;
67 }
```

```

68
69   getQuestions(): readonly QuestionEntity[] {
70       return this.questions;
71   }
72
73   toPersistent(): QuizPersistent {
74       return {
75           id: this.id,
76           courseId: this.courseId,
77           lessonId: this.lessonId,
78           title: this.title,
79           status: this.status,
80           createdAt: this.createdAt,
81           updatedAt: this.updatedAt,
82           publishedAt: this.publishedAt,
83           attemptLimit: this.attemptLimit,
84       };
85   }
86 }

```

Listing D.5: QuizEntity aggregate root

Key design decisions in this entity:

- **Private constructor.** The entity cannot be instantiated directly; it must be created via the static `create()` factory method or reconstituted from persistence via a repository.
- **Immutable question list.** The `getQuestions()` method returns a `readonly` array to prevent external mutation of the aggregate's internal state.
- **State machine enforcement.** The `publish()` method checks both the current status (only `DRAFT` can be published) and the business invariant (at least one question required).
- **Separation of domain and persistence.** The `toPersistent()` method produces a plain object suitable for Prisma, while the entity itself remains free of ORM decorators.

D.3. Use Case Structure

Use cases follow the NestJS application service pattern. Each use case is a single class with an `execute()` method that accepts a typed DTO and returns a typed result. The `HandleGenerationCallbackUseCase` demonstrates the pattern for webhook-driven AI pipeline interactions.

```

1  @Injectable()
2  export class HandleGenerationCallbackUseCase {

```

```
3 constructor(
4   private readonly quizRepository: QuizRepository,
5   private readonly generationJobRepository:
6     GenerationJobRepository,
7   private readonly ledger: ProjectionEventLedger,
8   private readonly outbox: OutboxPublisher,
9   private readonly prisma: PrismaService,
10 ) {}
11
12 async execute(dto: CallbackDto): Promise<void> {
13   const eventId = dto.jobId;
14   const handlerName = 'HandleGenerationCallback';
15
16   const isNew = await this.ledger.tryBegin(eventId, handlerName);
17   if (!isNew) {
18     return;
19   }
20
21   try {
22     const job = await this.generationJobRepository.findById(
23       dto.jobId,
24     );
25
26     if (!job) {
27       throw new NotFoundException(
28         `Generation job ${dto.jobId} not found`,
29       );
30     }
31
32     job.complete(dto.generatedQuiz);
33
34     const quiz = QuizEntity.create({
35       courseId: job.getCourseId(),
36       lessonId: job.getLessonId(),
37       title: `AI-Generated Quiz: ${job.getLessonTitle()}`,
38     });
39
40     for (const q of dto.generatedQuiz.questions) {
41       quiz.addQuestion(
42         QuestionEntity.create({
43           text: q.text,
```

```
43         type: q.type,
44         options: q.options,
45         correctAnswer: q.correctAnswer,
46     })),
47 );
48 }
49
50 await this.prisma.$transaction(async (tx) => {
51     await this.generationJobRepository.save(job, tx);
52     await this.quizRepository.save(quiz, tx);
53     await this.outbox.publish(
54         'Quiz',
55         quiz.id,
56         'QuizGeneratedByAI',
57         { quizId: quiz.id, lessonId: quiz.lessonId },
58         tx,
59     );
60 });
61
62 await this.ledger.complete(eventId, handlerName);
63 } catch (err) {
64     await this.ledger.fail(eventId, handlerName, err.message);
65     throw err;
66 }
67 }
68 }
```

Listing D.6: HandleGenerationCallbackUseCase

The use case pattern demonstrates several architectural principles:

- **Dependency injection.** All dependencies are injected through the constructor, making the use case testable with mocked repositories.
- **Idempotent entry.** The `ProjectionEventLedger.tryBegin()` guard ensures that duplicate callbacks are silently ignored.
- **Transaction boundary.** The repository saves and the outbox publish occur within the same database transaction, ensuring atomicity.
- **Error handling.** Failures are recorded in the ledger with the error message, enabling observability and manual retry.
- **Business logic delegation.** The use case delegates state mutations to domain entities (`job.complete()`, `quiz.addQuestion()`) rather than manipulating data directly.

Appendix E:

Feasibility Study

This appendix presents a multi-dimensional feasibility analysis of the NexaLearn project, examining technical, operational, economic, and schedule considerations, along with a risk assessment.

E.1. Technical Feasibility

The project is technically feasible for the following reasons:

- **Mature technology stack.** All technologies chosen (Node.js 22 LTS, NestJS 11, PostgreSQL 16, Prisma 6, Redis 7, Docker 27) are stable, well-documented, and widely adopted in production environments. Node.js has been the dominant server-side JavaScript runtime for over a decade, and NestJS has accumulated 75,000+ GitHub stars since its initial release in 2017.
- **Open-source licensing.** Every software component used in the project—from the runtime and framework to the database and caching layer—is open-source and free to use. No paid licenses are required for development or production deployment.
- **Team competency.** The project team possesses the required skills: TypeScript/Node.js development, relational database design, REST API design, Docker containerisation, and Git version control. The DDD and Clean Architecture patterns were learned through literature and applied iteratively during development.
- **Third-party API maturity.** Stripe, Mux, AWS S3, and Resend are well-established services with comprehensive SDKs, documentation, and free tiers suitable for development and small-scale deployment.
- **Cloud deployment readiness.** The Docker Compose-based deployment model is directly compatible with all major cloud providers (AWS ECS, Google Cloud Run, Azure Container Instances, DigitalOcean App Platform) and can be migrated to Kubernetes with minimal changes.

E.2. Operational Feasibility

The project is operationally feasible for the following reasons:

- **Containerised deployment.** The entire application, including the database and cache, runs in Docker containers managed by Docker Compose. This eliminates environment-specific configuration issues and simplifies deployment to any Linux server with Docker installed.
- **Automated CI/CD.** The GitHub Actions CI pipeline runs tests on every pull request, builds the Docker image, and pushes it to a container registry. Production deployment can be automated with a single command: `docker compose up -d`.
- **Monitoring and observability.** NestJS provides built-in logging via the `Logger` service, and the application exposes health check endpoints (`GET /health`) for container orchestration systems. Structured logging in JSON format is compatible with log aggregation tools (ELK Stack, Grafana Loki).
- **Backup and restore.** PostgreSQL backups are automated via `pg_dump` with daily cron jobs. Redis persistence is configured with AOF (Append-Only File) for crash recovery. The S3 bucket for certificate PDFs has cross-region replication enabled.
- **Scalability path.** The modular monolith architecture supports vertical scaling (larger server) for the initial deployment and horizontal scaling (multiple application instances behind a load balancer) for growth. The outbox pattern and Redis-backed rate limiting are designed to work across multiple instances.

E.3. Economic Feasibility

E.3.1. Development Costs

Development costs were limited to hardware and infrastructure:

- **Development hardware:** Existing laptops (MacBook Pro, Dell XPS) at no incremental cost.
- **Infrastructure:** GitHub Actions free tier (2,000 minutes/month, sufficient for the project). Docker Desktop free tier.
- **Third-party services:** Stripe test mode (free). Mux development tier (free up to 1,000 minutes of encoded video). Resend free tier (100 emails/day). AWS S3 free tier (5 GB storage, 20,000 GET requests).
- **Total development cost:** Approximately \$0 (using free tiers of all services).

E.3.2. Production Infrastructure Costs

Table E.1 estimates the monthly infrastructure cost for a medium-scale deployment serving 1,000 active learners.

Table E.1: Estimated Monthly Infrastructure Costs (Medium Deployment)

Component	Specification	Monthly Cost	
VPS (application + DB + cache)	4 vCPU, 8 GB RAM, 100 GB SSD	\$80	Hetzner CX32 or
Domain name	nexalearn.app	\$2	A
SSL certificate	Let's Encrypt	\$0	Automated
Mux Video	500 min encoded / 5 TB served	\$50	Estimated for 1,0
AWS S3	10 GB storage, 50 GB transfer	\$2	Certificate PDF
Resend Email	5,000 emails/month	\$15	Transacti
Stripe	2.9% + \$0.30 per transaction	Variable	Deducted fro
Total (excl. Stripe)		\$149	

At an average course price of \$49 and 200 enrollments per month, the gross revenue would be approximately \$9,800/month, with Stripe fees of approximately \$320/month. The net revenue comfortably covers the infrastructure costs.

E.3.3. Licensing

NexaLearn is released under the MIT open-source license. There are no paid framework licenses, royalty fees, or per-seat licensing costs. The MIT license permits commercial use, modification, and redistribution, enabling both self-hosted deployments and commercial SaaS offerings.

E.4. Schedule Feasibility

The project was completed over two academic semesters (approximately 10 months). Table E.2 presents the timeline with phases, milestones, and deliverables.

CHAPTER E: FEASIBILITY STUDY

Table E.2: Project Timeline and Milestones

Phase	Start	End	Deliverables
Semester 1			
Literature review	Week 1	Week 4	Survey of e-learning platforms, DDD literature review, technology comparison
Requirements analysis	Week 3	Week 6	Problem statement, target customer segments, use case specification
Architecture design	Week 5	Week 8	System architecture diagram, module dependency graph, state machine designs
Database schema design	Week 7	Week 10	Prisma schema (schema.prisma), migration plan, seed data scripts
Auth module implementation	Week 9	Week 12	User registration, JWT token issuance, Redis session store, login/logout
Course module implementation	Week 11	Week 14	Course CRUD, module/lesson hierarchy, category management, publishing flow
Media module implementation	Week 13	Week 16	Mux upload integration, video asset lifecycle, HLS playback URLs
Assessment module (basic)	Week 15	Week 18	Quiz CRUD, question types, manual quiz attempt grading
Semester 2			
AI pipeline integration	Week 19	Week 22	Whisper transcription, LLM quiz generation, HMAC callback auth
Payment module	Week 21	Week 24	Stripe payment intents, webhook processing, enrollment activation
Advanced assessment	Week 23	Week 26	Short-answer AI grading, per-question grading status, async grading flow
Discussion module	Week 25	Week 27	Thread creation, replies, instructor roles, rate limiting
Certificate module	Week 26	Week 28	PDF generation, S3 storage, verification URLs
Instructor analytics	Week 27	Week 28	Revenue dashboard, enrollment metrics, quiz score aggregation
Testing (unit + integration)	Week 20-28	Week 28	Jest test suite, Testcontainer setup, CI pipeline

The schedule was realistic: core modules (Auth, Courses) were completed in the first semester, while integration-heavy modules (AI Pipeline, Payments, Assessments with async grading) were completed in the second semester. Testing was conducted in parallel with development, following the schedule in Section 5.3.

E.5. Risk Analysis

Table E.3 identifies and analyses the key risks to project success, along with mitigation strategies.

Table E.3: Risk Assessment Matrix

Risk	Probability	Impact	Mitigation
R1: Third-party API breaking changes	Low	High	API version pinning (Stripe: 2024-06-20, Mux: v1). All external API calls are wrapped in adapter classes that can be updated independently of domain logic. Integration tests verify API contract compliance.
R2: Database performance degradation under load	Medium	High	PostgreSQL connection pooling via PgBouncer. Index analysis via <code>EXPLAIN ANALYZE</code> . Query optimisation using Prisma's <code>raw</code> queries for complex joins. Read replicas planned for analytics queries.
R3: AI pipeline latency or failure	Medium	High	Asynchronous processing with timeout (30 minutes). Dead-letter queue for failed jobs. Instructor notification and manual retry. Fallback to manual quiz creation if AI pipeline is unavailable.
R4: Security vulnerability (JWT theft, XSS, SQL injection)	Medium	Critical	HS256 JWT with short expiration (15 min). Refresh token rotation with reuse detection. Prisma parameterised queries prevent SQL injection. NestJS Helmet middleware for HTTP security headers. Regular dependency updates with Dependabot.

Risk Assessment Matrix (continued)

Risk	Probability	Impact	Mitigation
R5: Scope creep (feature expansion beyond available time)	High	Medium	Strict prioritisation: core DDD modules first (Auth, Courses, Assessment), AI and payments second, analytics and certificates third. Use cases scoped to minimum viable product. Non-critical features deferred to future work.
R6: Team member unavailability	Low	High	Git-based collaboration allows asynchronous contribution. Module ownership is documented in the architecture guide. Code review requirements ensure knowledge sharing across the team.
R7: Docker image size and cold start time	Medium	Low	Multi-stage Docker build with distroless production image (final image size: 180 MB). Node.js 22 startup time under 2 seconds. Docker layer caching in CI pipeline.
R8: Stripe/Mux webhook delivery reliability	Medium	Medium	At-least-once handling with idempotency via ProjectionEventLedger. Retry with exponential backoff (3 attempts). Dead-letter queue for events that exceed max retries. Manual replay available via admin dashboard.